

A Gentle Introduction to Isabelle and Isabelle HOL

Gunnar Teege

March 6, 2023

Contents

1	Isabelle System	4
1.1	Invoking Isabelle	4
1.1.1	Installation and Configuration	4
1.1.2	Theories and Sessions	5
1.1.3	Invocation as Editor	5
1.1.4	Invocation for Batch Processing	6
1.1.5	Invocation for Document Creation	7
1.2	Interactively Working with Isabelle	7
1.2.1	The Text Area	7
1.2.2	The Sidekick Panel	7
1.2.3	The Output Panel	8
1.2.4	The State Panel	8
1.2.5	The Symbols Panel	8
1.2.6	The Documentation Panel	8
1.2.7	The Query Panel	8
1.2.8	The Theories Panel	8
2	Isabelle Basics	9
2.1	Isabelle Theories	9
2.1.1	Theory Structure	9
2.1.2	Types	10
2.1.3	Terms	11
2.1.4	Definitions and Abbreviations	13
2.1.5	Overloading	14
2.1.6	Propositions	15
2.1.7	Theorems	18
2.1.8	Locales	21
2.2	Isabelle Proofs	23
2.2.1	Proof Context	23
2.2.2	Proof Procedure	24
2.2.3	Basic Proof Structure	26
2.2.4	Backward Reasoning Steps	29
2.2.5	Forward Reasoning Steps	30

2.2.6	Facts as Proof Input	32
2.2.7	Fact Chaining	33
2.2.8	Assuming Facts	36
2.2.9	Fixing Variables	39
2.2.10	Obtaining Variables	41
2.2.11	Term Abbreviations	43
2.2.12	Accumulating Facts	45
2.2.13	Equational Reasoning	47
2.3	Proof Methods	48
2.3.1	The empty Method	49
2.3.2	Rule Application	50
2.3.3	Composed Proof Methods	56
2.3.4	The Simplifier	57
2.3.5	Other Automatic Proof Methods	63
2.4	Case Based Proofs	64
2.4.1	Goal Cases	64
2.4.2	The <i>goal.cases</i> Method	66
2.4.3	Case Based Reasoning	67
2.4.4	Induction	70
3	Isabelle HOL Basics	78
3.1	Algebraic Types	78
3.1.1	Definition of Algebraic Types	78
3.1.2	Constructors	79
3.1.3	Destructors	81
3.1.4	Rules	84
3.2	Record Types	85
3.3	Subtypes	85
3.3.1	Subtype Definitions	86
3.3.2	Subtype Rules	86
3.4	Quotient Types	87
3.5	Type Independent Mechanisms	87
3.5.1	Undefined Value	87
3.5.2	Let Terms	87
4	Isabelle HOL Types	88
4.1	Boolean Values	88
4.1.1	Conditional Terms	88
4.1.2	Logic Rules	88
4.2	Natural Numbers	89
4.3	Tuple Types	89
4.3.1	Tuple Values	89
4.4	Optional Values	89
4.5	Sets	89

4.6	Lists	89
4.7	Fixed-Size Binary Numbers	89

Chapter 1

Isabelle System

Isabelle is a “proof assistant” for formal mathematical proofs. It supports a notation for propositions and their proofs, it can check whether a proof is correct, and it can even help to find a proof.

This introductory manual explains how to work with Isabelle to develop mathematical models. It does not presume prior knowledge about formal or informal proof techniques. It only assumes that the reader has a basic understanding of mathematical logics and the task of proving mathematical propositions.

1.1 Invoking Isabelle

After installation, Isabelle can be invoked interactively as an editor for entering propositions and proofs, or it can be invoked noninteractively to check a proof and generate a PDF document which displays the propositions and proofs.

1.1.1 Installation and Configuration

Isabelle is freely available from <https://isabelle.in.tum.de/> for Windows, Mac, and Linux. It is actively maintained, there is usually one release every year. Older releases are available in a distribution archive.

To install Isabelle, follow the instructions on

`https://isabelle.in.tum.de/installation.html`

Although there are many configuration options, there is no need for an initial configuration, interactive and noninteractive invocation is immediately possible.

1.1.2 Theories and Sessions

The propositions and proofs in Isabelle notation are usually collected in “theory files” with names of the form `name.thy`. A theory file must import at least one other theory file to build upon its content. For theories based on higher order logic (“HOL”), the usual starting point to import is the theory *Main*.

Several theory files can be grouped in a “session”. A session is usually stored in a directory in the file system. It consists of a file named `ROOT` which contains a specification of the session, and the theory files which belong to the session.

When Isabelle loads a session it loads and checks all its theory files. Then it can generate a “heap file” for the session which contains the processed session content. The heap file can be reloaded by Isabelle to avoid the time and effort for processing and checking the theory files.

A session always has a single parent session, with the exception of the Isabelle builtin session *Pure*. Thus, every session depends on a linear sequence of ancestor sessions which begins at *Pure*. The ancestor sessions have separate heap files. A session is always loaded together with all ancestor sessions.

Every session has a name of the form `chap/sess` where `chap` is an arbitrary “chapter name”, it defaults to `Unsorted`. The session name and the name of the parent session are specified in the `ROOT` file in the session directory. When a session is loaded by Isabelle, its directory and the directories of all ancestor sessions must be known by Isabelle.

The Isabelle distribution provides heap files for the session `HOL/HOL` and its parent session `Pure/Pure`, the session directories are automatically known.

Every session may be displayed in a “session document”. This is a PDF document generated by translating the content of the session theory files to \LaTeX . A frame \LaTeX document must be provided which includes all content generated from the theory files. The path of the frame document, whether a session document shall be generated and which theories shall be included is specified in the `ROOT` file.

The command

```
isabelle mkroot [OPTIONS] [Directory]
```

can be used to initialize the given directory (default is the current directory) as session directory. It creates an initial `ROOT` file to be populated with theory file names and other specification for the session, and it creates a simple frame \LaTeX document.

1.1.3 Invocation as Editor

Isabelle is invoked for editing using the command

```
isabelle jedit [OPTIONS] [Files ...]
```

It starts an interactive editor and opens the specified theory files. If no file is specified it opens the file `Scratch.thy` in the user's home directory. If that file does not exist, it is created as an empty file.

The editor also loads a session (together with its ancestors), the default session to load is `HOL`. If a heap file exists for the loaded session it is used, otherwise a heap file is created by processing all the session's theories.

The default session to load can be changed by the option

```
-l <session name>
```

Moreover the editor also loads (but does not open) theories which are transitively imported by the opened theory files. If these are Isabelle standard theories it finds them automatically. If they belong to the session in the current directory it also finds them. If they belong to other sessions, the option

```
-d <directory pathname>
```

must be used to make the session directory known to Isabelle. For every used session a separate option must be specified.

If an imported theory belongs to the loaded session or an ancestor, it is directly referenced there. Otherwise the theory file is loaded and processed.

1.1.4 Invocation for Batch Processing

Isabelle is invoked for batch processing of all theory files in one or more sessions using the command

```
isabelle build [OPTIONS] [Sessions ...]
```

It loads all theory files of the specified sessions and checks the contained proofs. It also loads all required ancestor sessions. If not known to Isabelle, the corresponding session directories must be specified using option `-d` as described in Section 1.1.3. Sessions required for other sessions are loaded from heap files if existent, otherwise the corresponding theories are loaded and a heap file is created.

If option `-b` is specified, heap files are also created for all sessions specified in the command. Option `-c` clears the specified sessions (removes their heap files) before processing them. Option `-n` omits the actual session processing, together with option `-c` it can be used to simply clear the heap files.

The specified sessions are only processed if at least one of their theory files has changed since the last processing or if the session is cleared using option

-c. If option `-v` is specified all loaded sessions and all processed theories are listed on standard output.

If specified for a session in its `ROOT` file (see Section 1.1.5), also the session document is generated when a session is processed.

1.1.5 Invocation for Document Creation

`** todo **`

1.2 Interactively Working with Isabelle

After invoking Isabelle as editor (see Section 1.1.3) it supports interactive work with theories.

The user interface consists of a text area which is surrounded by docking areas where additional panels can be displayed. Several panels can be displayed in the same docking area, using tabs to switch among them. Panels may also be displayed as separate undocked windows.

A panel can be displayed by selecting it in the `Plugins -> Isabelle` menu. Some of the panels are described in the following sections.

1.2.1 The Text Area

The text area displays the content of an open theory file and supports editing it. The font (size) used for display can be configured through the menu in `Utilities -> Global Options -> jEdit -> Text Area` together with many other options for display.

Moreover, in the default configuration, Isabelle automatically processes the theory text up to the current part visible in the text area window. This includes processing the content of all imported theory files, if the import statement is visible.

Whenever the window is moved forward the processing is continued if “Continuous checking” has not been disabled in the Theories panel. Whenever the content of the text area is modified the processing is set back and restarted at the modified position.

In the default configuration the progress of the processing is shown by shading the unprocessed text in red and by a bar on the right border of the text area which symbolizes the whole theory file and shows the unprocessed part by red shading as well.

1.2.2 The Sidekick Panel

`**todo**`

1.2.3 The Output Panel

The Output panel displays the result of the theory text processing when it reaches the cursor position in the text area.

The displayed information depends on the cursor position and may be an information about the current theorem or proof or it may be an error message.

1.2.4 The State Panel

The State panel displays a specific result of the theory text processing if the cursor position is in a proof. It is called the “goal state” (see Section 2.2.2), and describes what remains to be proven by the rest of the proof.

The Output panel can be configured to include the goal state in its display by checking the “Proof state” button.

1.2.5 The Symbols Panel

Isabelle uses a large set of mathematical symbols and other special symbols which are usually not on the keyboard. The Symbols panel can be used to input such symbols in the text area.

1.2.6 The Documentation Panel

A comprehensive set of documentation documents about Isabelle can be opened through the Documentation Panel. This manual refers to some of these documentations, if applicable.

For example, more information about the use of the interactive editor can be found in the Isabelle Reference Manual about `jedit`.

1.2.7 The Query Panel

**** todo****

1.2.8 The Theories Panel

The Theories panel displays the loaded session and the opened or imported theories which do not belong to the loaded session or its ancestors. The (parts of) theories which have not been processed are shaded in red.

If the check button next to a theory is checked, the theory file is processed independently of the position of the text area window.

Chapter 2

Isabelle Basics

The basic mechanisms of Isabelle mainly support defining types, constants, and functions and specifying and proving statements about them.

2.1 Isabelle Theories

A theory is the content of an Isabelle theory file.

2.1.1 Theory Structure

The content of a theory file has the structure

```
theory name
imports name1 ... namen
begin
  ...
end
```

where *name* is the theory name and *name*₁ ... *name*_{*n*} are the names of the imported theories. The theory name *name* must be the same which is used for the theory file, i.e., the file name must be *name.thy*.

The theory structure is a part of the Isabelle “outer syntax” which is mainly fixed and independent from the specific theories. Other kind of syntax is embedded into the outer syntax. The main embedded syntax is the “inner syntax” which is mainly used to denote types and terms. Content in inner syntax must always be surrounded by double quotes. The only exception is a single isolated identifier, for it the quotes may be omitted.

This introduction describes only a selected part of the outer syntax. The full outer syntax is described in the Isabelle/Isar Reference Manual.

Additionally, text written in \LaTeX syntax can be embedded into the outer syntax using the form `text(...)` and \LaTeX sections can be created using

chapter(...), **section**(...), **subsection**(...), **subsubsection**(...), **paragraph**(...), **subparagraph**(...). Note that the delimiters used here are not the “lower” and “greater” symbols, but the “cartouche delimiters” available in the editor’s Symbols panel in tab “Punctuation”.

Embedded \LaTeX text is intended for additional explanations of the formal theory content. It is displayed in the session document together with the formal theory content.

It is also possible to embed inner and outer syntax in the \LaTeX syntax (see Chapter 4 in the Isabelle/Isar Reference Manual).

Moreover, comments of the form

```
(* ... *)
```

can be embedded into the outer syntax. They are only intended for the reader of the theory file and are not displayed in the session document.

2.1.2 Types

As usual in formal logics, the basic building blocks of propositions are terms. Terms denote arbitrary objects like numbers, sets, functions, or boolean values. Isabelle is strongly typed, so every term must have a type. However, in most situations Isabelle can derive the type of a term automatically, so that it needs not be specified explicitly. Terms and types are always denoted using the inner syntax.

Types are usually specified by type names. In Isabelle HOL (see Chapter 3) there are predefined types such as *nat* and *bool* for natural numbers and boolean values. New types can be defined in the form

```
typedecl name
```

which introduces the *name* for a new type for which the values are different from the values of all existing types and the set of values is not empty. No other information about the values is given, that must be done separately. See Chapter 3 for ways of defining types with specifying more information about their values.

Types can be parameterized, then the type arguments are denoted *before* the type name, such as in *nat set* which is the type of sets of natural numbers. A type name with *n* parameters is declared in the form

```
typedecl ('name1, ..., 'namen') name
```

such as **typedecl** ('*a*) *set*. The type parameters are denoted by “type variables” which always have the form '*name* with a leading single quote character. Every use where the parameters are replaced by actual types, such as in *nat set*, is called an “instance” of the parameterized type.

Alternatively a type name can be introduced as a synonym for an existing type in the form

```
type_synonym name = type
```

such as in `type_synonym natset = nat set`. Type synonyms can also be parameterized as in

```
type_synonym ('name1, ..., 'namen) name = type
```

where the type variables occur in *type* in place of actual type specifications.

2.1.3 Terms

Constants and Variables

Terms are mainly built as syntactical structures based on constants and variables. Constants are usually denoted by names, using the same namespace as type names. Whether a name denotes a constant or a type depends on its position in a term. Predefined constant names of type *bool* are *True* and *False*.

Constants of number types, such as *nat*, may be denoted by number literals, such as *6* or *42*.

A constant can be defined by specifying its type. The definition

```
consts name1 :: type1 ... namen :: typen
```

introduces *n* constants with their names and types. No information is specified about the constant's values, in this respect the constants are "under-specified". The information about the values must be specified separately.

If the constant's type contains type variables the constant is called "polymorphic". Thus the declaration

```
consts myset :: "'a set"
```

declares the polymorphic constant *myset* which may be a set of elements of arbitrary type.

A (term) variable has the same form as a constant name, but it has not been introduced as a constant. Whenever a variable is used in a term it has a specific type which is either derived from its context or is explicitly specified in the form *varname* :: *type*.

Nested terms are generally written by using parentheses (...). There are many priority rules how to nest terms automatically, but if in doubt, it is always safe to use parentheses.

Functions

A constant name denotes an object, which, according to its type, may also be a function of arbitrary order. Functions basically have a single argument. The type of a function is written as *argtype* \Rightarrow *restype*.

Functions in Isabelle are always total, i.e., they map every value of type *argtype* to some value of type *restype*. However, a function may be “underspecified” so that no information is (yet) available about the result value for some or all argument values. A function defined by

```
consts mystery :: nat  $\Rightarrow$  nat
```

is completely underspecified: although it maps every natural number to a unique other natural number no information about these numbers is available. Functions may also be partially specified by describing the result value only for some argument values. This does not mean that the function is “partial” and has no value for the remaining arguments. The information about these values may always be provided later, this does not “modify” the function, it only adds information about it.

Functions with Multiple Arguments

The result type of a function may again be a function type, then it may be applied to another argument. This is used to represent functions with more than one argument. Function types are right associative, thus a type *argtype*₁ \Rightarrow *argtype*₂ \Rightarrow \dots \Rightarrow *argtype*_{*n*} \Rightarrow *restype* describes a function which can be applied to *n* arguments.

Function application terms for a function *f* and an argument *a* are denoted by *f a*, no parentheses are required around the argument. Function application terms are left associative, thus a function application to *n* arguments is written *f a*₁ ... *a*_{*n*}. Note that an application *f a*₁ ... *a*_{*m*} where *m* < *n* (a “partial application”) is a correct term and denotes a function taking the remaining *n-m* arguments.

For every constant alternative syntax forms may be defined for application terms. This is often used for binary functions to represent application terms in infix notation with an operator symbol. As an example, the name for the addition function is *plus*, so an application term is denoted in the form *plus 3 5*. For *plus* an alternative syntax is defined as infix notation using the operator symbol *+*, therefore the application term can also be denoted by *3 + 5*. Similar infix notations are supported for many basic functions, using operator symbols such as *-*, ****, *=*, *≠*, *≤*, or *∈*.

Lambda-Terms

Functions can be denoted by lambda terms of the form $\lambda x. \text{term}$ where x is a variable which may occur in the *term*. A function to be applied to n arguments can be denoted by the lambda term $\lambda x_1 \dots x_n. \text{term}$ where $x_1 \dots x_n$ are distinct variables. As usual, types may be specified for (some of) the variables in the form $\lambda(x_1::t_1) \dots (x_n::t_n). \text{term}$.

2.1.4 Definitions and Abbreviations

A constant name may be introduced together with information about its associated value by specifying a term for the value. There are two forms for introducing constant names in this way, definitions and abbreviations.

Definitions

A definition defines a new constant together with its type and value. It is denoted in the form

```
definition name :: type  
where "name  $\equiv$  term"
```

Note that the “defining equation” $\text{name} \equiv \text{term}$ is specified in inner syntax and, like *type* must be delimited by quotes. The *name* may not occur in the *term*, i.e., this form of definition do not support recursion.

If the type of the defined name is a function type, the *term* may be a lambda term. Alternatively, the definition for a function applicable to n arguments can be written in the form

```
definition name :: type  
where "name  $x_1 \dots x_n \equiv \text{term}$ "
```

with variable names $x_1 \dots x_n$ which may occur in the *term*. This form is mainly equivalent to

```
definition name :: type  
where "name  $\equiv \lambda x_1 \dots x_n. \text{term}$ "
```

A short form of a definition is

```
definition "name  $\equiv \text{term}$ "
```

Here, the type of the new constant is derived as the type of the *term*.

Usually, a constant defined in this way is fully specified, i.e., all information about its value is available. However, if the term does not provide this information, the constant is still underspecified. Consider the definition

definition *"mystery2 \equiv mystery"*

where *mystery* is defined as above. Then it is only known that *mystery2* has type $\text{nat} \Rightarrow \text{nat}$ and is the same total function as *mystery*, but nothing is known about its values.

Abbreviations

An abbreviation definition does not define a constant, it only introduces the name as a synonym for a term. Upon input the name is automatically expanded, and upon output it is used whenever a term matches its specification and the term is not too complex. An abbreviation definition is denoted in a similar form as a definition:

abbreviation *name* :: *type*
where *"name \equiv term"*

As for definitions, recursion is not supported, the *name* may not occur in the *term*. The short form is also available as for definitions.

The alternative form for functions is also available. The abbreviation definition

abbreviation *name* :: *type*
where *"name x₁ ... x_n \equiv term"*

introduces a “parameterized” abbreviation. An application term *name term₁ ... term_n* is replaced upon input by *term* where all occurrences of *x_i* have been substituted by *term_i*. Upon output terms are matched with the structure of *term* and if successful a corresponding application term is constructed and displayed.

2.1.5 Overloading

True Overloading

One way of providing information about the value of an underspecified constant is overloading. It provides the information with the help of another constant together with a definition for it.

Overloading depends on the type. Therefore, if a constant is polymorphic, different definitions can be associated for different type instances.

Overloading is only possible for constants which do not yet have a definition, i.e., they must have been defined by **consts** (see Section 2.1.3). Such a constant *name* is associated with *n* definitions by the following overloading specification:

```

overloading
  name1  $\equiv$  name
  ...
  namen  $\equiv$  name
begin
  definition name1 :: type1 where ...
  ...
  definition namen :: typen where ...
end

```

where all *type*_{*i*} must be instances of the type declared for *name*.

The auxiliary constants *name*₁ ... *name*_{*n*} are only introduced locally and cannot be used outside of the **overloading** specification.

Adhoc Overloading

There is also a form of overloading which achieves similar effects although it is implemented completely differently. It is only performed on the syntactic level, like abbreviations. To use it, the theory *HOL-Library.Adhoc_Overloading* must be imported by the surrounding theory:

```
imports "HOL-Library.Adhoc_Overloading"
```

(Here the theory name must be quoted because it contains a minus sign.)

Then a constant name *name* can be defined to be a “type dependent abbreviation” for *n* terms of different type instances by

```
adhoc_overloading name term1 ... termn
```

Upon input the type of *name* is determined from the context, then it is replaced by the corresponding *term*_{*i*}. Upon output terms are matched with the corresponding *term*_{*i*} and if successful *name* is displayed instead.

Although *name* must be the name of an existing constant, only its type is used. The constant is not affected by the adhoc overloading, however, it becomes inaccessible because its name is now used as term abbreviation.

Several constant names can be overloaded in a common specification:

```
adhoc_overloading name1 term11 ... term1n and ... and namek ...
```

2.1.6 Propositions

A proposition denotes a statement, which can be valid or not. Valid statements are called “facts”, they are the main content of a theory. Propositions are specific terms and are hence written in inner syntax and must be enclosed in quotes.

Formulas

In its simplest form a proposition is a single term of type `bool`, such as

$$6 * 7 = 42$$

Terms of type `bool` are also called “formulas”.

A proposition may contain free variables as in

$$2 * x = x + x$$

A formula as proposition is valid if it evaluates to `True` for all possible values substituted for the free variables.

Derivation Rules

More complex propositions can express, “derivation rules” used to derive propositions from other propositions. Derivation rules are denoted using a “metallogic language”. It is still written in inner syntax but uses a small set of “metallogic operators”.

Derivation rules consist of assumptions and a conclusion. They can be written using the metallogic operator \implies in the form

$$A_1 \implies \dots \implies A_n \implies C$$

where the $A_1 \dots A_n$ are the assumptions and C is the conclusion. The conclusion must be a formula. The assumptions may be arbitrary propositions. If an assumption contains metallogic operators parentheses can be used to delimit them from the rest of the derivation rule.

A derivation rule states that if the assumptions are valid, the conclusion can be derived as also being valid. So it can be viewed as a “meta implication” with a similar meaning as a boolean implication, but with a different use.

An example for a rule with a single assumption is

$$(x::nat) < c \implies n*x \leq n*c$$

Note that type `nat` is explicitly specified for variable `x`. This is necessary, because the constants `<`, `*`, and `≤` are overloaded and can be applied to other types than only natural numbers. Therefore the type of `x` cannot be derived automatically. However, when the type of `x` is known, the types of `c` and `n` can be derived to also be `nat`.

An example for a rule with two assumptions is

$$(x::nat) < c \implies n > 0 \implies n*x < n*c$$

In most cases the assumptions are also formulae, as in the example. However, they may also be again derivation rules. Then the rule is a “meta rule” which derives a proposition from other rules. This introductory manual usually does not take such meta rules into account.

Binding Free Variables

A proposition may contain universally bound variables, using the metalogic quantifier \bigwedge in the form

$$\bigwedge x_1 \dots x_n. P$$

where the $x_1 \dots x_n$ may occur free in the proposition P . As usual, types may be specified for (some of) the variables in the form $\bigwedge (x_1::t_1) \dots (x_n::t_n). P$. An example for a valid derivation rule with bound variables is

$$\bigwedge (x::nat) \ c \ n. \ x < c \implies n*x \leq n*c$$

If a standalone proposition contains free variables they are implicitly universally bound. Thus the example derivation rule above is equivalent to the single-assumption example rule in the previous section. Explicit binding of variables is only required to avoid name clashes with constants of the same name. In the proposition

$$\bigwedge (True::nat). \ True < c \implies n*True \leq n*c$$

the name `True` is used locally as a variable of type `nat` instead of the pre-defined constant of type `bool`. Of course, using well known constant names as variables is confusing and should be avoided.

Alternative Rule Syntax

An alternative, Isabelle specific syntax for derivation rules is

$$\bigwedge x_1 \dots x_m. \llbracket A_1; \dots; A_n \rrbracket \implies C$$

which is often considered as more readable, because it better separates the assumptions from the conclusion. In the interactive editor it may be necessary to switch to this form by setting `Print Mode` to `brackets` in `Plugin Options for Isabelle General`. The fat brackets are available for input in the editor's Symbols panel in tab "Punctuation".

Using this syntax the two-assumption example rule from the previous section is denoted by

$$\bigwedge (x::nat) \ c \ n. \llbracket x < c; \ n > 0 \rrbracket \implies n*x < n*c$$

or equivalently without quantifier by

$$\llbracket (x::nat) < c; \ n > 0 \rrbracket \implies n*x < n*c$$

Note that in the literature a derivation rule $\llbracket P; Q \rrbracket \implies P \wedge Q$ is often denoted in the form

$$\frac{P \quad Q}{P \wedge Q}$$

Another alternative, Isabelle specific syntax for a derivation rule $\bigwedge x_1 \dots x_m. \llbracket A_1; \dots; A_n \rrbracket \implies C$ is the “structured” proposition

"C" if "A₁" ... "A_n" for x₁ ... x_m

The assumptions and the variables may be grouped or separated for better readability by the keyword **and**. For every group of variables a type may be specified in the usual form, it applies to all variables in the group. Note that the keywords **if**, **and**, **for** belong to the outer syntax. Thus, the original rule must be quoted as a whole, whereas in the structured proposition only the sub-propositions C, A_1, \dots, A_n must be individually quoted. The x_1, \dots, x_m need not be quoted, but if a type is specified for a variable the type must be quoted, if it is not a single type name.

If written in this form, the two-assumption example rule from the previous subsections becomes

"n*x < n*c" if "x < c" and "n > 0" for x::nat and n c

2.1.7 Theorems

A theorem specifies a proposition together with a proof, that the proposition is valid. Thus it adds a fact to the enclosing theory. A simple form of a theorem is

theorem "prop" <proof>

where *prop* is a proposition in inner syntax and <proof> is a proof as described in Section 2.2. The keyword **theorem** can be replaced by one of the keywords **lemma**, **corollary**, **proposition** to give a hint about the use of the statement to the reader.

Unknowns

Whenever a theorem turns a proposition to a fact, the free (or universally bound) variables are replaced by “unknowns”. For a variable *name* the corresponding unknown is *?name*. This is only a technical difference, it signals to Isabelle that the unknowns can be consistently substituted by arbitrary terms, as long as the types are preserved.

When turned to a fact, the example rule from the previous sections becomes

?x < ?c \implies ?n*?x \leq ?n*?c

with type `nat` associated to all unknowns.

The result of such a substitution is always a special case of the fact and therefore also a fact. In this way a fact with unknowns gives rise to a (usually infinite) number of facts which are constructed by substituting unknowns by terms.

Isabelle can be configured to suppress the question mark when displaying unknowns, then this technical difference becomes invisible.

Named Facts

Facts are often used in proofs of other facts. For this purpose they can be named so that they can be referenced by name. A named fact is specified by a theorem of the form

```
theorem name: "prop" <proof>
```

The example rule from the previous sections can be turned into a fact named `example1` by

```
theorem example1: "(x::nat) < c  $\implies$  n*x  $\leq$  n*c" <proof>
```

It is also possible to introduce named collections of facts. A simple way to introduce such a named collection is

```
lemmas name = name1 ... namen
```

where `name1 ... namen` are names of existing facts or fact collections.

If there is a second rule stated as a named fact by

```
theorem example2: "(x::nat)  $\leq$  c  $\implies$  x + m  $\leq$  c + m" <proof>
```

a named collection can be introduced by

```
lemmas examples = example1 example2
```

Alternatively a “dynamic fact set” can be declared by

```
named_theorems name
```

It can be used as a “bucket” where facts can be added afterwards by specifying the bucket name in the theorem:

```
theorem [name]: "prop" <proof>
```

or together with specifying a fixed fact name `namef` by

```
theorem namef [name]: "prop" <proof>
```

A named fact or fact set (but not a dynamic fact set) can be displayed using the command

thm *name*

which may be entered outside of theorems and at most positions in a proof. The facts are displayed in the Output panel (see Section 1.2.3).

Alternative Theorem Syntax

There is an alternative syntax for theorems which have a derivation rule as their proposition. A theorem **theorem** " $\bigwedge x_1 \dots x_m. \llbracket A_1; \dots; A_n \rrbracket \implies C$ " *<proof>* can also be specified in the form

```
theorem
  fixes  $x_1 \dots x_m$ 
  assumes " $A_1$ " ... " $A_n$ "
  shows " $C$ "
  <proof>
```

Similar to the structured proposition form, the variables and assumptions may be grouped by **and**, the keywords belong to the outer syntax and the C , A_1 , ..., A_n must be individually quoted.

Using this syntax the two-assumption example rule from the previous sections can be written as

```
theorem
  fixes  $x::nat$  and  $c\ n$ 
  assumes " $x < c$ " and " $n > 0$ "
  shows " $n*x < n*c$ "
  <proof>
```

In contrast to the general theorem syntax this alternative syntax allows to specify names for some or all of the assumption groups as

```
assumes  $name_1$ : " $A_{11}$ " ... " $A_{1m_1}$ " and ... and  $name_n$ : " $A_{n1}$ " ... " $A_{nm_n}$ "
```

These names can (only) be used in the proof of the theorem. More consequences this syntax has for the proof are described in Section 2.2.

Note that a name specified for the conclusion as

```
shows  $name$ : " $C$ "
```

becomes the name for the whole fact introduced by the theorem, not only for the conclusion. It is not available in the proof of the theorem. Alternatively the name for the fact can be specified after the **theorem** keyword:

```

theorem name:
  fixes x::nat and c and n
  assumes "x < c" and "n > 0"
  shows "n*x < n*c"
  <proof>

```

Definitions as Facts

The definitions described in Section 2.1.4 also introduce facts in the enclosing theory. Every definition introduces a new constant and specifies a defining equation of the form *name* \equiv *term* for it. This equation is a proposition using the “meta equality” \equiv which is another metalogic operator. It is the initial information given for the new constant, thus it is valid “by definition” and is a fact in the theory.

These facts are automatically named. If *name* is the name of the defined constant, the defining equation is named *name_def*. Alternatively an explicit name can be specified in the form

```

definition name :: type
where fact_name: "name  $\equiv$  term"

```

Although the auxiliary constants used in an **overloading** specification (see Section 2.1.5) are not accessible outside the specification, their definitions are. So they can be referred by their names and used as information about the overloaded constant.

2.1.8 Locales

There are cases where theory content such as definitions and theorems occur which has similar structure but differs in some types or terms. Then it is useful to define a “template” and instantiate it several times. This can be done in Isabelle using a “locale”.

Simple Locales

A locale can be seen as a parameterized theory fragment, where the parameters are terms. A locale with *n* parameters is defined by

```

locale name =
  fixes x1 ... xn
begin
  ...
end

```

where the variables x_1, \dots, x_n are the parameters. Like the bound variables in a theorem they can be grouped by **and** and types can be specified for some or all groups. The content between **begin** and **end** may consist of definitions and theorems which may use the parameter names like constant names. Content may also be added to an existing locale in the form

```
context name
begin
  ...
end
```

Therefore the **begin** ... **end** block can also be omitted in the locale definition and the locale can be filled later.

An instance of the parameterized theory fragment is created by “interpreting” the locale in the form

```
interpretation name term1 ... termn .
```

where *term*₁ ... *term*_{*n*} are the terms to be substituted for the locale parameters, their types must match the parameter types, i.e., must be instances of them. The final dot in the interpretation is a rudimentary proof. An actual proof is needed, if the locale definition specifies additional assumptions for the parameters.

Locales with Assumptions

Additional assumptions for locale parameters can be specified as propositions in the form

```
locale name =
  fixes x1 ... xn
  assumes A1 ... Am
begin
  ...
end
```

where the *A*₁, ..., *A*_{*m*} are propositions. Like in a theorem, they can be grouped by **and** and named. The names can be used to reference the assumptions as facts in proofs in the locale content. When the locale is interpreted, all the assumptions must be proved with the actual terms substituted for the parameters. Therefore the more general form of an interpretation is

```
interpretation name term1 ... termn <proof>
```

Extending Locales

A locale may extend one or more other locales using the form

```
locale name = name1 + ... + namen +  
  fixes ...  
  assumes ...  
begin  
  ...  
end
```

where *name*₁ ... *name*_{*n*} are the names of the extended locales. Their parameters become parameters of the defined locale, inserted before the parameters declared by the **fixes** ... clause.

2.2 Isabelle Proofs

Every proposition stated as a fact in an Isabelle theory must be proven immediately by specifying a proof for it. A proof may have a complex structure of several steps and nested sub-proofs, its structure is part of the outer syntax.

2.2.1 Proof Context

Every proof is performed in a temporary environment which collects facts and other proof elements. This environment is called the “proof context”. At the end of the proof the proof context is disposed with all its content, only the proven fact remains in the enclosing entity.

The proof context may contain

- Facts: as usual, facts are valid propositions. However, they need not be globally valid, they can be assumed to be only valid locally during the proof.
- Goals: a goal is a proposition which has not yet been proven. Typically it is the duty of a proof to prove one or more goals in its proof context.
- Fixed variables: fixed variables are used to denote the “arbitrary but fixed” objects often used in a proof. They can be used in all facts and goals in the same proof context.
- Term abbreviations: these are names introduced locally for terms. Using such names for terms occurring in propositions it is often possible to denote propositions in a more concise form.

The initial proof context in a theorem of the form **theorem** "*prop*" $\langle proof \rangle$ has the proposition *prop* as the only goal and is otherwise empty.

The initial proof context in a theorem of the alternative form

theorem

fixes $x_1 \dots x_m$
assumes " A_1 " \dots " A_n "
shows " C "
 $\langle proof \rangle$

has the proposition C as the only goal and contains $x_1 \dots x_m$ as fixed variables and the A_1, \dots, A_n as assumed facts. Therefore, although the theorem is equivalent to the theorem **theorem** " $\bigwedge x_1 \dots x_m. \llbracket A_1; \dots; A_n \rrbracket \implies C$ " $\langle proof \rangle$ the initial proof contexts differ.

2.2.2 Proof Procedure

Assume you want to prove a derivation rule $A \implies C$ with a single assumption A and the conclusion C . The basic procedure to build a proof for it is to construct a sequence of the form $F_1 \implies F_2, F_2 \implies F_3, F_3 \implies \dots \implies F_{n-1}, F_{n-1} \implies F_n$ from rules $RA_i \implies RC_i$ for $i=1 \dots n-1$ which are already known to be valid (i.e., facts) where F_1 matches with A and RA_1, F_n matches with C and RC_{n-1} , and every other F_i matches with RA_i and RC_{i-1} .

The sequence can be constructed from left to right (called “forward reasoning”) or from right to left (called “backward reasoning”) or by a combination of both.

Consider the rule $(x::nat) < 5 \implies 2*x+3 \leq 2*5+3$. A proof can be constructed from the two example rules *example1* and *example2* from the previous sections as the sequence $(x::nat) < 5 \implies 2*x \leq 2*5, 2*x \leq 2*5 \implies 2*x+3 \leq 2*5+3$ consisting of three facts.

Forward reasoning starts by assuming A to be a local fact and incrementally constructs the sequence from it. An intermediate result is a part $F_1 \implies \dots \implies F_i$ of the sequence, here F_i is the “current fact”. A forward reasoning step consists of stating a proposition F_{i+1} and proving it to be a new local fact from the current fact F_i using a valid rule $RA_i \implies RC_i$. The step results in the extended sequence $F_1 \implies \dots \implies F_i, F_i \implies F_{i+1}$ and the new current fact F_{i+1} . When a step successfully proves a current fact F_n which matches the conclusion C the proof is complete.

Backward reasoning starts at the conclusion C of the single goal $A \implies C$ and incrementally constructs the sequence from it backwards. An intermediate result is a part $F_i \implies \dots \implies F_n$ of the sequence where F_n matches C , here the remaining part of the sequence $F_1 \implies F_i$ is the “current goal”. A backward reasoning step consists of applying a proof method to F_i which constructs a new current goal $F_1 \implies F_{i-1}$ and the extended sequence F_{i-1}

$\implies F_i$, $F_i \implies \dots \implies F_n$. When a step produces the new current goal $F_1 \implies F_1$, which is trivially valid, the proof is complete.

Note the slight difference in how the steps are specified: A forward step specifies the new current fact F_{i+1} and then proves it. A backward step specifies the proof method, the new current goal $F_1 \implies F_{i-1}$ is constructed by the method and is not an explicit part of the proof text. For that reason a proof constructed by forward reasoning is usually easier to read and write than a proof constructed by backward reasoning, since in the former case the sequence of the facts F_i is explicitly specified in the proof text, whereas in the latter case the sequence of the facts F_i is implicitly constructed and the proof text specifies only the methods.

However, since every forward reasoning step again requires a proof as its part (a “subproof”), no proof can be written using only forward reasoning steps. The main idea of writing “good” proofs is to use nested forward reasoning until every subproof is simple enough to be done in a single backward reasoning step, i.e., the proof method directly goes from the conclusion to the assumption.

Unification

The matching at the beginning and end of the sequence and when joining the used rules is done by “unification”. Two propositions P and Q are unified by substituting terms for unknowns in P and Q so that the results become syntactically equal.

Since only the $RA_i \implies RC_i$ are facts containing unknowns, only they are modified by the unification, A and C remain unchanged.

Note that when an unknown is substituted by a term in RA_i , the same unknown must be substituted by the same term in RC_i and vice versa, to preserve the validness of the rule $RA_i \implies RC_i$. In other words, the sequence is usually constructed from specializations of the facts $RA_i \implies RC_i$ where every conclusion is syntactically equal to the assumption of the next rule.

In the example the assumption $?x < ?c$ of rule `example1` is unified with $(x::nat) < 5$ by substituting the term 5 for the unknown $?c$, and the variable x for the unknown $?x$ resulting in the specialized rule $(x::nat) < 5 \implies n*x \leq n*5$. The conclusion $?x + ?m \leq ?c + ?m$ of rule `example2` is unified with $2*x+3 \leq 2*5+3$ by substituting the term $2*x$ for the unknown $?x$, the term $2*5$ for the unknown $?c$, and the term 3 for the unknown $?m$ resulting in the specialized rule $2*x \leq 2*5 \implies 2*x+3 \leq 2*5+3$. Now the two specialized rules can be joined by substituting the term 2 for the unknown $?n$ in the first, resulting in the sequence which constitutes the proof.

Multiple Assumptions

If the rule to be proven has more than one assumption A the sequence to be constructed becomes a tree where the branches start at (copies of) the assumptions A_1, \dots, A_n and merge to finally lead to the conclusion C . Two branches which end in facts F_{1n} and F_{2m} are joined by a step $\llbracket F_{1n}; F_{2m} \rrbracket \Rightarrow F_1$ to a common branch which continues from fact F_1 .

Now a forward reasoning step may use several current facts to prove a new current fact. Therefore all proven local facts are stored in the proof context for possible later use. Every forward reasoning step selects a subset of the stored local facts as the current facts and uses them to prove a new local fact from them.

A backward reasoning step may now produce several new current goals, which belong to different branches in the tree. A step always produces the goals for all branches, therefore the previous goal is never used again in a step and is removed from the proof context after the step. When a current goal has the form $A \Rightarrow A$ the proof method “*assumption*” removes it from the proof context without producing a new goal. Thus a proof ends when no goal remains in the proof context.

The set of current goals is called the “goal state” of the proof. Since it is not visible in the proof text, the interactive editor displays the current goal state in the separate State panel and optionally also in the Output panel (see Sections 1.2.3 and 1.2.4), according to the cursor position in the proof text.

Proving from External Facts

The branches in the fact tree need not always start at an assumption A_i , they may also start at an “external” fact which is not part of the local proof context. In such cases the used external facts are referenced by their names. In that way a proof can use facts from the enclosing theory and a subproof can use facts from the enclosing proof(s) and the enclosing toplevel theory. In particular, if the proposition of a theorem has no assumptions, i.e., the proposition is a formula and consists only of the conclusion C , every proof must start at one or more external facts (if C is no tautology which is valid by itself).

2.2.3 Basic Proof Structure

A proof is written in outer syntax and describes how the fact tree is constructed which leads from the assumptions or external facts to the conclusion.

Proof Modes

When writing a proof the “proof mode” determines the mode of operation: whether forward reasoning (mode: *proof(state)*) or backward reasoning (mode: *proof(prove)*) is used.

The mode names refer to what is done in the next steps: In mode *proof(state)* facts are “stated” which lead to the conclusion, whereas in mode *proof(prove)* the goals are “proven”, leading to assumptions and external facts.

At the beginning of a proof the mode is always *proof(prove)*, i.e., backward reasoning. In the course of the proof it is possible to switch to forward reasoning mode *proof(state)*, but not back again. After switching to forward reasoning the proof must be completed in forward reasoning mode, only at the end a last backward reasoning step may be applied.

However, in forward reasoning for every stated fact a (sub-)proof must be specified, which again starts in backward reasoning mode. This way it is possible to freely switch between both modes in the course of a proof with nested subproofs.

Proof Syntax

If BS_i denote backward reasoning steps and FS_i denote forward reasoning steps, the general form of a proof is

```
BS1 ... BSn
proof BSn+1
    FS1 ... FSm
qed BSn+2
```

The last step BS_{n+2} can be omitted if it is not needed.

The part **proof** BS_{n+1} switches from backward reasoning mode *proof(prove)* to forward reasoning mode *proof(state)*.

The part **proof** ... **qed** can be replaced by **done**, then the proof only consists of backward reasoning steps and has the form $BS_1 \dots BS_n$ **done**. Such proofs are called “proof scripts”.

If the backward reasoning steps $BS_1 \dots BS_n$ are omitted the proof only consists of the forward reasoning part and has the form

```
proof BS1
    FS1 ... FSm
qed BS2
```

where BS_2 can also be omitted. Such proofs are called “structured proofs”. A structured proof can be so simple, that it has no forward reasoning steps. For this case the syntax

by BS_1 BS_2

abbreviates the form **proof** BS_1 **qed** BS_2 . Again, BS_2 can be omitted which leads to the form

by BS_1

In this form the proof consists of a single backward reasoning step which directly leads from the conclusion C to the assumptions and used external facts.

Fake Proofs

A proof can also be specified as

sorry

This is a “fake proof” which turns the proposition to a fact without actually proving it.

A fake proof can be specified at any point in backward reasoning mode, so it can be used to abort a proof script in the form $BS_1 \dots BS_n$ **sorry**.

A structured proof in forward reasoning mode cannot be aborted in this way, however, subproofs can be specified as fake proofs. This makes it possible to interactively develop a structured proof in a top-down way, by first stating all required facts with fake proofs and then replacing the fake proofs by actual proofs.

Nested Proof Contexts

The proof contexts in a structured proof can be nested. In a nested context the content of the enclosing contexts is available together with the local content. When a nested context is ended, it is removed together with all its local content.

A nested proof context is created syntactically by enclosing forward reasoning steps in braces:

$FS_1 \dots FS_m \{ FS_{m+1} \dots FS_n \} FS_{n+1} \dots$

Note that according to the description until now the nested context is useless, because the facts introduced by its forward reasoning steps are removed at its end and cannot contribute to the proof. How the content of a nested context can be “exported” and preserved for later use will be explained further below.

For names, nested contexts behave like a usual block structure: A name can be redefined in a nested context, then the named object in the outer context becomes inaccessible (“shadowed”) in the inner context, but becomes accessible again when the inner context ends.

When two nested contexts follow each other immediately, this has the effect of “clearing” the content of the inner contexts: the content of the first context is removed and the second context starts being empty. This can be denoted by the keyword

next

which can be thought of being equivalent to a pair $\} \{$ of adjacent braces.

Moreover the syntax **proof** *method*₁ *FS*₁ ... *FS*_{*n*} **qed** *method*₂ automatically wraps the forward steps *FS*₁ ... *FS*_{*n*} in a nested context. Therefore it is possible to denote a structured proof which only consists of a sequence of nested contexts without braces as

```
proof method1
  FS11 ... FS1m1 next FS21 ... FS2m2 next ... next FSn1 ... FSnmn
qed methodn+2
```

where each occurrence of **next** clears the content of the context. The goal state of the proof is maintained in the enclosing outermost context, thus it is preserved over the whole sequence of nested contexts.

2.2.4 Backward Reasoning Steps

A backward reasoning step consist of applying a proof method.

Proof Methods

Proof methods are basically denoted by method names, such as *standard*, *simp*, or *rule*. A proof method name can also be a symbol, such as $-$.

A method may have arguments, then it usually must be delimited by parentheses such as in (*rule example1*) or (*simp add: example2*), where *example1* and *example2* are fact names.

Methods can be applied to the goal state, they modify the goal state by removing and adding goals.

The effect of applying a method is determined by its implementation and must be known to the proof writer. Isabelle supports a large number of proof methods. A selection of proof methods used in this manual are described in Section 2.3.

Method Application

A standalone method application step is denoted as

apply *method*

where *method* denotes the proof method to be applied.

The backward reasoning steps which follow **proof** and **qed** in a structured proof are simply denoted by the applied method. Hence the general form of a proof where all backward reasoning steps are method applications is

```
apply method1 ... apply methodn
proof methodn+1
  FS1 ... FSm
qed methodn+2
```

where *FS*₁ ... *FS*_{*m*} are forward reasoning steps. The method *method*_{*n*+1} is called the “initial method” of the structured proof part.

2.2.5 Forward Reasoning Steps

A forward reasoning step consist of stating and proving a fact.

Stating a Fact

A fact is stated in the form

have “*prop*” ⟨*proof*⟩

where *prop* is a proposition in inner syntax and ⟨*proof*⟩ is a (sub-) proof for it. This form is similar to the specification of a theorem in a theory and has a similar effect in the local proof context.

As for a theorem the fact can be named:

have *name*: “*prop*” ⟨*proof*⟩

Note that the alternative form of a theorem using **fixes**, **assumes**, and **shows** (see Section 2.1.7) is not available for stating facts in a proof.

The subproof ⟨*proof*⟩ uses a nested context, therefore all content of the enclosing proof context is available there and can be referenced by name, as long as the name is not shadowed by a redefinition in the subproof. Note that the name given to the fact to be proven cannot be used to access it in the subproof, because it is only assigned after the proof has been finished.

Proving a Goal

A forward reasoning proof ends, if the last stated fact F_n unifies with the conclusion C . Therefore a special form of stating a fact exists, which, after proving the fact, replaces free variables by unknowns (which is called “exporting the fact”) and tries to unify it with the conclusion of a goal in the goal state. If successful, it removes the goal from the goal state:

show "*prop*" $\langle proof \rangle$

The syntax is the same as for **have**. If the unification with some goal conclusion is not successful the step is erroneous and the proof cannot be continued, in the interactive editor an error message is displayed.

The **show** step tries to match and remove a goal from the innermost enclosing proof context which maintains a goal state. This is one way how a fact proven in a nested context can affect an enclosing context and thus contribute to the proof there. Since the forward reasoning steps in a structured proof are wrapped in a nested context, while the goal state is maintained in the enclosing outer context, the **show** step affects the goal state of the enclosing proof.

The **have** and **show** steps are called “goal statements”, because they state the proposition *prop* as a goal which is then proven by the $\langle proof \rangle$.

Note that the proposition *prop* in a **show** statement often is the same proposition which has been specified as conclusion C in the proposition $\llbracket A_1; \dots; A_n \rrbracket \implies C$ which should be proven by the proof. To avoid repeating it, Isabelle automatically provides the abbreviation *?thesis* for it. So in the simplest case the last step of a forward proof can be written as

show *?thesis* $\langle proof \rangle$

The abbreviation *?thesis* is a single identifier, therefore it needs not be quoted.

If, however, the application of the initial method *method* in a structured proof **proof** *method* ... modifies the original goal, this modification is not reflected in *?thesis*. So a statement **show** *?thesis* $\langle proof \rangle$ will usually not work, because *?thesis* no more unifies with the conclusion of the modified goal. Instead, the proof writer must know the modified goal and specify its conclusion explicitly as proposition in the **show** statement. If the *method* splits the goal into several new goals, several **show** statements may be needed to remove them.

To test whether a proposition unifies with the conclusion of a goal in the goal state, a **show** statement can be specified with a fake proof:

show "*prop*" **sorry**

If that statement is accepted, the proposition unifies with the conclusion of a goal and removes it.

2.2.6 Facts as Proof Input

If a linear fact sequence $F_1 \implies \dots \implies F_n$ is constructed in forward reasoning mode in the form

```
have "F1" <proof>1
...
have "Fn" <proof>n
```

every fact F_i needs to refer to the previous fact F_{i-1} in its proof $\langle proof \rangle_i$. This can be done by naming all facts

```
have name1: "F1" <proof>1
...
have namen: "Fn" <proof>n
```

and refer to F_{i-1} in $proof_i$ by $name_{i-1}$.

Isabelle supports an alternative way by passing facts as input to a proof.

Using Input Facts in a Proof

The input facts are passed as input to the first method applied in the proof. In a proof script it is the method applied in the first **apply** step, in a structured proof **proof method** ... it is the initial method *method*.

Every proof method accepts a set of facts as input. Whether it processes them and how it uses them depends on the kind of method. Therefore input facts for a proof only work in the desired way, if a corresponding method is used at the beginning of the proof. See Section 2.3 for descriptions how methods process input facts.

Inputting Facts into a Proof

In backward reasoning mode *proof(prove)* facts can be input to the remaining proof $\langle proof \rangle$ by

```
using name1 ... namen <proof>
```

where the $name_i$ are names of facts or fact sets. The union of all referred facts is input to the proof following the **using** specification. In a proof script it is input to the next **apply** step. If a structured proof follows, it is input to its initial method. Since in backward reasoning mode no local facts are stated by previous steps, only external facts can be input this way.

In forward reasoning mode *proof(state)* fact input is supported with the help of the special fact set name *this*. The statement

then

inputs the facts named *this* to the proof of the following goal statement.

The statement **then** must be immediately followed by a goal statement (**have** or **show**). This is enforced by a special third proof mode *proof(chain)*. In it only a goal statement is allowed, **then** switches to this mode, the goal statement switches back to mode *proof(state)* after its proof.

Note that **then** is allowed in forward reasoning mode, although it does not state a fact. There are several other such auxiliary statements allowed in mode *proof(state)* in addition to the goal statements **have** and **show**.

The fact set *this* can be set by the statement

```
note name1 ... namen
```

Therefore the statement sequence

```
note name1 ... namen  
then have "prop" <proof>
```

inputs the union of all facts referred by *name*₁ ... *name*_{*n*} to the <*proof*>, in the same way as **using** inputs them to the remaining proof following it.

The statement sequence

```
note name1 ... namen then
```

can be abbreviated by the statement

```
from name1 ... namen
```

Like **then** it switches to mode *proof(chain)* and it inputs the union of the facts referred by *name*₁ ... *name*_{*n*} to the proof of the following goal statement.

2.2.7 Fact Chaining

In both cases described for fact input until now, the facts still have been referred by names. This can be avoided by automatically using a stated fact as input to the proof of the next stated fact. That is called “fact chaining”.

Automatic Update of the Current Facts

Fact chaining is achieved, because Isabelle automatically updates the fact set *this*. Whenever a new fact is added to the proof context, the set *this* is redefined to contain (only) this fact. In particular, after every goal statement *this* names the new proven fact. Therefore the fact set *this* is also called the “current facts”.

Thus a linear sequence of facts can be constructed by

```

have "F1" <proof>1
then have "F2" <proof>2
...
then have "Fn" <proof>n

```

Now in every *proof*_{*i*} the fact *F*_{*i*-1} is available as input and can be used to prove *F*_{*i*}.

Chaining can be combined with explicit fact referral by a statement of the form

```

note this name1 ... namen

```

It sets *this* to the union of *this* and the name₁ ... name_n, i.e., it adds the name₁ ... name_n to *this*. In this way the current facts can be extended with other facts and then chained to the proof of the next stated fact.

The statement sequence

```

note this name1 ... namen then

```

can be abbreviated by the statement

```

with name1 ... namen

```

Like **then** it switches to mode *proof(chain)* and it inputs the union of the facts referred by name₁ ... name_n together with the current facts to the proof of the following goal statement.

If a proof consists of a fact tree with several branches, every branch can be constructed this way. Before switching to the next branch the last fact must be named, so that it can later be used to prove the fact where the branches join. A corresponding proof pattern for two branches which join at fact *F* is

```

have "F11" <proof>11
then have "F12" <proof>12
...
then have name1: "F1m" <proof>1m
have "F21" <proof>21
then have "F22" <proof>22
...
then have "F2n" <proof>2n
with name1 have "F" <proof>

```

Naming and Grouping Current Facts

Since the fact set built by a **note** statement is overwritten by the next stated fact, it is possible to give it an explicit name in addition to the name *this* in the form

note *name* = *name*₁ ... *name*_{*n*}

The *name* can be used later to refer to the same fact set again, when *this* has already been updated. Defining such names is only possible in the **note** statement, not in the abbreviated forms **from** and **with**.

The facts specified in **note**, **from**, **with**, and **using** can be grouped by separating them by **and**. Thus it is possible to write

from *name*₁ **and** ... **and** *name*_{*n*} **have** "*prop*" *<proof>*

In the case of a **note** statement every group can be given an additional explicit name as in

note *name*₁ = *name*₁₁ ... *name*_{1*m*1} **and** ... **and** *name*_{*n*} = *name*_{*n*1} ... *name*_{*n**m**n*}

Accessing Input Facts in a Structured Proof

At the beginning of a structured proof the set name *this* is undefined, the name cannot be used to refer to the input facts (which are the current facts in the enclosing proof). To access the input facts they must be named before they are chained to the goal statement, then they can be referenced in the subproof by that name. For example in

note *input* = *this*
then have "*prop*" *<proof>*

the input facts can be referenced by the name *input* in *<proof>*.

Exporting the Current Facts of a Nested Context

At the end of a nested context (see Section 2.2.3) the current facts are automatically exported to the enclosing context, i.e. they become available there as the fact set named *this*, replacing the current facts before the nested context. This is another way how facts from a nested context can contribute to the overall proof.

Basically, only the last fact is current at the end of a context. Arbitrary facts can be exported from the nested context by explicitly making them current at its end, typically using a **note** statement:

```
... {
  have f1: "prop1" <proof>1
  ...
  have fn: "propn" <proof>n
  note f1 ... fn
}
```

Here all facts are named and the **note** statement makes them current by referring them by their names. Note, that the names are only valid in the nested context and cannot be used to refer to the exported facts in the outer context.

The exported facts can be used in the outer context like all other current facts by directly chaining them to the next stated fact:

```
... { ... } then have "prop" <proof> ...
```

or by naming them for later use, with the help of a **note** statement:

```
... { ... } note name = this ...
```

2.2.8 Assuming Facts

In a theorem of the form **theorem** " $\bigwedge x_1 \dots x_m. \llbracket A_1; \dots; A_n \rrbracket \implies C$ " *<proof>* the assumptions A_1, \dots, A_n are needed as facts in the proof context to start the branches of the fact tree. However, they are not automatically inserted, that must be done explicitly.

Introducing Assumed Facts

An assumption is inserted in the proof context by a statement of the form

```
assume "prop"
```

Several assumptions can be inserted in a single **assume** statement of the form

```
assume "prop1" ... "propn"
```

As usual, the assumptions can be grouped by **and** and the groups can be named.

Like goal statements an **assume** statement makes the assumed facts current, i.e. it updates the set *this* to contain the specified propositions as facts, so that they can be chained to a following goal statement. This way the fact sequence $A \implies F_1, F_1 \implies \dots$ of a proof using fact chaining can be started:

```
assume "A"
then have "F1"
...
```

Alternatively, the assumed facts can be named:

```
assume name: "prop1" ... "propn"
```

so that they can be referred by name in the rest of the proof. Then the fact chain can be started in the form

```
assume a: "A"
from a have "F1"
...
```

This can be useful if the proof has several branches which all start at the same assumption.

Admissible Assumed Facts

In an **assume** statement no proof is needed, since these propositions are only “assumed” to be valid. Therefore, only the propositions occurring as assumptions in the goal $\llbracket A_1; \dots; A_n \rrbracket \implies C$ to be proven are allowed here.

Actually, this condition is not checked for the **assume** statement, an arbitrary proposition can be specified by it. The condition becomes only relevant in subsequent **show** statements. When the fact proven by a show statement is tried to be unified with the conclusion of a goal, additionally each proposition stated in an **assume** statement is tested for unifying with an assumption of the same goal. If that is not satisfied, the **show** statement fails. Therefore, if a proposition is used in an **assume** statement which does not unify with an assumption in a goal, the proof cannot be completed, because all **show** statements will fail.

Exporting Facts with Assumptions

More generally, whenever a local fact F is exported from a proof context, it is combined with all locally assumed facts AF_1, \dots, AF_n to the derivation rule $\llbracket AF_1; \dots; AF_n \rrbracket \implies F$. This reflects the intention of the local assumptions: They may have been used locally to prove F without knowing whether they are valid. So outside the local context F is only known to be valid if all the assumptions are valid.

If the fact F is itself a derivation rule $\llbracket A_1; \dots; A_n \rrbracket \implies C$ then the locally assumed facts are added, resulting in the exported rule $\llbracket AF_1; \dots; AF_n; A_1; \dots; A_n \rrbracket \implies C$.

If the fact F has been proven in a **show** statement it is also exported in this way, resulting in a derivation rule. It matches with a goal if the conclusion of the exported rule unifies with the goal conclusion and if every assumption of the exported rule unifies with an assumption in the goal. This way of matching a rule with a goal is called “refinement”. So the condition for a successful **show** statement can be stated as “the exported fact must refine a goal”.

Externally Assumed Facts

To make **show** statements succeed, an **assume** statement will usually repeat one or more assumptions from the proposition to be proven. This is similar to a **show** statement, which usually repeats the conclusion of that proposition. However, unlike *?thesis* for the conclusion, there is no abbreviation provided for the assumptions.

To avoid repeating propositions in **assume** statements, the proposition to be proven can be specified in the form (see Section 2.1.7)

theorem

```
assumes "A1" and ... and "An"  
shows "C"  
⟨proof⟩
```

As described in Section 2.2.1 this form automatically inserts the assumptions as assumed facts in the proof context. No **assume** statements are needed, thus the assumptions need not be repeated.

The assumptions can still be named and referred by name in the proof. A proof can be started at assumption A_1 in the form

theorem

```
assumes name1: "A1" and ... and namen: "An"  
shows "C"  
proof method  
  from name1 have "F1" ⟨proof⟩  
  ...
```

Additionally, the set of all assumptions specified in this form of a theorem is automatically named *assms*. Since unneeded assumptions usually do not harm in a proof, each proof branch can be started in the form

```
from assms have "F1"  
...
```

but it is usually clearer for the reader to specify only the relevant assumption(s) by explicit names.

In subproofs **assume** statements cannot be avoided in this way, because propositions in local goal statements cannot be specified using **assumes** and **shows**. However, goal statements usually specify only a fact as proposition without assumptions. Instead of assumed facts the subproof can either use facts provided as input, or use external facts from the enclosing proof context by referring to them by name.

Note that there is another difference when specifying a theorem using **assumes** and **shows**. As described in Section 2.2.1, the goal only consists of the conclusion C , the assumptions do not belong to the goal. So how can a fact

exported by a **show** statement refine the goal, if the assumed facts inserted by **assumes** are added to the fact upon export? For the answer remember from Section 2.2.3 that a structured proof always has two proof contexts, an outer and an inner one. Whereas facts inserted by **assume** belong to the inner context, the facts inserted by **assumes** are part of the *outer* context and are not added to facts exported by **show** statements in the inner context and therefore need not be present in refined goals.

Presuming Facts

It is also possible to use a proposition as assumed fact which does not unify with an assumption in a goal, but can be proven from them. In other words, the proof is started somewhere in the middle of the fact tree, works in forward reasoning mode, and when it reaches the conclusion the assumed fact remains to be proven. The statement

presume "*prop*"

inserts such a presumed fact into the proof context.

When a fact is exported from a context with presumed facts, they do not become a part of the exported rule. Instead, at the end of the context for each presumed fact F_p a new goal $\llbracket A_1; \dots; A_n \rrbracket \implies F_p$ is added to the enclosing goal state. So the proof has to continue after proving all original goals and is only finished when all such goals for presumed facts have been proven as well.

2.2.9 Fixing Variables

Variables occurring in the proposition of a theorem can be used in the proof as well, they are universally bound in the whole proof context, if they are not explicitly bound in the proposition, which restricts their use to the proposition itself. Thus in

theorem " $\bigwedge x::nat. x < 3 \implies x < 5$ " $\langle proof \rangle$

the variable x is restricted to the proposition and is not accessible in $\langle proof \rangle$, whereas in

theorem " $(x::nat) < 3 \implies x < 5$ " $\langle proof \rangle$

and

theorem " $x < 3 \implies x < 5$ " **for** $x::nat$ $\langle proof \rangle$

the variable x is accessible in $\langle proof \rangle$.

Local Variables

Additional local variables can be introduced (“fixed”) in a proof context in mode *proof(state)* by the statement

```
fix  $x_1 \dots x_n$ 
```

As usual the variables can be grouped by **and** and types can be specified for (some of) the groups.

If a variable name is used in a proof context without explicitly fixing it, it either refers to a variable in an enclosing context or in the proposition to be proven, or it is free. If it is explicitly fixed it names a variable which is different from all variables with the same name in enclosing contexts and the proposition to be proven.

A fixed local variable is common to the whole local context. If it occurs in several local facts it always is the same variable, it is not automatically restricted to the fact, as for toplevel theorems. Hence in

```
fix  $x :: \text{nat}$   
assume  $a$ : " $x < 3$ "  
have " $x < 5$ "  $\langle \text{proof} \rangle$ 
```

the $\langle \text{proof} \rangle$ may refer to fact a because the x is the same variable in both facts.

By convention variable names are often short consisting of one or two letters, whereas constants defined on toplevel in a theory have longer and more descriptive names. Therefore it is usually not necessary to explicitly fix the variables in the proposition of a theorem to prevent name clashes with constants. By contrast, in a nested proof context there may be other variables with the same name in enclosing contexts, therefore it is recommended to explicitly fix all local variables.

Exporting Facts with Local Variables

Explicitly fixing variables in a proof context is not only important for avoiding name clashes. If a fact is exported from a proof context, all fixed local variables are replaced by unknowns, other variables remain unchanged. Since unification only works for unknowns, it makes a difference whether a fact uses a local variable or a variable which originates from an enclosing context or is free.

The proposition $x < 3 \implies x < 5$ can be proven by the statements

```
fix  $y :: \text{nat}$   
assume " $y < 3$ "  
then show " $y < 5$ "  $\langle \text{proof} \rangle$ 
```

because when the fact $y < 5$ is exported, the assumption is added (as described in Section 2.2.8) and then variable y is replaced by the unknown $?y$ because y has been locally fixed. The result is the rule $?y < 3 \implies ?y < 5$ which unifies with the proposition.

If, instead, y is not fixed, the sequence

```
assume "(y : nat) < 3"
then have "y < 5" <proof>
```

still works and the local fact $y < 5$ can be proven, but it cannot be used with the **show** statement to prove the proposition $x < 3 \implies x < 5$, because the exported rule is now $y < 3 \implies y < 5$ which does not unify with the proposition, it contains a different variable instead of an unknown.

Externally Fixed Variables

Like assumed facts inserted by **assumes** (see Section 2.2.8), local variables added from outside by **fixes** instead of locally by **fix** are part of the *outer* proof context of a theorem. Therefore they are not converted to unknowns when a fact is exported by a **show** statement. Actually, this is not necessary, because variables fixed by **fixes** may directly occur in assumptions specified by **assumes** and in the goal specified by **shows**.

In the theorem

```
fixes x : nat
assumes "x < 3"
shows "x < 5" <proof>
```

the $\langle proof \rangle$ can be directly written using the local variable x .

2.2.10 Obtaining Variables

Local variables may also be introduced together with a fact which allows to determine their values. This is done using a statement of the form

```
obtain x1 ... xm where "prop" <proof>
```

where *prop* is a proposition in inner syntax which contains the variables $x_1 \dots x_m$. Like for variables introduced by **fix** the variables can be grouped by **and** and types can be specified for (some of) the groups.

The proposition usually relates the values of the new variables to values of existing variables (which may be local or come from the environment). In the simplest case the proposition directly specifies terms for the new variables, such as in

```
fix x::nat
obtain y z where "y = x + 3 ∧ z = x + 5" ⟨proof⟩
```

But it is also possible to specify the values indirectly:

```
fix x::nat
obtain y z where "x = y - 3 ∧ y + z = 2*x + 8" ⟨proof⟩
```

Here the proposition may be considered to be an additional assumption which is added to the proof context.

Proving *obtain* Statements

Actually, several propositions may be specified in an **obtain** statement:

```
obtain x1 ... xm where "prop1" ... "propn" ⟨proof⟩
```

The propositions may be grouped by **and** and the groups can be named as usual. This **obtain** statement has a similar meaning as the statements

```
fix x1 ... xm
assume "prop1" ... "propn"
```

but there is one important difference: the propositions in an **obtain** statement must be redundant in the local proof context.

That is the reason why an **obtain** statement is a goal statement and includes a proof. The proof must prove the redundancy of the propositions, which may be stated in the following way: if any other proposition can be derived from them in the local proof context it must be possible to also derive it without the propositions. This can be stated formally as

$$(\bigwedge x_1 \dots x_m. \llbracket prop_1; \dots; prop_n \rrbracket \implies P) \implies P$$

which is exactly the goal to be proven for the **obtain** statement.

Consider the statements

```
fix x::nat
obtain y where "x = 2*y" ⟨proof⟩
```

This proposition is not redundant, because it implies that **x** must be even. Therefore no proof exists.

Note that after a successful proof of an **obtain** statement the current facts are the propositions specified in the statement, not the proven redundancy statement. Input facts may be passed to **obtain** statements. Like for the other goal statements, they are input to the *⟨proof⟩*.

Exporting Facts after Obtaining Variables

Unlike facts assumed by an **assume** statement (see Section 2.2.8) the propositions in an **obtain** statement are *not* added as assumptions when a fact F is exported from the local context. This is correct, since they have been proven to be redundant, therefore they can be omitted.

However, that implies that an exported fact F may not refer to variables introduced by an **obtain** statement, because the information provided by the propositions about them gets lost during the export.

2.2.11 Term Abbreviations

A term abbreviation is a name for a proposition or a term in it.

Defining Term Abbreviations

A term abbreviation can be defined by a statement of the form

```
let ?name = "term"
```

Afterwards the name is “bound” to the term and can be used in place of the term in propositions and other terms, as in:

```
let ?lhs = "2*x+3"  
let ?rhs = "2*5+3"  
assume "x < 5"  
have "?lhs ≤ ?rhs" ⟨proof⟩
```

The name *?thesis* (see Section 2.2.5) is a term abbreviation of this kind.

A **let** statement can define several term abbreviations in the form

```
let ?name1 = "term1" and ... and ?namen = "termn"
```

A **let** statement can occur everywhere in mode *proof(state)*. However, it does not preserve the current facts, the fact set *this* becomes undefined by it.

Pattern Matching

Note that term abbreviations have the form of “unknowns” (see Section 2.1.7), although they are defined (“bound”). The reason is that they are actually defined by unification.

The more general form of a **let** statement is

```
let "pattern" = "term"
```

where *pattern* is a term which may contain unbound unknowns. As usual, if the pattern consists of a single unknown, the quotes may be omitted. The **let** statement unifies *pattern* and *term*, i.e., it determines terms to substitute for the unknowns, so that the pattern becomes syntactically equal to *term*. If that is not possible, an error is signaled, otherwise the unknowns are bound to the substituting terms. Note that the equals sign belongs to the outer syntax, therefore both the pattern and the term must be quoted separately. The **let** statement

```
let "?lhs ≤ ?rhs" = "2*x+3 ≤ 2*5+3"
```

binds the unknowns to the same terms as the two **let** statements above.

The *term* may contain unknowns which are already bound. They are substituted by their bound terms before the pattern matching is performed. Thus the term can be constructed with the help of abbreviation which have been defined previously. A useful example is matching a pattern with *?thesis*:

```
theorem "x < 5 ⇒ 2*x+3 ≤ 2*5+3"
```

```
proof method
```

```
  let "?lhs ≤ ?rhs" = ?thesis
```

```
  ...
```

to reuse parts of the conclusion in the proof.

Note that the unknowns are only bound at the end of the whole **let** statement. In the form

```
let "pattern1" = "term1" and ... and "patternn" = "termn"
```

the unknowns in *pattern_i* cannot be used to build *term_{i+1}* because they are not yet bound. In contrast, in the sequence of **let** statements

```
let "pattern1" = "term1"
```

```
  ...
```

```
let "patternn" = "termn"
```

the unknowns in *pattern_i* can be used to build *term_{i+1}* because they are already bound.

If a bound unknown occurs in the pattern its bound term is ignored and the unknown is rebound according to the pattern matching. In particular, it does not imply that the old and new bound terms must be equal, they are completely independent.

If a part of the term is irrelevant and need not be bound the dummy unknown “_” (underscore) can be used to match it in the pattern without binding an unknown to it:

```
let "_ ≤ ?rhs" = "2*x+3 ≤ 2*5+3"
```

will only bind *?rhs* to the term on the righthand side.

If the term internally binds variables which are used in a subterm, the subterm cannot be matched separately by an unknown because then the variable bindings would be lost. Thus the statement

```
let "λx. ?t" = "λx. x+1"
```

will fail to bind *?t* to *x+1* whereas

```
let "λx. x+?t" = "λx. x+1"
```

will successfully bind *?t* to *1* since the bound variable *x* does not occur in it.

2.2.12 Accumulating Facts

Instead of giving individual names to facts in the proof context, facts can be collected in named fact sets. Isabelle supports the specific fact set named *calculation* and provides statements for managing it.

The fact set *calculation* is intended to accumulate current facts for later use. Therefore it is typically initialized by the statement

```
note calculation = this
```

and afterwards it is extended by several statements

```
note calculation = calculation this
```

After the last extension the facts in the set are chained to the next proof:

```
note calculation = calculation this then have ...
```

Support for Fact Accumulation

Isabelle supports this management of *calculation* with two statements. The statement

```
moreover
```

is equivalent to `note calculation = this` when it occurs the first time in a context, afterwards it behaves like `note calculation = calculation this` but without making *calculation* current, instead, it leaves the current fact(s) unchanged. The statement

```
ultimately
```

is equivalent to **note** *calculation* = *calculation this then*, i.e., it adds the current facts to the set, makes the set current, and chains it to the next goal statement.

Due to the block structure of nested proofs, the *calculation* set can be reused in nested contexts without affecting the set in the enclosing context. The first occurrence of **moreover** in the nested context initializes a fresh local *calculation* set. Therefore fact accumulation is always local to the current proof context.

Accumulating Facts in a Nested Context

Fact accumulation can be used for collecting the facts in a nested context for export (see Section 2.2.7) without using explicit names for them:

```
... {
  have "prop1" <proof>1
  moreover have "prop2" <proof>2
  ...
  moreover have "propn" <proof>n
  moreover note calculation
}
```

Accumulating Facts for Joining Branches

Fact accumulation can also be used for collecting the facts at the end of joined fact branches in a proof and inputting them to the joining step. A corresponding proof pattern for two branches which join at fact *F* is

```
have "F11" <proof>11
then have "F12" <proof>12
...
then have "F1m" <proof>1m
moreover have "F21" <proof>21
then have "F22" <proof>22
...
then have "F2n" <proof>2n
ultimately have "F" <proof>
```

The **moreover** statement starts the second branch and saves the fact *F_{1m}* to *calculation*. The **ultimately** statement saves the fact *F_{2n}* to *calculation* and then inputs the set to the proof of *F*.

Note that **moreover** does not chain the current facts to the following goal statement.

Using nested contexts sub-branches can be constructed and joined in the same way.

2.2.13 Equational Reasoning

Often informal proofs on paper are written in the style

$$2*(x+3) = 2*x+6 \leq 3*x+6 < 3*x+9 = 3*(x+3)$$

to derive the fact $2*(x+3) < 3*(x+3)$. Note that the formula shown above is not a wellformed proposition because of several occurrences of the toplevel relation symbols $=$, \leq and $<$. Instead, the formula is meant as an abbreviated notation of the fact sequence

$$2*(x+3) = 2*x+6, 2*x+6 \leq 3*x+6, 3*x+6 < 3*x+9, 3*x+9 = 3*(x+3)$$

which sketches a proof for $2*(x+3) < 3*(x+3)$. This way of constructing a proof is called “equational reasoning” which is a specific form of forward reasoning.

Transitivity Rules

The full proof needs additional facts which must be inserted into the sequence. From the first two facts the fact $2*(x+3) \leq 3*x+6$ is derived, then with the third fact the fact $2*(x+3) < 3*x+9$ is derived, and finally with the fourth fact the conclusion $2*(x+3) < 3*(x+3)$ is reached. The general pattern of these additional derivations can be stated as the derivation rules $\llbracket a = b; b \leq c \rrbracket \implies a \leq c$, $\llbracket a \leq b; b < c \rrbracket \implies a < c$, and $\llbracket a < b; b = c \rrbracket \implies a < c$.

Rules of this form are called “transitivity rules”. They are valid for relations such as equality, equivalence, orderings, and combinations thereof.

This leads to the general form of a proof constructed by equational reasoning: every forward reasoning step starts at a fact F_i of the form $a \ r \ b$ where r is a relation symbol. It proves an intermediate fact H_i of the form $b \ r \ c$ where r is the same or another relation symbol and uses a transitivity rule to prove the fact F_{i+1} which is $a \ r \ c$. In this way it constructs a linear sequence of facts which leads to the conclusion.

The intermediate facts H_i are usually proven from assumptions or external facts, or they may have a more complex proof which forms an own fact branch which ends at H_i and is joined with the main branch at F_{i+1} with the help of the transitivity rule.

Support for Equational Reasoning

Isabelle supports equational reasoning in the following form. It provides the statement

also

which expects that the set *calculation* contains the fact F_i and the current fact *this* is the fact H_i . It automatically selects an adequate transitivity rule, uses it to derive the fact F_{i+1} and replaces F_i in *calculation* by it. Upon its first use in a proof context **also** simply stores the current fact *this* in *calculation*. The statement

finally

behaves in the same way but additionally makes the resulting fact F_{i+1} current, i.e., puts it into the set *this*, and chains it into the next goal statement. In other words, **finally** is equivalent to **also from calculation**.

Note that **also** behaves like **moreover** and **finally** behaves like **ultimately**, both with additional application of the transitivity rule.

Additionally, Isabelle automatically maintains the term abbreviation ... (which is the three-dot-symbol available for input in the Symbols panel (see Section 1.2.5) of the interactive editor in tab “Punctuation”) for the term on the right hand side of the current fact. Together, the example equational reasoning proof from above can be written

```
have "2*(x+3) = 2*x+6" <proof>
also have "... ≤ 3*x+6" <proof>
also have "... < 3*x+9" <proof>
also have "... = 3*(x+3)" <proof>
finally show ?thesis <proof>
```

where *?thesis* abbreviates the conclusion $2*(x+3) < 3*(x+3)$. This form is quite close to the informal paper style of the proof.

Determining Transitivity Rules

To automatically determine the transitivity rule used by **also** or **finally**, Isabelle maintains the dynamic fact set (see Section 2.1.7) named *trans*. It selects a rule from that set according to the relation symbols used in the facts in *calculation* and *this*.

A transitivity rule which is not in *trans* can be explicitly specified by name in the form

```
also (name)
finally (name)
```

2.3 Proof Methods

The basic building blocks of Isabelle proofs are the proof methods which modify the goal state. If there are several goals in the goal state it depends

on the specific method which goals are affected by it. In most cases only the first goal is affected.

2.3.1 The empty Method

The empty method is denoted by a single minus sign

-

If no input facts are passed to it, it does nothing, it does not alter the goal state.

The empty method is useful at the beginning of a structured proof of the form

proof method $FS_1 \dots FS_n$ **qed**

If the forward reasoning steps $FS_1 \dots FS_n$ shall process the unmodified original goal the empty method must be specified for *method*, thus the structured proof becomes

proof - $FS_1 \dots FS_n$ **qed**

Note that it is possible to syntactically omit the *method* completely, but then it defaults to the method named *standard* which alters the goal state (see below).

If input facts are passed to the empty method, it affects all goals by inserting the input facts as assumptions after the existing assumptions. If the input facts are F_1, \dots, F_m a goal of the form $\llbracket A_1; \dots; A_n \rrbracket \implies C$ is replaced by the goal $\llbracket A_1; \dots; A_n; F_1, \dots, F_m \rrbracket \implies C$. This makes it possible to assume the input facts using **assume** statements in the proof for the goal:

have $"F_i"$ $\langle proof \rangle$
then have $"F_{i+1}"$ **proof - assume** $"F_i"$ **then ... qed**

If the input facts would not have been inserted into the goal, a subsequent **show** statement could not refine it.

However, an **assume** statement for an input fact usually has to repeat the corresponding proposition which has already been specified in an enclosing context where it has been proven as a fact. Then it is easier to name the fact in the enclosing context and refer to it by name instead of assuming it:

have f_i : $"F_i"$ $\langle proof \rangle$
have $"F_{i+1}"$ **proof - from** f_i **... qed**

Since there are no other ways to use input facts in a structured proof with empty initial method than re-specifying them in **assume** statements, it is normally useless to input facts into such a proof.

2.3.2 Rule Application

As described in Section 2.2.2 the basic step in the construction of a proof is to establish the connection between a fact F_i and a fact F_{i+1} in the fact sequence. Assume that there is already a valid derivation rule $RA_i \implies RC_i$ named r_i where RA_i unifies with F_i and RC_i unifies with F_{i+1} . Then the connection can be established by applying that rule.

The *rule* Method

A rule is applied by the method

rule name

where *name* is the name of a valid rule. The method only affects the first goal. If that goal has the form $\llbracket A_1; \dots; A_n \rrbracket \implies C$ and the rule referred by *name* has the form $\llbracket RA_1; \dots; RA_m \rrbracket \implies RC$ the method first unifies RC with the goal conclusion C . That yields the specialized rule $\llbracket RA_1'; \dots; RA_m' \rrbracket \implies RC'$ where RC' is syntactically equal to C and every RA_j' results from RA_j by substituting unknowns by the same terms as in RC' . Note that the goal normally does not contain unknowns, therefore C is not modified by the unification. If the unification fails the method cannot be executed on the goal state and an error is signaled. Otherwise the method replaces the goal by the m new goals $\llbracket A_1; \dots; A_n \rrbracket \implies RA_j'$.

If the rule has the form $RA \implies RC$ with only one assumption the method replaces the goal by the single new goal $\llbracket A_1; \dots; A_n \rrbracket \implies RA'$. If the rule is a formula RC without any assumptions the method removes the goal without introducing a new goal.

Using the *rule* Method for Backward Reasoning Steps

Assume that during a proof as described in Section 2.2.2 the intermediate fact sequence $F_{i+1} \implies \dots \implies F_n$ has already been constructed by backward reasoning and the current goal is $F_1 \implies F_{i+1}$. The backward reasoning step

apply (*rule* r_i)

will replace that goal by $F_1 \implies F_i$ and thus extend the fact sequence to $F_1 \implies \dots \implies F_n$. The fact F_i is the specialized assumption RA_i' constructed by the method from the assumption RA_i of rule r_i .

Therefore the fact sequence $F_1 \implies \dots \implies F_n$ of the complete proof can be constructed by the proof script consisting of the backward reasoning steps

apply (*rule* r_{n-1})

```
...
apply(rule  $r_1$ )
```

Note however, that this proof script does not complete the proof, since it results in the goal $F_1 \Rightarrow F_1$. The proof method

```
assumption
```

must be used to process it. The method only affects the first goal. If that goal has the form $\llbracket A_1; \dots; A_n \rrbracket \Rightarrow C$ and one assumption A_i is syntactically equal to C the method removes the goal, otherwise an error is signaled.

Together, the full proof script has the form

```
apply (rule  $r_{n-1}$ )
...
apply(rule  $r_1$ )
apply(assumption)
done
```

If the example from Section 2.2.2 is proven this way the theorem is written together with its proof as

```
theorem " $x < 5 \Rightarrow 2*x+3 \leq 2*5 + 3$ " for  $x :: \text{nat}$ 
  apply(rule example2)
  apply(rule example1)
  apply(assumption)
done
```

Note that the assumption A of the initial goal must be reached exactly by the sequence of rule applications. If it is replaced in the example by the stronger assumption $x < 3$ the rule applications will lead to the goal $x < 3 \Rightarrow x < 5$ which is trivial for the human reader but not applicable to the *assumption* method.

Using the *rule* Method for Forward Reasoning Steps

Assume that during a proof as described in Section 2.2.2 the intermediate fact sequence $F_1 \Rightarrow \dots \Rightarrow F_i$ has already been constructed by forward reasoning, so that the next step is to state fact F_{i+1} and prove it. Using method *rule* the step can be started by

```
have " $F_{i+1}$ " proof (rule  $r_i$ )
```

The goal of this subproof is simply F_{i+1} , so applying the *rule* method with r_i will result in the new goal RA_i' which is F_i , as above. The subproof is not finished, since its goal state is not empty. But the goal is an already known fact. The proof method

fact name

can be used to remove that goal. The method only affects the first goal. If the fact referred by *name* unifies with it, the goal is removed, otherwise an error is signaled.

Using this method the forward reasoning step can be completed as

```
have "Fi+1" proof (rule ri) qed (fact fi)
```

if F_i has been named f_i . This can be abbreviated (see Section 2.2.3) to

```
have "Fi+1" by (rule ri) (fact fi)
```

Therefore the fact sequence $F_1 \Rightarrow \dots \Rightarrow F_n$ of the complete proof can be constructed by the structured proof of the form

```
proof -
  assume f1: "F1"
  have f2: "F2" by (rule r1) (fact f1)
  ...
  have fn-1: "Fn-1" by (rule rn-2) (fact fn-2)
  show "Fn" by (rule rn-1) (fact fn-1)
qed
```

where the last fact F_n is usually the conclusion \mathcal{C} and can be specified as *?thesis*.

If the example from Section 2.2.2 is proven this way the theorem is written together with its proof as

```
theorem "x < 5  $\implies$  2*x+3  $\leq$  2*5 + 3" for x :: nat
proof -
  assume f1: "x < 5"
  have f2: "2*x  $\leq$  2*5" by (rule example1) (fact f1)
  show ?thesis by (rule example2) (fact f2)
qed
```

The *fact* method can be specified in the form

fact

without naming the fact to be used. Then it selects a fact automatically. It uses the first fact from the proof context which unifies with the goal. If there is no such fact in the proof context an error is signaled.

Thus the example can be written without naming the facts as

```
theorem "x < 5  $\implies$  2*x+3  $\leq$  2*5 + 3" for x :: nat
proof -
```

```

    assume "x < 5"
    have "2*x ≤ 2*5" by (rule example1) fact
    show ?thesis by (rule example2) fact
qed

```

Input Facts for the *rule* Method

If input facts F_1, \dots, F_n are passed to the *rule* method, they are used to remove assumptions from the rule applied by the method. If the rule has the form $\llbracket RA_1; \dots; RA_{n+m} \rrbracket \implies RC$ and every fact F_i unifies with assumption RA_i the first n assumptions are removed and the rule becomes $\llbracket RA_{n+1}; \dots; RA_{n+m} \rrbracket \implies RC$. Then it is applied to the first goal in the usual way. If there are more facts than assumptions or if a fact does not unify, an error is signaled. This allows to establish the connection from a fact F_i to F_{i+1} in a fact chain by a forward reasoning step of the form

```

from  $f_i$  have " $F_{i+1}$ " by (rule  $r_i$ )

```

where f_i names the fact F_i . When it is input to the goal statement it is passed to the *rule* method and removes the assumption from the applied rule $RA_i \implies RC_i$, resulting in the assumption-less “rule” RC_i . When it is applied to the goal F_{i+1} it unifies and removes the goal, thus the subproof is complete.

For the fact sequence chaining can be used to write a structured proof without naming the facts:

```

proof -
  assume " $F_1$ "
  then have " $F_2$ " by (rule  $r_1$ )
  ...
  then have " $F_{n-1}$ " by (rule  $r_{n-2}$ )
  then show " $F_n$ " by (rule  $r_{n-1}$ )
qed

```

If the example from Section 2.2.2 is proven this way the theorem is written together with its proof as

```

theorem "x < 5  $\implies$  2*x+3 ≤ 2*5 + 3" for x :: nat
proof -
  assume "x < 5"
  then have "2*x ≤ 2*5" by (rule example1)
  then show ?thesis by (rule example2)
qed

```

The Method *this*

Rule application can also be done by the method

this

Instead of applying a named method, it applies the input fact as rule to the first goal.

If several input facts are given, the method applies them exactly in the given order. Therefore the fact sequence can also be constructed by a structured proof of the form:

```
proof -
  assume "F1"
  with r1 have "F2" by this
  ...
  with rn-2 have "Fn-1" by this
  with rn-1 show "Fn" by this
qed
```

The **with** statement inserts the explicitly named facts *before* the current facts. Therefore every goal statement for F_i gets as input the rule r_{i-1} followed by the chained fact F_{i-1} . The method *this* first applies the rule which replaces the goal by F_{i-1} . Then it applies the fact F_{i-1} as rule to this goal which removes it and finishes the subproof.

The proof

by *this*

can be abbreviated by . (a single dot).

Therefore the example from Section 2.2.2 can also be proven in the form

```
theorem "x < 5  $\implies$  2*x+3  $\leq$  2*5 + 3" for x :: nat
proof -
  assume "x < 5"
  with example1 have "2*x  $\leq$  2*5" .
  with example2 show ?thesis .
qed
```

Automatic Rule Selection

The *rule* method can be specified in the form

rule

without naming the rule to be applied. Then it selects a rule automatically. It uses the first rule from the dynamic fact set *intro* for which the conclusion unifies with the goal conclusion. If there is no such rule in the set an error is signaled.

If the rules *example1* and *example2* would be in the *intro* set, the example proof could be written as

```
theorem "x < 5  $\implies$  2*x+3  $\leq$  2*5 + 3" for x :: nat
proof -
  assume "x < 5"
  then have "2*x  $\leq$  2*5" by rule
  then show ?thesis by rule
qed
```

However, the set *intro* is intended by Isabelle for a specific kind of rules called “introduction rules”. In such a rule the toplevel operator of the conclusion does not occur in any assumption, hence it is “introduced” by the rule. When an introduction rule is applied backwards, the operator is removed from the goal. This can be iterated to “deconstruct” the goal, some proofs can be written using this technique, however, the content of the set must be designed very carefully to not run into cycles.

Since in the rule *example2* the toplevel operator \leq occurs in the assumption it is not an introduction rule and should not be added to *intro*. Rule *example1* is an introduction rule but would interfere with predefined rules in the set.

The *standard* Method

The method

standard

is a method alias which can be varied for different Isabelle applications. Usually it is an alias for the *rule* method.

The *standard* method is the default, if no method is specified as the initial step in a structured proof. Thus

```
proof FS1 ... FSn qed
```

is an abbreviation for

```
proof standard FS1 ... FSn qed
```

Note that the *standard* method will usually affect the goal by applying an introduction rule to it. That may be useful in some cases, but it has to be taken into account when writing the forward reasoning steps of the proof.

For example, in Isabelle HOL there is an introduction rule in *intro* which splits a logical conjunction into two subgoals (see Section ??). Therefore the proof of a conjunction $P \wedge Q$ can be written

```
proof
show P <proof>
show Q <proof>
qed
```

because the goal is split by the *standard* method. If the empty method is used instead, the proof has to be of the form

```
proof -
show "P ∧ Q" <proof>
qed
```

because the goal is not modified.

In the abbreviated form **by** *method* of a structured proof the method cannot be omitted, but the proof **by** *standard* can be abbreviated to

```
..
```

(two dots). It can be used as complete proof for a proposition which can be proven by a single automatic rule application. Hence, if the rules *example1* and *example2* would be in the *intro* set, the example proof could be further abbreviated as

```
theorem "x < 5  $\implies$  2*x+3  $\leq$  2*5 + 3" for x :: nat
proof -
  assume "x < 5"
  then have "2*x  $\leq$  2*5" ..
  then show ?thesis ..
qed
```

2.3.3 Composed Proof Methods

Proof methods can be composed from simpler methods with the help of “method expressions”. A method expression has one of the following forms:

- m_1, \dots, m_n : a sequence of methods which are applied in their order,
- $m_1; \dots; m_n$: a sequence of methods where each is applied to the goals created by the previous method,
- $m_1 / \dots / m_n$: a sequence of methods where only the first applicable method is applied,

- $m[n]$: the method m is applied to the first n goals,
- $m?$: the method m is applied if it is applicable,
- $m+$: the method m is applied once and then repeated as long as it is applicable.

Parentheses are used to structure and nest composed methods.

Composed methods can be used to combine backward reasoning steps to a single step. Using composed methods the example backward reasoning proof from Section 2.3.2 can be written as

```
theorem "x < 5  $\implies$  2*x+3  $\leq$  2*5 + 3" for x :: nat
  apply (rule example2, rule example1, assumption)
  done
```

In particular, it is possible to apply an arbitrarily complex backward reasoning step as initial method in a structured proof. Using composed methods the first example forward reasoning proof can be written

```
theorem "x < 5  $\implies$  2*x+3  $\leq$  2*5 + 3" for x :: nat
proof -
  assume "x < 5"
  have "2*x  $\leq$  2*5" by (rule example1, fact)
  show ?thesis by (rule example2, fact)
qed
```

2.3.4 The Simplifier

A common proof technique is “rewriting”. If it is known that a term a is equal to a term b , some occurrences of a in a proposition can be replaced by b without changing the validity of the proposition.

Equality of two terms a and b can be expressed by the proposition $a = b$. If that proposition has been proven to be valid, i.e., is a fact, a can be substituted by b and vice versa in goals during a proof.

The *subst* Method

Rewriting is performed by the method

```
subst name
```

where *name* references an equality fact. The method only affects the first goal. If the referenced fact has the form $a = b$ the method replaces the first occurrence of *a* in the goal conclusion by *b*. The order of the terms in the equality fact matters, the method always substitutes the term on the left by that on the right.

If the equality contains unknowns unification is used: *a* is unified with every sub-term of the goal conclusion, the first match is replaced by *b'* which is *b* after substituting unknowns in the same way as in *a*. If there is no match of *a* in the goal conclusion an error is signaled.

For a goal $\llbracket A_1; \dots; A_n \rrbracket \Rightarrow C$ the method only rewrites in the conclusion *C*. The first match in the assumptions $A_1 \dots A_n$ can be substituted by the form

`subst (asm) name`

If not only the first match shall be substituted, a number of the match or a range of numbers may be specified in both forms as in

`subst (asm) (i..j) name`

The equality fact can also be a meta equality of the form $a \equiv b$. Therefore the method can be used to expand constant definitions. After the definition

definition `"inc x \equiv x + 1"`

the method `subst inc_def` will rewrite the first occurrence of a function application (`inc t`) in the goal conclusion to $(t + 1)$. Remember from Section 2.1.4 that the defining equation is automatically named `inc_def`. Note the use of unification to handle the actual argument term *t*.

The equality fact may be conditional, i.e., it may be a derivation rule with assumptions of the form $\llbracket RA_1; \dots; RA_m \rrbracket \Rightarrow a = b$. When the `subst` method applies a conditional equation of this form to a goal $\llbracket A_1; \dots; A_n \rrbracket \Rightarrow C$, it adds the goals $\llbracket A_1; \dots; A_n \rrbracket \Rightarrow RA_i'$ to the goal state after rewriting, where RA_i' result from RA_i by the unification of *a* in *C*. These goals are inserted before the original goal, so the next method application will usually process the goal $\llbracket A_1; \dots; A_n \rrbracket \Rightarrow RA_1'$.

As an example if there are theorems

theorem `eq1: "n = 10 \Rightarrow n+3 = 13" for n::nat <proof>`
theorem `eq2: "n = 5 \Rightarrow 2*n = 10" for n::nat <proof>`

the method `subst (2) eq2` replaces the goal $(x::nat) < 5 \Rightarrow 2*x+3 \leq 2*5 + 3$ by the goals

$x < 5 \Rightarrow 5 = 5$
 $x < 5 \Rightarrow 2 * x + 3 \leq 10 + 3$

where the first is trivial (but still must be removed by applying a rule). The second goal is replaced by the method `subst (2) eq1` by

$$\begin{aligned} x < 5 &\implies 10 = 10 \\ x < 5 &\implies 2 * x + 3 \leq 13 \end{aligned}$$

Note that the method `subst eq2` would unify $2*n$ with the first match $2*x$ in the original goal and replace it by

$$\begin{aligned} x < 5 &\implies x = 5 \\ x < 5 &\implies 10 + 3 \leq 2 * 5 + 3 \end{aligned}$$

where the first goal cannot be proven because it is invalid.

Simplification

If the term b in an equation $a = b$ is in some sense “simpler” than a , the goal will also become simpler by successful rewriting with the equation. If there are several such equations a goal can be replaced by successively simpler goals by rewriting with these equations. This technique can contribute to the goal’s proof and is called “simplification”.

Basically, simplification uses a set of equations and searches an equation in the set where the left hand side unifies with a sub-term in the goal, then substitutes it. This step is repeated until no sub-term in the goal unifies with a left hand side in an equation in the set.

It is apparent that great care must be taken when populating the set of equations, otherwise simplification may not terminate. If two equations $a = b$ and $b = a$ are in the set simplification will exchange matching terms forever. If an equation $a = a+0$ is in the set, a term matching a will be replaced by an ever growing sum with zeroes.

Simplification with a set of definitional equations from constant definitions (see Section 2.1.4) always terminates. Since constant definitions cannot be recursive, every substitution removes one occurrence of a defined constant from the goal. Simplification terminates if no defined constant from the set remains in the goal. Although the resulting goal usually is larger than the original goal, it is simpler in the sense that it uses fewer defined constants.

If the set contains conditional equations, simplification may produce additional goals. Then simplification is applied to these goals as well. Together, simplification may turn a single complex goal into a large number of simple goals, but it cannot reduce the number of goals. Therefore simplification is usually complemented by methods which remove trivial goals like $x = x$, $A \implies A$, and True . Such an extended simplification may completely solve and remove the goal to which it is applied.

The *simp* Method

Isabelle supports simplification by the method

simp

which is also called “the simplifier”. It uses the dynamic fact set *simp* as the set of equations, which is also called “the simpset”. The method only affects the first goal. If no equation in the simpset is applicable to it or it is not modified by the applicable equations an error is signaled.

The *simp* method simplifies the whole goal, i.e., it applies rewriting to the conclusion and to all assumptions.

The simpset may contain facts which are not directly equations, but can be converted to an equation. In particular, an arbitrary derivation rule $\llbracket A_1; \dots; A_n \rrbracket \implies C$ can always be converted to the equation $\llbracket A_1; \dots; A_n \rrbracket \implies C = \text{True}$. The simplifier performs this conversion if no other conversion technique applies, therefore the simpset may actually contain arbitrary facts.

The *simp* method also detects several forms of trivial goals and removes them. Thus a complete proof may be performed by a single application of the simplifier in the form

by *simp*

In Isabelle HOL (see Section 3) the simpset is populated with a large number of facts which make the simplifier a very useful proof tool. Actually all examples of facts used in the previous sections can be proven by the simplifier:

```
theorem example1: "(x::nat) < c  $\implies$  n*x  $\leq$  n*c" by simp
theorem example2: "(x::nat)  $\leq$  c  $\implies$  x + m  $\leq$  c + m" by simp
theorem "(x::nat) < 5  $\implies$  2*x+3  $\leq$  2*5 + 3" by simp
theorem eq1: "n = 10  $\implies$  n+3 = 13" for n::nat by simp
theorem eq2: "n = 5  $\implies$  2*n = 10" for n::nat by simp
```

Configuring the Simplifier

The simplifier can be configured by modifying the equations it uses. The form

simp add: *name*₁ ... *name*_n

uses the facts *name*₁, ..., *name*_n in addition to the facts in the simpset for its rewriting steps. The form

simp del: *name*₁ ... *name*_n

uses only the facts from the simpset without the facts *name*₁, ..., *name*_{*n*}, and the form

*simp only: name*₁ ... *name*_{*n*}

uses only the facts *name*₁, ..., *name*_{*n*}. The three forms can be arbitrarily combined.

As usual, a theorem may be added permanently to the simpset as described in Section 2.1.7 by specifying it as

theorem [*simp*]: "*prop*" <*proof*>

and the defining equation of a definition can be added by

definition *name::type* **where** [*simp*]: "*name* ≡ *term*"

Adding own constant definitions to the simplifier is a common technique to expand the definition during simplification. However, this may also have a negative effect: If an equation has been specified using the defined constant, it is no more applicable for rewriting after expanding the definition. Note that the facts in the simpset and the facts provided by *add:*, *del:*, and *only:* are not simplified themselves, the defined constant will not be expanded there.

Therefore it is usually not recommended to add defining equations to the simpset permanently. Instead, they can be specified by *add:* when they really shall be expanded during simplification.

Splitting Terms

There are certain terms in which the simplifier will not apply its simpset rules. A typical example are terms with an internal case distinction (see Section 3.1.3). To process such terms in a goal conclusion the terms must be split. Splitting a term usually results in several new goals with simpler terms which are then further processed by the simplifier.

Term splitting is done by applying specific rules to the goal. These rules are called "split rules". Usually split rules are not automatically determined and applied by the simplifier, this must be configured explicitly in the form

*simp split: name*₁ ... *name*_{*n*}

where the *name*_{*i*} are the names of the split rules to use. This configuration can be arbitrarily combined with the other simplifier configuration options.

Input Facts for the *simp* Method

As usual, facts may be input to the *simp* method. Like the empty method (see Section 2.3.1) it inserts these facts as assumptions into the goal, before it starts simplification. Since simplification is also applied to the assumptions, the input facts will be simplified as well.

As a possible effect of this behavior, after simplifying an input fact and the goal conclusion the results may unify, leading to the situation where the goal is removed by the *assumption* method (see Section 2.3.2). This is also done by the simplifier, hence in this way the input fact may contribute to prove the goal.

The *simp_all* Method

The method

simp_all

behaves like the *simp* method but processes all goals. It inserts input facts to all goals in the goal state and simplifies them. If it fails for all goals an error is signaled. Otherwise it simplifies only the goals for which it does not fail. If it achieves to remove all goals the proof is finished.

The *simp_all* method can be configured by *add:*, *del:*, *only:*, and *split:* in the same way as the *simp* method.

The *simp_all* method is useful, if first a method *method* is applied to the goal which splits it into several subgoals which all can be solved by simplification. Then the complete proof can be written as

by *method simp_all*

Debugging the Simplifier

If the simplifier fails, it may be difficult to find out the reason. There are several debugging techniques which may help.

The content of the simpset can be displayed by the command

print_simpset

which may be specified in the proof text in modes *proof(prove)* and *proof(state)* and outside of proofs. In the interactive editor the result is displayed in the Output panel (see Section 1.2.3).

There is also a simplifier trace which displays the successful rewrite steps. It is activated by the command

declare *[[simp_trace_new depth=n]]*

outside a theorem or definition. The number *n* should be atleast 2. When the cursor is positioned on an application of the *simp* method the button “Show trace” can be used in the Simplifier Trace Panel to display the trace in a separate window. See the documentation for more information about how to use the trace.

Another technique is to replace the *simp* method by a sequence of *subst* method applications and explicitly specify the equations which should have been used. To do this for a structured proof, replace it by a proof script for the *subst* methods.

2.3.5 Other Automatic Proof Methods

Isabelle provides several other proof methods which internally perform several steps, like the simplifier.

Automatic Methods

The following list contains automatic methods other than *simp*:

- *blast* mainly applies logical rules and can be used to solve complex logical formulas.
- *clarify* is similar but does not split goals and does not follow unsafe paths. It can be used to show the problem if *blast* fails.
- *auto* combines logical rule application with simplification. It processes all goals and leaves those it cannot solve.
- *clarsimp* combines *clarify* with simplification. It processes only the first goal and usually does not split goals.
- *fastforce* uses more techniques than *auto*, but processes only the first goal.
- *force* uses even more techniques and tries to solve the first goal.

The methods which do simplification can be configured like the simplifier by adding specifications *simp add:*, *simp del:*, *simp only:*, and *split:*. For example, additional simplification rules can be specified for the *auto* method in the form

```
auto simp add: name1 ... namen
```

For more information about these methods see the Isabelle documentation.

Trying Methods

Instead of manually trying several automatic methods it is possible to specify the command

try

anywhere in mode *proof* (*prove*), i.e. at the beginning of a proof or in a proof script. It will try many automatic proof methods and describe the result in the Output window. It may take some time until results are displayed, in particular, if the goal is invalid and cannot be proven.

If **try** finds a proof for one or more goals it displays it as a single (composed) proof method, which, by clicking on it can be copied to the cursor position in the text area. The **try** command must be removed manually.

If **try** tells you that the goal can be “directly solved” by some fact, you can prove it by the *fact* method, but that also means that there is already a fact of the same form and your theorem is redundant.

It may also be the case that **try** finds a counterexample, meaning that the goal is invalid and cannot be proven.

2.4 Case Based Proofs

If a proof method splits a goal at the beginning of a structured proof the resulting subgoals must be proven separately, each with an own **show** statement. Some proof methods provide support for this by initializing a proof context for each goal resulting from the method application.

2.4.1 Goal Cases

If a method supports this, it creates a “case” for each corresponding goal. The cases are named. Using these names a proof context can be populated with elements prepared by the method.

Named Contexts

Every case name is associated with proof context elements, thus it can be seen as a named context which has been prepared by the method for later use. Such named contexts may contain facts, local variables, and term abbreviations. The actual content depends on the proof method and the goals for which the cases are created.

When the named context is used in a proof, its content is “injected” into the current proof context. Usually a named context is used in this way to initialize a new nested context immediately after its beginning.

The case Command

A case may be injected into the current proof context by the command

```
case name
```

where *name* is the case name. It mainly has the effect of the sequence

```
fix  $x_1 \dots x_k$   
let  $?a_1 = t_1$  and  $\dots ?a_m = t_m$   
assume name: " $A_1$ "  $\dots$  " $A_n$ "
```

where $x_1 \dots x_k$ are the local variables, $?a_1, \dots, ?a_m$ are the term abbreviations, and A_1, \dots, A_n are the facts in the named context of the case. The facts are injected as assumptions and the set of these assumptions is named using the case name. Moreover, like the **assume** statement, the **case** command makes the assumed facts current.

Instead of using the case name for naming the assumptions an explicit assumption name *aname* may be specified:

```
case aname: name
```

The local variables $x_1 \dots x_k$ are fixed by the **case** command but are hidden, they cannot be used in the subsequent proof text. If they should be used, explicit names must be specified for them in the form

```
case (name  $y_1 \dots y_j$ )
```

Then the names $y_1 \dots y_j$ can be used to reference the fixed variables in the current proof context. If fewer names are specified only the first variables are named, if more names are specified than there are local variables in the case an error is signaled.

When methods create named contexts they usually only define the term abbreviation *?case* for the conclusion of the corresponding goal.

Proof Structure with Cases

The usual proof structure using cases consists of a sequence of nested contexts (see Section 2.2.3). At its beginning each context is initialized by a **case** command, at its end it uses a **show** statement to remove the corresponding goal:

```
proof method  
case name1  
 $\dots$   
show ?case  $\langle proof \rangle$ 
```

```

next
case name2
...
show ?case ⟨proof⟩
next
...
next
case namen
...
show ?case ⟨proof⟩
qed

```

Every **show** command uses the local term abbreviation *?case* to refer to the conclusion of the corresponding goal.

In the interactive editor, when the cursor is positioned on **proof method** where the method supports cases, a skeleton of such a proof using the specific cases provided by the method is displayed in the Output panel. By clicking on it it may be copied into the text area immediately after the method specification.

2.4.2 The *goal_cases* Method

The simplest method with support for cases is

goal_cases

Without modifying the goal state it creates a named case for every existing goal. Input facts are ignored.

For a goal $\bigwedge x_1 \dots x_m. \llbracket A_1; \dots; A_n \rrbracket \implies C$ the created named context contains the local variables $x_1 \dots x_m$, the facts A_1, \dots, A_n , and the term abbreviation *?case* bound to C . If the goal contains variables which are not explicitly bound by \bigwedge these variables are not added to the context.

The effect is that if no other variables are fixed and no other facts are assumed a statement **show** *?case* after the corresponding **case** command will refine the goal.

The cases are named by numbers starting with 1. If other names should be used they can be specified as arguments to the method:

goal_cases name₁ ... name_n

If fewer names are specified than goals are present, only for the first n goals cases are created. If more names are specified an error is signaled.

When *goal_cases* is used in a composed proof method it can provide cases for the goals produced by arbitrary other methods:

proof (*method*, *goal_cases*)

provides cases for all goals existing after *method* has been applied. If *method* does not split the goal there will be only one case. This can be useful to work with a goal produced by *method*. In particular, the conclusion of that goal is available as *?case*.

Note that the proof state(s) resulting from *goal_cases* are not visible for the reader of the proof. Therefore it should only be applied if the goals produced by *method* are apparent. This can be supported by defining an abbreviated form of the conclusion by

let "pattern" = *?case*

2.4.3 Case Based Reasoning

In “case based reasoning” a goal is proven by processing “all possible cases” for an additional assumption separately. If the conclusion can be proven for all these cases and the cases cover all possibilities, the conclusion holds generally.

In the simplest form a single proposition is used as additional assumption. Then there are only two cases: if the proposition is *True* or if it is *False*.

Consider the derivation rule $(x::nat) < c \implies n*x \leq n*c$ from Section 2.1.6. As additional assumption the proposition that *n* is zero can be used. Then there are the two cases $n = 0$ and $n \neq 0$ and clearly these cover all possibilities. Using the first case as assumption implies that $n*x$ and $n*c$ are both zero and thus $n*x = n*c$. Using the second case as assumption together with the original assumption implies that $n*x < n*c$. Together the conclusion $n*x \leq n*c$ follows.

Case Rules

To prove a goal in this way it must be split into a separate goal for each case. All these goals must have the same conclusion but differ in the additional assumptions. This splitting can be done by applying a meta-rule of the form

$$\llbracket Q_1 \implies ?P; \dots; Q_n \implies ?P \rrbracket \implies ?P$$

where $Q_1 \dots Q_n$ are all cases of the additional assumption. Such rules are called “case rules”.

When this case rule is applied to a goal $\bigwedge x_1 \dots x_m. \llbracket A_1; \dots; A_n \rrbracket \implies C$ as described in Section 2.3.2, it unifies *?P* with the conclusion *C* and replaces the goal by the *n* goals

$$\begin{aligned} & \bigwedge x_1 \dots x_m. \llbracket A_1; \dots; A_n; Q_1 \rrbracket \implies C \\ & \dots \\ & \bigwedge x_1 \dots x_m. \llbracket A_1; \dots; A_n; Q_n \rrbracket \implies C \end{aligned}$$

where every goal has one of the propositions Q_i as additional assumption.

The case rule is only valid, if the Q_i together cover all possibilities, i.e., if $Q_1 \vee \dots \vee Q_n$ holds. If this has been proven the case rule is available as a fact which can be applied. Since the whole conclusion is the single unknown $?P$ it unifies with every proposition used as conclusion in a goal, hence a case rule can always be applied to arbitrary goals. It depends on the Q_i whether splitting a specific goal with the case rule is useful for proving the goal.

A case rule for testing a natural number for being zero would be

$$\llbracket ?n = 0 \implies ?P; ?n \neq 0 \implies ?P \rrbracket \implies ?P$$

It contains the number to be tested as the unknown $?n$, so that an arbitrary term can be substituted for it. This is not automatically done by unifying $?P$ with the goal's conclusion, thus the rule must be “prepared” for application to a specific goal. To apply it to the goal $(x::nat) < c \implies n*x \leq n*c$ in the intended way the unknown $?n$ must be substituted by the variable n from the goal conclusion. If the prepared rule is then applied to the goal it splits it into the goals

$$\begin{aligned} & \llbracket (x::nat) < c; n = 0 \rrbracket \implies n*x \leq n*c \\ & \llbracket (x::nat) < c; n \neq 0 \rrbracket \implies n*x \leq n*c \end{aligned}$$

which can now be proven separately.

Actually, the much more general case rule

$$\llbracket ?Q \implies ?P; \neg ?Q \implies ?P \rrbracket \implies ?P$$

is used for this purpose. Here the unknown $?Q$ represents the complete proposition to be used as additional assumption, therefore the rule can be used for arbitrary propositions. By substituting the term $n = 0$ for $?Q$ the rule is prepared to be applied in the same way as above.

Case rules may even be more general than shown above. Instead of a single proposition Q_i every case may have locally bound variables and an arbitrary number of assumptions, resulting in the most general form

$$\begin{aligned} & \llbracket \bigwedge x_{11} \dots x_{1k1}. \llbracket Q_{11}; \dots; Q_{1m1} \rrbracket \implies ?P; \\ & \dots; \\ & \llbracket \bigwedge x_{n1} \dots x_{nkn}. \llbracket Q_{n1}; \dots; Q_{nmn} \rrbracket \implies ?P \rrbracket \implies ?P \end{aligned}$$

of a case rule. When applied it splits the goal and adds the variables $x_{i1} \dots x_{iki}$ and the assumptions $Q_{i1} \dots Q_{imi}$ to the i th case.

The cases Method

Case based reasoning can be performed in a structured proof using the method `cases` in the form

```
cases "term" rule: name
```

where `name` is the name of a valid case rule. The method prepares the rule by substituting the specified `term` for the unknown in the assumptions, and applies the rule to the first goal in the goal state.

Additionally, the method creates a named context for every goal resulting from the rule application. The context contains the variables and assumptions specified in the corresponding case in the case rule. For the most general form depicted above the context for the i th case contains the variables $x_{i1} \dots x_{iki}$ and the assumptions $Q_{i1} \dots Q_{imi}$. No term abbreviation `?case` is defined, because the conclusion of every new goal is the same as that of the original goal, thus the existing abbreviation `?thesis` can be used instead.

Often a case rule has only one unknown in the case assumptions. If there are more, several terms may be specified in the `cases` method for preparing the rule. If less terms are specified than there are unknowns in the case assumptions the resulting goals will contain unbound unknowns and cannot be proven. If more terms are specified an error is signaled.

The `cases` method treats input facts like the empty method (see Section 2.3.1) by inserting them as assumptions into the original goal before splitting it.

Like the `rule` method (see Section 2.3.2) the `cases` method supports automatic rule selection for the case rule and may be specified in the form

```
cases "term"
```

Normally the rule is selected according to the type of the specified `term`. In Isabelle HOL (see Section 3) most types have an associated case rule. Only case rules with a single unknown in the case assumptions can be automatically selected in this way.

The rule $\llbracket ?Q \implies ?P; \neg ?Q \implies ?P \rrbracket \implies ?P$ depicted above is associated with type `bool`. Therefore the case splitting for comparing `n` with zero can be done by the method

```
cases "n = 0"
```

The names used for the contexts created by the `cases` method can be specified by attributing the case rule. The case rule for `bool` is attributed to use the case names `True` and `False`. Note that these are names, not the constants for the values of type `bool`.

The proof writer may not know the case names specified by the automatically selected case rule. However, they can be determined from the proof skeleton which is displayed in the interactive editor when the cursor is positioned on the `cases` method (see Section 2.4.1).

Together, a structured proof for the goal $(x : \text{nat}) < c \implies n * x \leq n * c$ with case splitting may have the form

```
proof (cases "n = 0")
case True
...
show ?thesis <proof>
next
case False
...
show ?thesis <proof>
qed
```

The `cases` method adds the assumptions $n=0$ and $n \neq 0$ respectively to the goals of the cases, the `case` commands add them as assumed facts to the local context, so that they are part of the rule exported by the `show` statement and match the assumption in the corresponding goal.

Note that the `case` command adds only the assumptions originating from the case rule. The other assumptions in the original goal (here $x < c$) must be added to the context in the usual ways (see Section 2.2.8) if needed for the proof.

2.4.4 Induction

With induction a goal is proven by processing “all possible cases” for certain values which occur in it. If the goal can be proven for all these cases and the cases cover all possibilities, the goal holds generally. A specific technique is to assume the goal for some values and then prove it for other values. In this way it is possible to cover infinite value sets by proofs for only a finite number of values and steps from values to other values.

Perhaps the best known example of induction is a proposition which is proven for the natural number 0 and the step from a number n to its successor $n+1$, which covers the whole infinite set of natural numbers.

As a (trivial) example consider the proposition $0 \leq n$. To prove that it is valid for all natural numbers n we prove the “base case” where n is 0, which is true because $0 \leq 0$. Then we prove the “induction step”, by assuming that $0 \leq n$ (the “induction hypothesis”) and proving that $0 \leq n+1$ follows, which is true because addition increases the value.

Induction Rules

Like for case based reasoning (see Section 2.4.3) a goal is split into these cases by applying a meta-rule. For induction the splitting can be done by a meta-rule of the form

$$\llbracket P_1 ; \dots ; P_n \rrbracket \implies ?P \text{ ?a}$$

where every P_i is a rule of the form

$$\bigwedge y_{i1} \dots y_{ipi} \cdot \llbracket Q_{i1} ; \dots ; Q_{iqi} \rrbracket \implies ?P \text{ term}_i$$

where the assumptions Q_{ij} may contain the unknown $?P$ but no other unknowns, in particular not $?a$. A rule for a base case usually has no bound variables y_{ij} and no assumptions Q_{ij} , at least the Q_{ij} do not contain $?P$. The remaining rules mostly have only a single assumption Q_{ij} which contains $?P$. Note that the unknown $?a$ only occurs once in the conclusion of the meta-rule and nowhere else. Like the case rules induction rules must be “prepared” for use, this is done by replacing $?a$ by a specific term term . This is the term for which all possible cases shall be processed in the goal. It must have the same type as all term_i in the P_i .

Usually, “all possible cases” means all values of the type of term , then term consists of a single variable which may adopt any values of its type. There are also other forms of induction where more complex terms are used, but they are not presented in this introduction, refer to other Isabelle documentations for them. In the following the unknown $?a$ will always be replaced by a variable x .

When a prepared induction rule is applied to a goal C without bound variables and assumptions as described in Section 2.3.2, it unifies $?P \ x$ with the conclusion C . This has the effect of abstracting C to a (boolean) function PC by identifying all places where x occurs in C and replacing it by the function argument. The function PC is then bound to the unknown $?P$, so that applying $?P$ to the argument x again yields C . The function PC is the property to be proven for all possible argument values. Therefore the cases of the proof can be described by applying $?P$ to the terms term_i for the specific values in the rules P_i for the cases.

The application of the prepared rule results in the n goals

$$\begin{aligned} &\bigwedge y_{11} \dots y_{1p1} \cdot \llbracket Q_{11} ; \dots ; Q_{1q1} \rrbracket \implies PC \text{ term}_1 \\ &\dots \\ &\bigwedge y_{n1} \dots y_{nqn} \cdot \llbracket Q_{n1} ; \dots ; Q_{nqn} \rrbracket \implies PC \text{ term}_n \end{aligned}$$

The induction rule is only valid if the terms term_i cover all possible values of their associated type. If this has been proven the case rule is available as a fact which can be applied. After preparing the induction rule for application,

its conclusion $?P\ x$ matches all propositions which contain the variable x in one or more copies. It depends on the P_i in the rule whether splitting a specific goal with the induction rule is useful for proving the goal.

The real power of induction rules emerges, when a Q_{ij} contains the unknown $?P$. Due to the type associated with $?P$ it must be applied to an argument $term_{ij}$ of the same type as x and the $term_i$. Then the goal resulting from P_i states the property that if Q_{ij} holds when specialized to $PC\ term_{ij}$, PC holds for $term_i$ (an “induction step”). Thus, for covering the possible values of x , the step from $term_{ij}$ to $term_i$ can be repeated arbitrarily often which allows to cover some types with infinite value sets.

An induction rule for the natural numbers is

$$\llbracket ?P\ 0; \bigwedge y. ?P\ y \implies ?P\ (y+1) \rrbracket \implies ?P\ ?a$$

P_1 is the base case, it has no variables and assumptions and only consists of the conclusion $?P\ 0$. P_2 binds the variable y , has one assumption $?P\ y$ and the conclusion $?P\ (y+1)$. P_1 covers the value 0 , P_2 covers the step from a value y to its successor $y+1$, together they cover all possible values of type *nat*.

To apply the rule to the goal $0 \leq n$, it must be prepared by substituting the variable n for the unknown $?a$. Then the rule conclusion $?P\ n$ is unified with the goal which abstracts the goal to the boolean function $PC = (\lambda i. 0 \leq i)$ and substitutes it for all occurrences of $?P$. This results in the rule instance $\llbracket (\lambda i. 0 \leq i)\ 0; \bigwedge y. (\lambda i. 0 \leq i)\ y \implies (\lambda i. 0 \leq i)\ (y+1) \rrbracket \implies (\lambda i. 0 \leq i)\ n$. By substituting the arguments in the function applications its assumption part yields the two goals

$$\begin{aligned} 0 &\leq 0 \\ \bigwedge y. 0 &\leq y \implies 0 \leq (y+1) \end{aligned}$$

which correspond to the base case and induction step as described above.

Induction rules may even be more general than shown above. Instead of applying $?P$ to a single argument it may have several arguments and the conclusion becomes $?P\ ?a_1 \dots ?a_r$. Also in the P_i every occurrence of $?P$ then has r terms as arguments. Such an induction rule is valid if it covers all possible cases for all combinations of the r argument values.

The *induction* Method

Induction can be performed in a structured proof using the method *induction* in the form

induction x *rule*: *name*

where *name* is the name of a valid induction rule. The method prepares the rule by substituting the specified variable *x* for the unknown *?a* and applies the rule to the first goal in the goal state.

Additionally, the method creates a named context for every goal resulting from the rule application. The context contains the variables and assumptions specified in the corresponding case in the induction rule. For the general form depicted above the context for the *i*th case contains the variables $y_{i1} \dots y_{ipi}$ and the assumptions $Q_{i1}; \dots; Q_{iqi}$. The term abbreviation *?case* is defined for the case conclusion *PC term_i* which is to be proven for the case.

The *induction* method treats input facts like the empty method (see Section 2.3.1) and the *cases* method (see Section 2.4.3) by inserting them as assumptions into the original goal before splitting it.

Also like the *cases* method the *induction* method supports automatic rule selection for the induction rule. This is only possible if *?P* is applied to a single argument, which means that only one variable is specified in the method:

```
induction x
```

Then the rule is selected according to the type of *x*. In Isabelle HOL (see Section 3) most types have an associated induction rule.

The rule $\llbracket ?P \text{ True}; ?P \text{ False} \rrbracket \implies ?P \text{ ?a}$ is associated with type *bool*. Therefore induction can be applied to every proposition which contains a variable of type *bool*, such as the goal $b \wedge \text{False} = \text{False}$. Applying the method

```
induction b
```

will split the goal into the goals

```
False ∧ False = False
True ∧ False = False
```

which cover all possible cases for *b*. Here, the type has only two values, therefore induction is not really needed.

Like for the *cases* method (see Section 2.4.3) the names used for the contexts created by the *induction* method can be specified by attributing the induction rule. They can be determined from the proof skeleton which is displayed in the interactive editor when the cursor is positioned on the *induction* method (see Section 2.4.1).

If the induction rule $\llbracket ?P \text{ 0}; \bigwedge y. ?P \text{ y} \implies ?P \text{ (y+1)} \rrbracket \implies ?P \text{ ?a}$ for the natural numbers has been proven and named *induct_nat* with case names *Start* and *Step*, a structured proof for the goal $0 \leq n$ may have the form

```
proof (induction n rule: induct_nat)
```

```

case Start
...
show ?case <proof>
next
case Step
...
show ?case <proof>
qed

```

The *induction* method creates the named contexts *Start* and *Step*. The former has no local variables and assumptions and binds ?*case* to the proposition $0 \leq 0$, the latter has the local variable *y*, the assumption $0 \leq y$ named *Step* and binds ?*case* to the proposition $0 \leq y + 1$.

If the rule *induct_nat* has been associated with type *nat* the rule specification may be omitted in the method:

```

proof (induction n)
...

```

Case Assumption Naming and the *induct* Method

As usual, the **case** command uses the case name as name for the assumptions $Q_{i1} \dots Q_{iq_i}$ in the *i*th case or an explicit name may be specified for them. Additionally, the *induction* method arranges the named context for a case so that the set of assumptions is split into those which in the rule contain the unknown ?*P* and those which do not. These sets are separately named, so that they can be referenced individually.

The set of assumptions which originally contained ?*P* now contain an application of *PC* to a value $term_{ij}$ and allow the step from this value to value $term_i$ by induction. These assumptions are called “induction hypothesis” and are named “*cname.IH*” where *cname* is the case name or the explicit name for the case assumptions. The other assumptions are independent from *PC*, they are additional hypotheses and are named “*cname.hyps*”. Both forms of names must be enclosed in quotes because the dot is not a normal name constituent.

For an example consider the induction rule $\llbracket ?P\ 0; ?P\ 1; \bigwedge y. \llbracket y \geq 1; ?P\ y \rrbracket \implies ?P\ (y+1) \rrbracket \implies ?P\ ?a$ with an additional base case for the value 1 and a step which always starts at value 1 or greater. If applied to the goal $0 \leq n$ the *induction* method produces the three goals

```

0 ≤ 0
0 ≤ 1
 $\bigwedge y. \llbracket y \geq 1; 0 \leq y \rrbracket \implies 0 \leq (y+1)$ 

```

If the default case name 3 is used for the third case, the induction hypothesis $0 \leq y$ is named "*3.IH*" and the additional hypothesis $y \geq 1$ is named "*3.hyps*". There is a method *induct* which behaves like *induction* with the only difference that it does not introduce the name "*cname.IH*", it uses "*cname.hyps*" for all assumptions $Q_{i1} \dots Q_{iqi}$, whether they contain $?P$ or not.

Goals with Assumptions

If the *induction* method would apply the prepared induction rule in the same way as the *rule* method to a goal $\llbracket A_1; \dots; A_m \rrbracket \implies C$ with assumptions it would unify $?P \ x$ only with the conclusion C and copy the assumptions A_1, \dots, A_m to all resulting goals unchanged. However, if x also occurs in one or more of the A_l this connection with C is lost after applying the prepared induction rule.

Consider the goal

$$4 < n \implies 5 \leq n$$

which is of the form

$$A \implies C$$

When applying the prepared induction rule for the natural numbers $\llbracket ?P \ 0; \bigwedge y. ?P \ y \implies ?P \ (y+1) \rrbracket \implies ?P \ n$ in the way of the *rule* method the conclusion will be matched which leads to the abstracted function $PC \equiv (\lambda i. 5 \leq i)$ and the resulting goals are

$$\begin{aligned} 4 < n &\implies 5 \leq 0 \\ \bigwedge y. \llbracket 4 < n; 5 \leq y \rrbracket &\implies 5 \leq (y+1) \end{aligned}$$

where the first goal is invalid. Although the second goal is valid, it shows that the relation between the variable n in the assumption and the variable y used in the induction rule has been lost.

Therefore the *induction* method works in a different way. It unifies $?P \ x$ with the conclusion C and separately with every assumption A_l and thus additionally determines an abstracted function PA_l for every A_l . Instead of using A_l directly in the resulting goals the method replaces it by $PA_l \ term_i$ in the i th goal. Moreover, if a Q_{ij} contains a sub-term $?P \ term_{ij}$ the assumptions $PA_1 \ term_{ij}; \dots; PA_m \ term_{ij}$ are added to Q_{ij} . Thus the resulting goals actually are

$$\begin{aligned} \bigwedge y_{11} \dots y_{1p1}. \llbracket Q_{11}'; \dots; Q_{1q1}'; PA_1 \ term_1; \dots; PA_m \ term_1 \rrbracket &\implies PC \ term_1 \\ \dots & \\ \bigwedge y_{n1} \dots y_{nqn}. \llbracket Q_{n1}'; \dots; Q_{nqn}'; PA_1 \ term_n; \dots; PA_m \ term_n \rrbracket &\implies PC \ term_n \end{aligned}$$

where Q_{ij}' is Q_{ij} with assumptions added as described above.

Moreover, the *induction* method (and also the *induct* method) arranges the named contexts in a way that the assumptions $PA_1 \text{ term}_n; \dots; PA_m \text{ term}_n$ which originate from the goal are named by "*cname.prem*s" and can thus be referenced separate from the Q_{ij}' which are named "*cname.hyps*" and possibly "*cname.IH*" as described above.

In the example above the *induction* method additionally unifies $?P \ n$ with the assumption $4 < n$ which yields the abstracted function $PA \equiv (\lambda i. 4 < i)$ and produces the goals

$$4 < 0 \implies 5 \leq 0$$

$$\bigwedge y. \llbracket 4 < y \implies 5 \leq y; 4 < (y+1) \rrbracket \implies 5 \leq (y+1)$$

Here $4 < (y+1)$ results from applying PA to $(y+1)$ and $4 < y$ results from adding PA applied to y as assumption to the assumption $?P \ y$ from the rule. If the default case name 2 is used for the second case, the case assumption $4 < y \implies 5 \leq y$ will be named "*2.IH*" and the case assumption $4 < (y+1)$ will be named "*2.prem*s" by the *case* command.

The *induction* method can only process assumptions A_1, \dots, A_m in the described way, if they are part of the goal to which the *induction* method is applied. If the assumptions are only present as assumed facts in the proof context, they cannot be processed by the method. In particular, if a theorem is specified in the form (see Section 2.1.7)

```
theorem
  assumes "A1" ... "An"
  shows "C"
  <proof>
```

the initial goal in *<proof>* is only C and does not contain the assumptions A_1, \dots, A_m .

In such situations, to apply the *induction* method, the required assumptions must first be inserted into the goal. This can be done by using fact input (see Section 2.2.6), since the *induction* method inserts input facts as assumptions before splitting the goal.

Thus, in the form

```
theorem
  assumes "A1" ... "An"
  shows "C"
using assms
proof (induction ...) ... qed
```

induction can be applied in the same way as if the theorem had been specified as

theorem " $\llbracket A_1; \dots; A_n \rrbracket \implies C$ "

because the assumptions A_1, \dots, A_m together are automatically named **assms** (see Section 2.2.8). Alternatively, explicit names could be specified for single assumptions to insert only some of them into the goal.

If induction is applied in a subproof, additional assumptions can also be passed to it by fact chaining (see Section 2.2.7).

Arbitrary Variables

****todo**** If some of the y_{ij} collide with some of the x_i they are consistently renamed.

Chapter 3

Isabelle HOL Basics

The basic mechanisms described in Chapter 2 can be used for working with different “object logics”. An object logic defines the types of objects available, constants of these types, and facts about them. An object logic may also extend the syntax, both inner and outer syntax.

The standard object logic for Isabelle is the “Higher Order Logic” HOL, it covers a large part of standards mathematics and is flexibly extensible. This chapter introduces some basic mechanisms defined by HOL which are used to populate HOL with many of its mathematical objects and functions and which can also be used to extend HOL to additional kinds of objects. Basically these mechanisms support the definition of new types.

HOL extends the basic type introduction mechanisms of Isabelle (see Section 2.1.2) by several ways of specifying the set of values of a new type. This section introduces four of them: algebraic types, records, subtypes, and quotient types. Additionally it introduces “let”-terms which can be used with arbitrary types.

3.1 Algebraic Types

Roughly an algebraic type is equivalent to a union of cartesian products with support for recursion. In this way most data types used in programming languages can be covered, such as records, unions, enumerations, and pointer structures. Therefore Isabelle also uses the notion “datatype” for algebraic types.

3.1.1 Definition of Algebraic Types

Basically, an algebraic type is defined in the form

datatype *name* = *alt*₁ / ... / *alt*_{*n*}

where *name* is the name of the new algebraic type and every alternative *alt_i* is a “constructor specification” of the form

cname_i "type_{i1}" ... "type_{iki}"

The *cname_i* are names and the *type_{ij}* are types. The types are specified in inner syntax and must be quoted, if they are not a single type name. All other parts belong to the outer syntax.

Recursion is supported for the types, i.e., the name *name* of the defined type may occur in the type specifications *type_{ij}*. However, there must be atleast one constructor specification which is not recursive, otherwise the definition does not “terminate”. Isabelle checks this condition and signals an error if it is not satisfied.

As a convention, capitalized names are used in Isabelle HOL for the *cname_i*. An example for a datatype definition with two constructor specifications is

```
datatype coord =
  Dim2 nat nat
| Dim3 nat nat nat
```

Its value set is equivalent to the union of pairs and triples of natural numbers.

3.1.2 Constructors

Every *cname_i* is used by the definition to introduce a “(value) constructor function”, i.e., a constant

cname_i :: "type_{i1} ⇒ ... ⇒ type_{iki} ⇒ name"

which is a function with *ki* arguments mapping their arguments to values of the new type *name*.

Every datatype definition constitutes a separate namespace for the functions it introduces. Therefore the same names may be used in constructor specifications of different datatype definitions. If used directly, a name refers to the constructor function of the nearest preceding datatype definition. To refer to constructor functions with the same name of other datatypes the name may be qualified by prefixing it with the type name in the form *name.cname_i*.

The definition of type *coord* above introduces the two constructor functions *Dim2* :: *nat* ⇒ *nat* ⇒ *coord* and *Dim3* :: *nat* ⇒ *nat* ⇒ *nat* ⇒ *coord*. Their qualified names are *coord.Dim2* and *coord.Dim3*.

Constructing Values

These constructor functions are assumed to be injective, thus their result values differ if atleast one argument value differs. This implies that the set

of all values of the constructor function $cname_i$ is equivalent to the cartesian product of the value sets of $type_{i1} \dots type_{iki}$: for every tuple of arguments there is a constructed value and vice versa. Note, however, that as usual the values of the new type are distinct from the values of all other types, in particular, they are distinct from the argument tuples.

Moreover the result values of different constructor functions are also assumed to be different. Together the set of all values of the defined type is equivalent to the (disjoint) union of the cartesian products of all constructor argument types. Moreover, every value of the type may be denoted by a term

$cname_i \ term_1 \dots term_{ki}$

where each $term_j$ is of type $type_{ij}$ and specifies an argument for the constructor function application.

Values of type *coord* as defined above are denoted by terms such as *Dim2 0 1* and *Dim3 10 5 21*.

Constant Constructors and Enumeration Types

A constructor specification may consist of a single constructor name $cname_i$, then the constructor functions has no arguments and always constructs the same single value. The constructor is equivalent to a constant of type *name*. As a consequence an “enumeration type” can be defined in the form

`datatype three = Zero | One | Two`

This type *three* has three values denoted by *Zero*, *One*, and *Two*.

Types with a Single Constructor

If a datatype definition consists of a single constructor specification its value set is equivalent to the corresponding cartesian product. The corresponding tuples have a separate component for every constructor argument type. As a consequence a “record type” can be defined in the form

`datatype recrd = MkRecrd nat "nat set" bool`

Its values are equivalent to triples where the first component is a natural number, the second component is a set of natural numbers, and the third component is a boolean value. An example value is denoted by *MkRecrd 5 {1,2,3} True*.

Since there must be atleast one nonrecursive constructor specification, definitions with a single constructor specification cannot be recursive.

3.1.3 Destructors

Since constructor functions are injective it is possible to determine for every value of the defined type the value of each constructor argument used to construct it. Corresponding mechanisms are called “destructors”, there are three different types of them.

Selectors

The most immediate form of a destructor is a selector function. For the constructor argument specified by $type_{ij}$ the selector function is a function of type $name \Rightarrow type_{ij}$. For every value constructed by $cname_i \ term_1 \dots \ term_{ki}$ it returns the value denoted by $term_j$.

The names of selector functions must be specified explicitly. This is done using the extended form of a constructor specification

$cname_i \ (sname_{i1} : "type_{i1}") \dots (sname_{iki} : "type_{iki}")$

where the $sname_{ij}$ are the names used for the corresponding selector functions. Selector names may be specified for all or only for some constructor arguments. As for constructors, selector names belong to the namespace of the defined type and may be qualified by prefixing the type name.

An example datatype definition with selectors is

datatype *recrd* = *MkRecrd* (*n:nat*) (*s:"nat set"*) (*b:bool*)

It shows that the selector functions correspond to the field names used in programming languages in record types to access the components. For every term r of type *recrd* the selector term $s \ r$ denotes the set component of r .

An example for a datatype with multiple constructor specifications is

datatype *coord* =
 Dim2 (*x:nat*) (*y:nat*)
 / *Dim3* (*x:nat*) (*y:nat*) (*z:nat*)

Note that the selectors x and y are specified in both alternatives. Therefore a single selector function $x :: coord \Rightarrow nat$ is defined which yields the first component both for a two-dimensional and a three-dimensional coordinate and analogously for y . If instead the definition is specified as

datatype *coord* =
 Dim2 (*x2:nat*) (*y:nat*)
 / *Dim3* (*x3:nat*) (*y:nat*) (*z:nat*)

two separate selector functions $x2$ and $x3$ are defined where the first one is only applicable to two-dimensional coordinates and the second one only to three-dimensional coordinates.

If a selector name does not occur in all constructor specifications, the selector function is still total, like all functions in Isabelle, but it is underspecified (see Section 2.1.3). It maps values constructed by other constructors to a unique value of its result type, even if that other constructor has no argument of this type. However, no information is available about that value.

For the type *coord* the selector function $z :: \text{coord} \Rightarrow \text{nat}$ is also applicable to two-dimensional coordinates, however, the values it returns for them is not specified.

Such selector values are called “default selector values”. They may be specified in the extended form of a datatype definition

```
datatype name = alt1 / ... / altn
where "prop1" / ... / "propm"
```

where every *prop_p* is a proposition of the form

```
snameij (cnameq var1 ... varkq) = termp
```

and specifies *term_p* as the default value of selector *sname_{ij}* for values constructed by *cname_q*.

The definition

```
datatype coord =
  Dim2 (x:nat) (y:nat)
  / Dim3 (x:nat) (y:nat) (z:nat)
where "z (Dim2 a b) = 0"
```

specifies 0 as default value for selector *z* if applied to a two-dimensional coordinate.

Discriminators

If an underspecified selector is applied to a datatype value it may be useful to determine which constructor has been used to construct the value. This is supported by discriminator functions. For every constructor specification for *cname_i* the discriminator function has type *name* \Rightarrow *bool* and returns true for all values constructed by *cname_i*. Like selector names, discriminator names must be explicitly specified using the extended form of a datatype definition

```
datatype name = dname1: alt1 / ... / dnamen: altn
```

Discriminator names may be specified for all alternatives or only for some of them. Note that for a datatype with a single constructor the discriminator returns always *True* and for a datatype with two constructors one discriminator is the negation of the other.

An example datatype definition with discriminators is

```

datatype coord =
  is_2dim: Dim2 nat nat
| is_3dim: Dim3 nat nat nat

```

In a datatype definition both discriminators and selectors may be specified.

The *case* Term

Additionally to using discriminators and selectors Isabelle HOL supports *case* terms. A *case* term specifies depending on a datatype value a separate term variant for every constructor of the datatype. In these variants the constructor arguments are available as bound variables.

A *case* term for a datatype *name* defined as in Section 3.1.1 has the form

```

case term of
  cname1 var11 ... var1k1 ⇒ term1
| ...
| cnamen varn1 ... varnkn ⇒ termn

```

where *term* is of type *name* and the *term_i* have an arbitrary but common type which is also the type of the *case* term. In the alternative for constructor *cname_i* the *var₁₁* ... *var_{1k₁}* must be distinct variables, they are bound to the constructor arguments and may be used in *term_i* to access them. The value of *var_{ij}* is the same as the value selected by *sname_{ij}* *term*.

If *cv* is a variable or constant of type *coord* an example *case* term for it is

```

case cv of
  Dim2 a b ⇒ a + b
| Dim3 a b c ⇒ a + b + c

```

It denotes the sum of the coordinates of *cv*, irrespective whether *cv* is two-dimensional or three-dimensional.

A *case* term is useful even for a datatype with a single constructor. If *rv* is of type *recrd* as defined in Section 3.1.3 the *case* term

```

case rv of MkRecrd nv sv bv ⇒ term

```

makes the components of *rv* locally available in *term* as *nv*, *sv*, *bv*. It is equivalent to *term* where *nv*, *sv*, and *bv* have been substituted by the selector applications (*n rv*), (*s rv*), and (*b rv*).

The variant terms in a *case* term cannot be matched directly by a **let** statement in a proof (see Section 2.2.11). The statement

```

let "case rv of MkRecrd nv sv bv ⇒ ?t"
  = "case rv of MkRecrd nv sv bv ⇒ term"

```

will fail to bind $?t$ to $term$ because then the variables nv , sv , and bv would occur free in it and the relation to the constructor arguments would be lost. Instead, the statement

```
let "case rv of MkRecrd nv sv bv  $\Rightarrow$  ?t nv sv bv"
    = "case rv of MkRecrd nv sv bv  $\Rightarrow$  term"
```

successfully binds $?t$ to the lambda term $\lambda nv\ sv\ bv. term$ which denotes the function which results in $term$ when applied to the constructor arguments.

3.1.4 Rules

A datatype definition also introduces a large number of named facts about the constructors and destructors of the defined type. All fact names belong to the namespace of the datatype definition. Since the fact names cannot be specified explicitly, all datatype definitions use the same fact names, therefore the fact names must always be qualified by prefixing the type name.

Several rules are configured for automatic application, e.g., they are added to the simpset for automatic application by the simplifier (see Section 2.3.4). Other rules must be explicitly used by referring them by their name.

Only some basic rules are described here, for more information refer to the Isabelle documentation about datatypes.

Simplifier Rules

The rules added by a datatype definition to the simpset (see Section 2.3.4) support many ways for the simplifier to process terms with constructors and destructors. An example are rules of the form

$$sname_{ij} (cname_i\ x_1 \dots x_{ki}) = x_j$$

The set of all rules added to the simpset is named $name.simps$. By displaying it using the **thm** command (see Section 2.1.7) it can be inspected to get an idea how the simplifier processes terms for datatypes.

Case Rule

Every datatype definition introduces a case rule (see Section 2.4.3) of the form

```
lemma name.exhaust:
  "[[ $\bigwedge x_1 \dots x_{k1}. y = cname_1\ x_1 \dots x_{k1} \Rightarrow P;$ 
   ... ;
    $\bigwedge x_1 \dots x_{kn}. y = cname_n\ x_1 \dots x_{kn} \Rightarrow P]$ ]  $\Rightarrow P$ "
```

It is valid because the constructor applications cover all possibilities of constructing a value y of the datatype.

This rule is associated with the datatype for use by the `cases` method (see Section 2.4.3). Therefore the application of the method

```
cases "term"
```

where `term` is of type `name` splits an arbitrary goal into n subgoals where every subgoal uses a different constructor to construct the `term`.

The names for the named contexts created by the `cases` method are simply the constructor names `cnamei`. Therefore a structured proof using case based reasoning for a `term` of datatype `name` has the form

```
proof (cases "term")
  case (cname1 x1 ... xk1) ... show ?thesis <proof>
next
...
next
  case (cnamen x1 ... xkn) ... show ?thesis <proof>
qed
```

The names x_i of the locally fixed variables can be freely selected, they denote the constructor arguments of the corresponding constructor. Therefore the case specification `(cnamei x1 ... xki)` looks like a constructor application to variable arguments, although it is actually a context name together with locally fixed variables.

Split Rule

```
**todo**
```

Induction Rule

```
**todo**
```

3.2 Record Types

```
**todo**
```

3.3 Subtypes

A subtype specifies the values of a type simply as a set of values of an existing type. However, since the values of different types are always disjoint,

the values in the set are not directly the values of the new type, instead, there is a 1-1 relation between them, they are isomorphic. The values in the set are called “representations”, the values in the new type are called “abstractions”.

3.3.1 Subtype Definitions

A subtype is defined in the form

```
typedef name = "term" <proof>
```

where *name* is the name of the new type and *term* is a term for the representing set. See Section 4.5 for how to denote such terms. The <*proof*> must prove that the representing set is not empty.

A simple example is the type

```
typedef three = "{1::nat,2,3}" by auto
```

which has three values. The representations are natural numbers, as usual, the type *nat* must be specified because the constants 1, 2, 3 may also denote values of other types. However, they do not denote the values of the new type *three*, the type definition does not introduce constants for them.

Instead, a subtype definition **typedef** *t* = "*term*" <*proof*> introduces two functions *Abs_t* and *Rep_t*. These are morphisms between the set and the new type, *Abs_t* maps from the set to type *t*, *Rep_t* is its inverse. Both functions are injective, together they provide the 1-1 mapping between the subtype and the representing set. The function *Abs_t* can be used to denote the values of the subtype.

In the example the morphisms are *Abs_three* :: *nat* \Rightarrow *three* and *Rep_three* :: *three* \Rightarrow *nat*. The values of type *three* may be denoted as (*Abs_three* 1), (*Abs_three* 2), and (*Abs_three* 3).

Alternative names may be specified for the morphisms in the form

```
typedef t = "term" morphisms rname aname <proof>
```

where *rname* replaces *Rep_t* and *aname* replaces *Abs_t*.

Like declared types subtypes may be parameterized (see Section 2.1.2):

```
typedef ('name1, ..., 'namen') name = "term" <proof>
```

where the '*name*_{*i*}' are the type parameters. They may occur in the type of the *term*, i.e., the *term* may be polymorphic (see Section 2.1.3).

3.3.2 Subtype Rules

****todo****

3.4 Quotient Types

****todo****

3.5 Type Independent Mechanisms

****todo****

3.5.1 Undefined Value

****todo****

3.5.2 Let Terms

****todo****

Chapter 4

Isabelle HOL Types

This chapter introduces a small basic part of the types available in HOL. The selected types are considered useful for some forms of program verification. Most of the types are algebraic types (see Section 3.1). Although some of them are defined differently for technical reasons, they are configured afterwards to behave as if they have been defined as algebraic types. Therefore they are described here using the corresponding datatype definition.

4.1 Boolean Values

todo

- $\wedge, \vee, \neg, \longrightarrow$
- $=, \neq, \longleftrightarrow$
- \forall, \exists
- intro

4.1.1 Conditional Terms

todo

4.1.2 Logic Rules

todo

- conjI, conjE, disjI1, disjI2, disjE, implI, mp

- `contrapos_*`
- `iffI, iffE, iffD1, iffD2`
- `allI, allE, exI, exE`

4.2 Natural Numbers

****todo****

4.3 Tuple Types

Tuple types are constructed by cartesian product of existing types. Tuple types can be directly denoted by type expressions of the form

$t_1 \times \dots \times t_n$

in inner syntax. It is not necessary to introduce a name to use a tuple type, however, this is possible by defining a type synonym (see Section 2.1.2):

type_synonym $t = t_1 \times \dots \times t_n$

4.3.1 Tuple Values

****todo****

4.4 Optional Values

****todo****

4.5 Sets

****todo****

4.6 Lists

****todo****

4.7 Fixed-Size Binary Numbers

****todo****