

A Gentle Introduction to Isabelle and Isabelle HOL

Gunnar Teege

January 9, 2023

Contents

1	Isabelle Basics	3
1.1	Invoking Isabelle	3
1.1.1	Installation and Configuration	3
1.1.2	Theories and Sessions	4
1.1.3	Invocation as Editor	4
1.1.4	Invocation for Batch Processing	5
1.1.5	Invocation for Document Creation	6
1.2	Interactively Working with Isabelle	6
1.2.1	The Text Area	6
1.2.2	The Output Panel	6
1.2.3	The State Panel	7
1.2.4	The Symbols Panel	7
1.2.5	The Documentation Panel	7
1.2.6	The Query Panel	7
1.2.7	The Theories Panel	7
1.3	Isabelle Theories	7
1.3.1	Theory Structure	8
1.3.2	Types	9
1.3.3	Terms	9
1.3.4	Definitions	11
1.3.5	Overloading	12
1.3.6	Propositions	13
1.3.7	Theorems	15
1.3.8	Locales	18
1.4	Isabelle Proofs	20
1.4.1	Proof Context	20
1.4.2	Proof Procedure	21
1.4.3	Basic Proof Structure	23
1.4.4	Backward Reasoning Steps	26
1.4.5	Forward Reasoning Steps	27
1.4.6	Facts as Proof Input	28
1.4.7	Fact Chaining	30
1.4.8	Assuming Facts	33

1.4.9	Fixing Variables	36
1.4.10	Term Abbreviations	37
1.4.11	Accumulating Facts	39
1.4.12	Equational Reasoning	41
1.5	Proof Methods	43
1.5.1	The empty Method	43
1.5.2	Rule Application	43
1.5.3	Composed Proof Methods	50
1.5.4	The Simplifier	51
1.5.5	Other Automatic Proof Methods	57
2	Isabelle HOL	59
2.1	Type Definitions	59
2.2	Terms	59
2.3	Types	59
2.3.1	Boolean Values	59
2.3.2	Sets	60
2.3.3	Natural Numbers	60
2.3.4	Tuples	60
2.3.5	Lists	60
2.3.6	Fixed-Size Binary Numbers	60

Chapter 1

Isabelle Basics

Isabelle is a “proof assistant” for formal mathematical proofs. It supports a notation for propositions and their proofs, it can check whether a proof is correct, and it can even help to find a proof.

This introductory manual explains how to work with Isabelle to develop mathematical models. It does not presume prior knowledge about formal or informal proof techniques. It only assumes that the reader has a basic understanding of mathematical logics and the task of proving mathematical propositions.

1.1 Invoking Isabelle

After installation, Isabelle can be invoked interactively as an editor for entering propositions and proofs, or it can be invoked noninteractively to check a proof and generate a PDF document which displays the propositions and proofs.

1.1.1 Installation and Configuration

Isabelle is freely available from <https://isabelle.in.tum.de/> for Windows, Mac, and Linux. It is actively maintained, there is usually one release every year. Older releases are available in a distribution archive.

To install Isabelle, follow the instructions on

`https://isabelle.in.tum.de/installation.html`

Although there are many configuration options, there is no need for an initial configuration, interactive and noninteractive invocation is immediately possible.

1.1.2 Theories and Sessions

The propositions and proofs in Isabelle notation are usually collected in “theory files” with names of the form `name.thy`. A theory file must import at least one other theory file to build upon its content. For theories based on higher order logic (“HOL”), the usual starting point to import is the theory *Main*.

Several theory files can be grouped in a “session”. A session is usually stored in a directory in the file system. It consists of a file named `ROOT` which contains a specification of the session, and the theory files which belong to the session.

When Isabelle loads a session it loads and checks all its theory files. Then it can generate a “heap file” for the session which contains the processed session content. The heap file can be reloaded by Isabelle to avoid the time and effort for processing and checking the theory files.

A session always has a single parent session, with the exception of the Isabelle builtin session `Pure`. Thus, every session depends on a linear sequence of ancestor sessions which begins at `Pure`. The ancestor sessions have separate heap files. A session is always loaded together with all ancestor sessions.

Every session has a name of the form `chap/sess` where `chap` is an arbitrary “chapter name”, it defaults to `Unsorted`. The session name and the name of the parent session are specified in the `ROOT` file in the session directory. When a session is loaded by Isabelle, its directory and the directories of all ancestor sessions must be known by Isabelle.

The Isabelle distribution provides heap files for the session `HOL/HOL` and its parent session `Pure/Pure`, the session directories are automatically known. Every session may be displayed in a “session document”. This is a PDF document generated by translating the content of the session theory files to \LaTeX . A frame \LaTeX document must be provided which includes all content generated from the theory files. The path of the frame document, whether a session document shall be generated and which theories shall be included is specified in the `ROOT` file.

The command

```
isabelle mkroot [OPTIONS] [Directory]
```

can be used to initialize the given directory (default is the current directory) as session directory. It creates an initial `ROOT` file to be populated with theory file names and other specification for the session, and it creates a simple frame \LaTeX document.

1.1.3 Invocation as Editor

Isabelle is invoked for editing using the command

```
isabelle jedit [OPTIONS] [Files ...]
```

It starts an interactive editor and opens the specified theory files. If no file is specified it opens the file `Scratch.thy` in the user's home directory. If that file does not exist, it is created as an empty file.

The editor also loads a session (together with its ancestors), the default session to load is HOL. If a heap file exists for the loaded session it is used, otherwise a heap file is created by processing all the session's theories.

The default session to load can be changed by the option

```
-l <session name>
```

Moreover the editor also loads (but does not open) theories which are transitively imported by the opened theory files. If these are Isabelle standard theories it finds them automatically. If they belong to the session in the current directory it also finds them. If they belong to other sessions, the option

```
-d <directory pathname>
```

must be used to make the session directory known to Isabelle. For every used session a separate option must be specified.

If an imported theory belongs to the loaded session or an ancestor, it is directly referenced there. Otherwise the theory file is loaded and processed.

1.1.4 Invocation for Batch Processing

Isabelle is invoked for batch processing of all theory files in one or more sessions using the command

```
isabelle build [OPTIONS] [Sessions ...]
```

It loads all theory files of the specified sessions and checks the contained proofs. It also loads all required ancestor sessions. If not known to Isabelle, the corresponding session directories must be specified using option `-d` as described in Section 1.1.3. Sessions required for other sessions are loaded from heap files if existent, otherwise the corresponding theories are loaded and a heap file is created.

If option `-b` is specified, heap files are also created for all sessions specified in the command. Option `-c` clears the specified sessions (removes their heap files) before processing them. Option `-n` omits the actual session processing, together with option `-c` it can be used to simply clear the heap files.

The specified sessions are only processed if at least one of their theory files has changed since the last processing or if the session is cleared using option

-c. If option -v is specified all loaded sessions and all processed theories are listed on standard output.

If specified for a session in its ROOT file (see Section 1.1.5), also the session document is generated when a session is processed.

1.1.5 Invocation for Document Creation

**** todo ****

1.2 Interactively Working with Isabelle

After invoking Isabelle as editor (see Section 1.1.3) it supports interactive work with theories.

The user interface consists of a text area which is surrounded by docking areas where additional panels can be displayed. Several panels can be displayed in the same docking area, using tabs to switch among them. Panels may also be displayed as separate undocked windows.

A panel can be displayed by selecting it in the **Plugins -> Isabelle** menu. Some of the panels are described in the following sections.

1.2.1 The Text Area

The text area displays the content of an open theory file and supports editing it. The font (size) used for display can be configured through the menu in **Utilities -> Global Options -> jEdit -> Text Area** together with many other options for display.

Moreover, in the default configuration, Isabelle automatically processes the theory text up to the current cursor position. This includes processing the content of all imported theory files, if the cursor is after the import statement.

Whenever the cursor is moved, the processing is continued or set back, if “Continuous checking” has not been disabled in the Theories panel.

In the default configuration the progress of the processing is shown by shading the unprocessed text in red and by a bar on the right border of the text area which symbolizes the whole theory file and shows the unprocessed part by red shading as well.

1.2.2 The Output Panel

The Output panel displays the result of the theory text processing when it reaches the cursor position in the text area.

The displayed information depends on the cursor position and may be an information about the current theorem or proof or it may be an error message.

1.2.3 The State Panel

The State panel displays a specific result of the theory text processing if the cursor position is in a proof. It is called the “goal state” (see Section 1.4.2), and describes what remains to be proven by the rest of the proof.

The Output panel can be configured to include the goal state in its display by checking the “Proof state” button.

1.2.4 The Symbols Panel

Isabelle uses a large set of mathematical symbols and other special symbols which are usually not on the keyboard. The Symbols panel can be used to input such symbols in the text area.

1.2.5 The Documentation Panel

A comprehensive set of documentation documents about Isabelle can be opened through the Documentation Panel. This manual refers to some of these documentations, if applicable.

For example, more information about the use of the interactive editor can be found in the Isabelle Reference Manual about `jedit`.

1.2.6 The Query Panel

`** todo**`

1.2.7 The Theories Panel

The Theories panel displays the loaded session and the opened or imported theories which do not belong to the loaded session or its ancestors. The (parts of) theories which have not been processed are shaded in red.

If the check button next to a theory is checked, the theory file is processed independently of the cursor position in the text area.

1.3 Isabelle Theories

Remark: This document uses a pretty printed syntax for Isabelle theories. A major difference to the actual source syntax is that underscores in names are

displayed as hyphens (minus signs). So for example the keyword displayed as **type-synonym** must be written `type_synonym` in the theory source text.

1.3.1 Theory Structure

The content of a theory file has the structure

```
theory name
imports name1 ... namen
begin
  ...
end
```

where *name* is the theory name and *name*₁ ... *name*_{*n*} are the names of the imported theories. The theory name *name* must be the same which is used for the theory file, i.e., the file name must be **name.thy**.

The theory structure is a part of the Isabelle “outer syntax” which is mainly fixed and independent from the specific theories. Other kind of syntax is embedded into the outer syntax. The main embedded syntax is the “inner syntax” which is mainly used to denote types and terms. Content in inner syntax must always be surrounded by double quotes. Note that in the pretty printed forms in this document these quotes are omitted.

This introduction describes only a selected part of the outer syntax. The full outer syntax is described in the Isabelle/Isar Reference Manual.

Additionally, text written in L^AT_EX syntax can be embedded into the outer syntax using the form **text**(...) and L^AT_EX sections can be created using **chapter**(...), **section**(...), **subsection**(...), **subsubsection**(...), **paragraph**(...), **subparagraph**(...). Note that the delimiters used here are not the “lower” and “greater” symbols, but the “cartouche delimiters” available in the editor’s Symbols panel in tab “Punctuation”.

This text is intended for additional explanations of the formal theory content. It is displayed in the session document together with the formal theory content.

It is also possible to embed inner and outer syntax in the L^AT_EX syntax (see Chapter 4 in the Isabelle/Isar Reference Manual).

Moreover, comments of the form

```
(* ... *)
```

can be embedded into the outer syntax. They are only intended for the reader of the theory file and are not displayed in the session document.

1.3.2 Types

As usual in formal logics, the basic building blocks of propositions are terms. Terms denote arbitrary objects like numbers, sets, functions, or boolean values. Isabelle is strongly typed, so every term must have a type. However, in most situations Isabelle can derive the type of a term automatically, so that it needs not be specified explicitly. Terms and types are always denoted using the inner syntax.

Types are usually specified by type names. There are predefined type names such as *nat* and *bool* for natural numbers and boolean values. New type names can be declared in the form

typedecl *name*

which introduces the *name* for a new type for which the values are different from the values of all existing types (and no other information about the values is given).

Types can be parameterized, then the type arguments are denoted *before* the type name, such as in *nat set* which is the type of sets of natural numbers. A type name with *n* parameters is declared in the form

typedecl (*'name₁, ..., 'name_n*) *name*

such as **typedecl** (*'a*) *set*. The type parameters are denoted by “type variables” which always have the form *'name* with a leading single quote character. Every use where the parameters are replaced by actual types, such as in *nat set*, is called an “instance” of the parameterized type.

Alternatively a type name can be introduced as a synonym for an existing type in the form

type-synonym *name* = *type*

such as in **type-synonym** *natset* = *nat set*. Type synonyms can also be parameterized as in

type-synonym (*'name₁, ..., 'name_n*) *name* = *type*

where the type variables occur in *type* in place of actual type specifications.

1.3.3 Terms

Constants and Variables

Terms are mainly built as syntactical structures based on constants and variables. Constants are usually denoted by names, using the same namespace as type names. Whether a name denotes a constant or a type depends on

its position in a term. Predefined constant names of type *bool* are *True* and *False*.

Constants of number types, such as *nat*, may be denoted by number literals, such as *6* or *42*.

A constant name can be introduced by declaring it together with its type. The declaration

consts *name*₁ :: *type*₁ ... *name*_{*n*} :: *type*_{*n*}

declares *n* constant names with their types. Constant names declared in this way are “atomic”, no other information is given about them.

If the constant’s type contains type variables the constant is called “polymorphic”. Thus the declaration

consts *myset* :: 'a *set*

declares the polymorphic constant *myset* which may be a set of elements of arbitrary type.

A (term) variable has the same form as a constant name, but it has not been introduced as a constant. Whenever a variable is used in a term it has a specific type which is either derived from its context or is explicitly specified in the form *varname* :: *type*.

Functions

A constant name denotes an object, which, according to its type, may also be a function of arbitrary order. Functions basically have a single argument. The type of a function is written as *argtype* \Rightarrow *restype*. The result type of a function may again be a function type, then it may be applied to another argument. This is used to represent functions with more than one argument. Function types are right associative, thus a type *argtype*₁ \Rightarrow *argtype*₂ \Rightarrow ... \Rightarrow *argtype*_{*n*} \Rightarrow *restype* describes a function which can be applied to *n* arguments.

Function application terms for a function *f* and an argument *a* are denoted by *f a*, no parentheses are required around the argument. Function application terms are left associative, thus a function application to *n* arguments is written *f a*₁ ... *a*_{*n*}. Note that an application *f a*₁ ... *a*_{*m*} where *m* < *n* (a “partial application”) is a correct term and denotes a function taking the remaining *n* − *m* arguments.

A constant name for a binary function can also be a symbol such as +, **, =, ≠, ≤, or ∈ which is applied in infix notation, i.e., an application term is denoted in the form *a*₁ + *a*₂. For all infix symbols *s* the form *op s* can be used like a normal function name in prefix notation, so an application of +

can also be denoted by $(op +) a_1 a_2$ and $(op +) 5$ is a partial application and denotes the function which increments its argument by 5.

Functions can be denoted by lambda terms of the form $\lambda x. term$ where x is a variable which may occur in the $term$. A function to be applied to n arguments can be denoted by the lambda term $\lambda x_1 \dots x_n. term$ where $x_1 \dots x_n$ are distinct variables. As usual, types may be specified for (some of) the variables in the form $\lambda(x_1::t_1) \dots (x_n::t_n). term$.

Nested terms are generally written by using parentheses (...). There are many priority rules how to nest terms automatically, but if in doubt, it is always safe to use parentheses.

1.3.4 Definitions

Constant names can also be introduced as “synonyms” for terms. There are two forms for introducing constant names in this way, definitions and abbreviations.

A definition introduces the name as a new entity in the logic, in the same way as a declaration. A definition is denoted in the form

definition $name :: type$
where $name \equiv term$

Note that the “defining equation” $name \equiv term$ is specified in inner syntax and, like $type$ must be delimited by quotes in the source text.

If the type of the defined name is a function type, the $term$ may be a lambda term. Alternatively, the definition for a function applicable to n arguments can be written in the form

definition $name :: type$
where $name x_1 \dots x_n \equiv term$

with variable names $x_1 \dots x_n$ which may occur in the $term$. This form is mainly equivalent to

definition $name :: type$
where $name \equiv \lambda x_1 \dots x_n. term$

A short form of a definition is

definition $name \equiv term$

Here, the type of the new constant is derived as the type of the $term$.

An abbreviation definition introduces the name only as a syntactical item which is not known by the internal logic system, upon input it is automatically expanded, and upon output it is used whenever a term matches its specification and the term is not too complex. An abbreviation definition is denoted in a similar form as a definition:

abbreviation *name* :: *type*
where *name* \equiv *term*

The alternative form for functions and the short form are also available like for definitions.

In both cases the defined constant may not occur in the *term*, i.e., these forms of definition do not support recursion.

1.3.5 Overloading

A declared constant name can be associated with a definition afterwards by overloading. Overloading depends on the type. Therefore, if a constant name is polymorphic, different definitions can be associated for different type instances.

A declared constant name *name* is associated with *n* definitions by the following overloading specification:

```
overloading
  name1  $\equiv$  name
  ...
  namen  $\equiv$  name
begin
  definition name1 :: type1 where ...
  ...
  definition namen :: typen where ...
end
```

where all *type*_{*i*} must be instances of the type declared for *name*.

There is also a form of overloading which is only performed on the syntactic level, like abbreviations. To use it, the theory *HOL-Library.Adhoc-Overloading* must be imported by the surrounding theory:

```
imports HOL-Library.Adhoc-Overloading
```

Then constant *name* can be associated with *n* terms of different type instances by

```
adhoc-overloading name term1 ... termn
```

The *term*_{*i*} need not be single constants. If an arbitrary term is specified, the constant *name* becomes an abbreviation for that term. A possible use case is to specify a partial application as term such as in

```
adhoc-overloading incby5 ((op +) 5)
```

Several names can be overloaded in a common specification:

```
adhoc-overloading name1 term11 ... term1n and ... and namek ...
```

1.3.6 Propositions

A proposition denotes a statement, which can be valid or not. Valid statements are called “facts”, they are the main content of a theory. Propositions are specific terms and are hence written in inner syntax and must be enclosed in quotes.

Formulas

In its simplest form a proposition is a single term of type *bool*, such as

$$6 * 7 = 42$$

Terms of type *bool* are also called “formulas”.

A proposition may contain free variables as in

$$2 * x = x + x$$

A formula as proposition is valid if it evaluates to *True* for all possible values substituted for the free variables.

Derivation Rules

More complex propositions can express, “derivation rules” used to derive propositions from other propositions. Derivation rules are denoted using a “metallogic language”. It is still written in inner syntax but uses a small set of “metallogic operators”.

Derivation rules consist of assumptions and a conclusion. They can be written using the metallogic operator \implies in the form

$$A_1 \implies \dots \implies A_n \implies C$$

where the $A_1 \dots A_n$ are the assumptions and C is the conclusion. The conclusion must be a formula. The assumptions may be arbitrary propositions. If an assumption contains metallogic operators parentheses can be used to delimit them from the rest of the derivation rule.

A derivation rule states that if the assumptions are valid, the conclusion can be derived as also being valid. So it can be viewed as a “meta implication” with a similar meaning as a boolean implication, but with a different use.

An example for a rule with a single assumption is

$$(x::nat) < c \implies n*x \leq n*c$$

Note that type *nat* is explicitly specified for variable x . This is necessary, because the constants $<$, $*$, and \leq are overloaded and can be applied to

other types than only natural numbers. Therefore the type of x cannot be derived automatically. However, when the type of x is known, the types of c and n can be derived to also be *nat*.

An example for a rule with two assumptions is

$$(x::nat) < c \implies n > 0 \implies n*x < n*c$$

In most cases the assumptions are also formulae, as in the example. However, they may also be again derivation rules. Then the rule is a “meta rule” which derives a proposition from other rules. This introductory manual usually does not take such meta rules into account.

Binding Free Variables

A proposition may contain universally bound variables, using the metalogic quantifier \bigwedge in the form

$$\bigwedge x_1 \dots x_n. P$$

where the $x_1 \dots x_n$ may occur free in the proposition P . As usual, types may be specified for (some of) the variables in the form $\bigwedge (x_1::t_1) \dots (x_n::t_n). P$. An example for a valid derivation rule with bound variables is

$$\bigwedge (x::nat) c \ n. \ x < c \implies n*x \leq n*c$$

If a standalone proposition contains free variables they are implicitly universally bound. Thus the example derivation rule above is equivalent to the single-assumption example rule in the previous section. Explicit binding of variables is only required to avoid name clashes with constants of the same name. In the proposition

$$\bigwedge (True::nat). \ True < c \implies n*True \leq n*c$$

the name *True* is used locally as a variable of type *nat* instead of the pre-defined constant of type *bool*. Of course, using well known constant names as variables is confusing and should be avoided.

Alternative Rule Syntax

An alternative, Isabelle specific syntax for derivation rules is

$$\bigwedge x_1 \dots x_n. [A_1; \dots; A_n] \implies C$$

which is often considered as more readable, because it better separates the assumptions from the conclusion. In the interactive editor it may be necessary to switch to this form by setting **Print Mode** to **brackets** in **Plugin**

Options for Isabelle General. The fat brackets are available for input in the editor's Symbols panel in tab "Punctuation".

Using this syntax the two-assumption example rule from the previous section is denoted by

$$\bigwedge (x::nat) \ c \ n. \llbracket x < c; \ n > 0 \rrbracket \Longrightarrow n*x < n*c$$

or equivalently without quantifier by

$$\llbracket (x::nat) < c; \ n > 0 \rrbracket \Longrightarrow n*x < n*c$$

Note that in the literature a derivation rule $\llbracket P; \ Q \rrbracket \Longrightarrow P \wedge Q$ is often denoted in the form

$$\frac{P \quad Q}{P \wedge Q}$$

Another alternative, Isabelle specific syntax for a derivation rule $\bigwedge x_1 \dots x_n. \llbracket A_1; \dots; A_n \rrbracket \Longrightarrow C$ is the "structured" proposition

$$C \text{ if } A_1 \dots A_n \text{ for } x_1 \dots x_n$$

Again, types can be specified for some variables in the usual form. The assumptions and the variables may be grouped or separated for better readability by the keyword **and**. Note that the keywords **if**, **and**, **for** belong to the outer syntax. Thus, the original rule must be quoted as a whole, whereas in the structured proposition only the sub-propositions C , A_1 , \dots , A_n must be individually quoted. The x_1, \dots, x_n need not be quoted, but if a type is specified for a variable the type must be quoted, if it is not a single type name.

If written in this form, the two-assumption example rule from the previous subsections becomes

$$n*x < n*c \text{ if } x < c \text{ and } n > 0 \text{ for } x::nat \ n \ c$$

1.3.7 Theorems

A theorem specifies a proposition together with a proof, that the proposition is valid. Thus it adds a fact to the enclosing theory. A simple form of a theorem is

theorem *prop* $\langle proof \rangle$

where *prop* is a proposition in inner syntax and $\langle proof \rangle$ is a proof as described in Section 1.4. The keyword **theorem** can be replaced by one of the keywords **lemma**, **corollary**, **proposition** to give a hint about the use of the statement to the reader.

Unknowns

Whenever a theorem turns a proposition to a fact, the free (or universally bound) variables are replaced by “unknowns”. For a variable *name* the corresponding unknown is *?name*. This is only a technical difference, it signals to Isabelle that the unknowns can be consistently substituted by arbitrary terms, as long as the types are preserved.

When turned to a fact, the example rule from the previous sections becomes

$$?x < ?c \implies ?n * ?x \leq ?n * ?c$$

with type *nat* associated to all unknowns.

The result of such a substitution is always a special case of the fact and therefore also a fact. In this way a fact with unknowns gives rise to a (usually infinite) number of facts which are constructed by substituting unknowns by terms.

Isabelle can be configured to suppress the question mark when displaying unknowns, then this technical difference becomes invisible.

Named Facts

Facts are often used in proofs of other facts. For this purpose they can be named so that they can be referenced by name. A named fact is specified by a theorem of the form

theorem *name*: *prop* *<proof>*

The example rule from the previous sections can be turned into a fact named *example1* by

theorem *example1*: $(x::nat) < c \implies n * x \leq n * c$ *<proof>*

It is also possible to introduce named collections of facts. A simple way to introduce such a named collection is

lemmas *name* = *name*₁ ... *name*_{*n*}

where *name*₁ ... *name*_{*n*} are names of existing facts or fact collections.

If there is a second rule stated as a named fact by

theorem *example2*: $(x::nat) \leq c \implies x + m \leq c + m$ *<proof>*

a named collection can be introduced by

lemmas *examples* = *example1 example2*

Alternatively a “dynamic fact set” can be declared by

named-theorems *name*

It can be used as a “bucket” where facts can be added afterwards by specifying the bucket name in the theorem:

theorem [*name*]: *prop* $\langle proof \rangle$

or together with specifying a fixed fact name *name_f* by

theorem *name_f*[*name*]: *prop* $\langle proof \rangle$

A named fact or fact set (but not a dynamic fact set) can be displayed using the command

thm *name*

which may be entered outside of theorems and at most positions in a proof. The facts are displayed in the Output panel (see Section 1.2.2).

Alternative Theorem Syntax

There is an alternative syntax for theorems which have a derivation rule as their proposition. A theorem **theorem** $\wedge x_1 \dots x_n. \llbracket A_1; \dots; A_n \rrbracket \implies C$ $\langle proof \rangle$ can also be specified in the form

theorem
 fixes $x_1 \dots x_n$
 assumes $A_1 \dots A_n$
 shows C
 $\langle proof \rangle$

Similar to the structured proposition form, the variables and assumptions may be grouped by **and**, the keywords belong to the outer syntax and the C , A_1 , \dots , A_n must be individually quoted.

Using this syntax the two-assumption example rule from the previous sections can be written as

theorem
 fixes $x::nat$ **and** c **and** n
 assumes $x < c$ **and** $n > 0$
 shows $n*x < n*c$
 $\langle proof \rangle$

In contrast to the general theorem syntax this alternative syntax allows to specify names for some or all of the assumption groups as

assumes $name_1: A_{11} \dots A_{1m_1}$ **and** \dots **and** $name_n: A_{n1} \dots A_{nm_n}$

These names can (only) be used in the proof of the theorem. More consequences this syntax has for the proof are described in Section 1.4.

Note that a name specified for the conclusion as

shows *name*: *C*

becomes the name for the whole fact introduced by the theorem, not only for the conclusion. It is not available in the proof of the theorem.

Definitions as Facts

The definitions described in Section 1.3.4 also introduce facts in the enclosing theory. Every definition introduces a new constant and specifies a defining equation of the form *name* \equiv *term* for it. This equation is a proposition using the “meta equality” \equiv which is another metalogic operator. It is the initial information given for the new constant, thus it is valid “by definition” and is a fact in the theory.

These facts are automatically named. If *name* is the name of the defined constant, the defining equation is named *name-def*. Alternatively an explicit name can be specified in the form

definition *name* :: *type*
where *fact-name*: *name* \equiv *term*

1.3.8 Locales

There are cases where theory content such as definitions and theorems occur which has similar structure but differs in some types or terms. Then it is useful to define a “template” and instantiate it several times. This can be done in Isabelle using a “locale”.

Simple Locales

A locale can be seen as a parameterized theory fragment, where the parameters are terms. A locale with *n* parameters is defined by

locale *name* =
 fixes *x*₁ ... *x*_{*n*}
begin
 ...
end

where the variables *x*₁, ..., *x*_{*n*} are the parameters. Like the bound variables in a theorem they can be grouped by **and** and types can be specified for

some or all of them. The content between **begin** and **end** may consist of definitions and theorems which may use the parameter names like constant names. Content may also be added to an existing locale in the form

```
context name
begin
  ...
end
```

Therefore the **begin** ... **end** block can also be omitted in the locale definition and the locale can be filled later.

An instance of the parameterized theory fragment is created by “interpreting” the locale in the form

```
interpretation name term1 ... termn .
```

where *term*₁ ... *term*_{*n*} are the terms to be substituted for the locale parameters, their types must match the parameter types, i.e., must be instances of them. The final dot in the interpretation is a rudimentary proof. An actual proof is needed, if the locale definition specifies additional assumptions for the parameters.

Locales with Assumptions

Additional assumptions for locale parameters can be specified as propositions in the form

```
locale name =
  fixes x1 ... xn
  assumes A1 ... Am
begin
  ...
end
```

where the *A*₁, ..., *A*_{*m*} are propositions. Like in a theorem, they can be grouped by **and** and named. The names can be used to reference the assumptions as facts in proofs in the locale content. When the locale is interpreted, all the assumptions must be proved with the actual terms substituted for the parameters. Therefore the more general form of an interpretation is

```
interpretation name term1 ... termn <proof>
```

Extending Locales

A locale may extend one or more other locales using the form

```

locale name = name1 + ... + namen +
  fixes ...
  assumes ...
begin
  ...
end

```

where *name*₁ ... *name*_{*n*} are the names of the extended locales. Their parameters become parameters of the defined locale, inserted before the parameters declared by the **fixes** ... clause.

1.4 Isabelle Proofs

Every proposition stated as a fact in an Isabelle theory must be proven immediately by specifying a proof for it. A proof may have a complex structure of several steps and nested sub-proofs, its structure is part of the outer syntax.

1.4.1 Proof Context

Every proof is performed in a temporary environment which collects facts and other proof elements. This environment is called the “proof context”. At the end of the proof the proof context is disposed with all its content, only the proven fact remains in the enclosing entity.

The proof context may contain

- Facts: as usual, facts are valid propositions. However, they need not be globally valid, they can be assumed to be only valid locally during the proof.
- Goals: a goal is a proposition which has not yet been proven. Typically it is the duty of a proof to prove one or more goals in its proof context.
- Fixed variables: fixed variables are used to denote the “arbitrary but fixed” objects often used in a proof. They can be used in all facts and goals in the same proof context.
- Term abbreviations: these are names introduced locally for terms. Using such names for terms occurring in propositions it is often possible to denote propositions in a more concise form.

The initial proof context in a theorem of the form **theorem** *prop* *<proof>* has the proposition *prop* as the only goal and is otherwise empty.

1.4.2 Proof Procedure

Assume you want to prove a derivation rule $A \implies C$ with a single assumption A and the conclusion C . The basic procedure to build a proof for it is to construct a sequence of the form $F_1 \implies F_2, F_2 \implies F_3, F_3 \implies \dots \implies F_{n-1}, F_{n-1} \implies F_n$ from rules $RA_i \implies RC_i$ for $i=1 \dots n-1$ which are already known to be valid (i.e., facts) where F_1 matches with A and RA_1 , F_n matches with C and RC_{n-1} , and every other F_i matches with RA_i and RC_{i-1} .

The sequence can be constructed from left to right (called “forward reasoning”) or from right to left (called “backward reasoning”) or by a combination of both.

Consider the rule $(x::nat) < 5 \implies 2*x+3 \leq 2*5+3$. A proof can be constructed from the two example rules *example1* and *example2* from the previous sections as the sequence $(x::nat) < 5 \implies 2*x \leq 2*5, 2*x \leq 2*5 \implies 2*x+3 \leq 2*5+3$ consisting of three facts.

Forward reasoning starts by assuming A to be a local fact and incrementally constructs the sequence from it. An intermediate result is a part $F_1 \implies \dots \implies F_i$ of the sequence, here F_i is the “current fact”. A forward reasoning step consists of stating a proposition F_{i+1} and proving it to be a new local fact from the current fact F_i using a valid rule $RA_i \implies RC_i$. The step results in the extended sequence $F_1 \implies \dots \implies F_i, F_i \implies F_{i+1}$ and the new current fact F_{i+1} . When a step successfully proves a current fact F_n which matches the conclusion C the proof is complete.

Backward reasoning starts at the conclusion C of the single goal $A \implies C$ and incrementally constructs the sequence from it backwards. An intermediate result is a part $F_i \implies \dots \implies F_n$ of the sequence where F_n matches C , here the remaining part of the sequence $F_1 \implies F_i$ is the “current goal”. A backward reasoning step consists of applying a proof method to F_i which constructs a new current goal $F_1 \implies F_{i-1}$ and the extended sequence $F_{i-1} \implies F_i, F_i \implies \dots \implies F_n$. When a step produces the new current goal $F_1 \implies F_1$, which is trivially valid, the proof is complete.

Note the slight difference in how the steps are specified: A forward step specifies the new current fact F_{i+1} and then proves it. A backward step specifies the proof method, the new current goal $F_1 \implies F_{i-1}$ is constructed by the method and is not an explicit part of the proof text. For that reason a proof constructed by forward reasoning is usually easier to read and write than a proof constructed by backward reasoning, since in the former case the sequence of the facts F_i is explicitly specified in the proof text, whereas in the latter case the sequence of the facts F_i is implicitly constructed and the proof text specifies only the methods.

However, since every forward reasoning step again requires a proof as its part (a “subproof”), no proof can be written using only forward reasoning

steps. The main idea of writing “good” proofs is to use nested forward reasoning until every subproof is simple enough to be done in a single backward reasoning step, i.e., the proof method directly goes from the conclusion to the assumption.

Unification

The matching at the beginning and end of the sequence and when joining the used rules is done by “unification”. Two propositions P and Q are unified by substituting terms for unknowns in P and Q so that the results become syntactically equal.

Since only the $RA_i \implies RC_i$ are facts containing unknowns, only they are modified by the unification, A and C remain unchanged.

Note that when an unknown is substituted by a term in RA_i , the same unknown must be substituted by the same term in RC_i and vice versa, to preserve the validness of the rule $RA_i \implies RC_i$. In other words, the sequence is usually constructed from specializations of the facts $RA_i \implies RC_i$ where every conclusion is syntactically equal to the assumption of the next rule.

In the example the assumption $?x < ?c$ of rule *example1* is unified with $(x::nat) < 5$ by substituting the term 5 for the unknown $?c$, and the variable x for the unknown $?x$ resulting in the specialized rule $(x::nat) < 5 \implies n*x \leq n*5$. The conclusion $?x + ?m \leq ?c + ?m$ of rule *example2* is unified with $2*x+3 \leq 2*5+3$ by substituting the term $2*x$ for the unknown $?x$, the term $2*5$ for the unknown $?c$, and the term 3 for the unknown $?m$ resulting in the specialized rule $2*x \leq 2*5 \implies 2*x+3 \leq 2*5+3$. Now the two specialized rules can be joined by substituting the term 2 for the unknown $?n$ in the first, resulting in the sequence which constitutes the proof.

Multiple Assumptions

If the rule to be proven has more than one assumption A the sequence to be constructed becomes a tree where the branches start at (copies of) the assumptions A_1, \dots, A_n and merge to finally lead to the conclusion C . Two branches which end in facts F_{1n} and F_{2m} are joined by a step $\llbracket F_{1n}; F_{2m} \rrbracket \implies F_1$ to a common branch which continues from fact F_1 .

Now a forward reasoning step may use several current facts to prove a new current fact. Therefore all proven local facts are stored in the proof context for possible later use. Every forward reasoning step selects a subset of the stored local facts as the current facts and uses them to prove a new local fact from them.

A backward reasoning step may now produce several new current goals, which belong to different branches in the tree. A step always produces the

goals for all branches, therefore the previous goal is never used again in a step and is removed from the proof context after the step. When a current goal has the form $A \implies A$ the proof method “*assumption*” removes it from the proof context without producing a new goal. Thus a proof ends when no goal remains in the proof context.

The set of current goals is called the “goal state” of the proof. Since it is not visible in the proof text, the interactive editor displays the current goal state in the separate State panel and optionally also in the Output panel (see Sections 1.2.2 and 1.2.3), according to the cursor position in the proof text.

Proving from External Facts

The branches in the fact tree need not always start at an assumption A_i , they may also start at an “external” fact which is not part of the local proof context. In such cases the used external facts are referenced by their names. In that way a proof can use facts from the enclosing theory and a subproof can use facts from the enclosing proof(s) and the enclosing toplevel theory. In particular, if the proposition of a theorem has no assumptions, i.e., the proposition is a formula and consists only of the conclusion C , every proof must start at one or more external facts.

1.4.3 Basic Proof Structure

A proof is written in outer syntax and describes how the fact tree is constructed which leads from the assumptions or external facts to the conclusion.

Proof Modes

When writing a proof the “proof mode” determines the mode of operation: whether forward reasoning (mode: *proof(state)*) or backward reasoning (mode: *proof(prove)*) is used.

The mode names refer to what is done in the next steps: In mode *proof(state)* facts are “stated” which lead to the conclusion, whereas in mode *proof(prove)* the goals are “proven”, leading to assumptions and external facts.

At the beginning of a proof the mode is always *proof(prove)*, i.e., backward reasoning. In the course of the proof it is possible to switch to forward reasoning mode *proof(state)*, but not back again. After switching to forward reasoning the proof must be completed in forward reasoning mode, only at the end a last backward reasoning step may be applied.

However, in forward reasoning for every stated fact a (sub-)proof must be specified, which again starts in backward reasoning mode. This way it is

possible to freely switch between both modes in the course of a proof with nested subproofs.

Proof Syntax

If BS_i denote backward reasoning steps and FS_i denote forward reasoning steps, the general form of a proof is

$$\begin{array}{l} BS_1 \dots BS_n \\ \mathbf{proof} \ BS_{n+1} \\ \quad FS_1 \dots FS_m \\ \mathbf{qed} \ BS_{n+2} \end{array}$$

The last step BS_{n+2} can be omitted if it is not needed.

The part **proof** BS_{n+1} switches from backward reasoning mode *proof*(*prove*) to forward reasoning mode *proof*(*state*).

The part **proof** ... **qed** can be replaced by **done**, then the proof only consists of backward reasoning steps and has the form $BS_1 \dots BS_n$ **done**. Such proofs are called “proof scripts”.

If the backward reasoning steps $BS_1 \dots BS_n$ are omitted the proof only consists of the forward reasoning part and has the form

$$\begin{array}{l} \mathbf{proof} \ BS_1 \\ \quad FS_1 \dots FS_m \\ \mathbf{qed} \ BS_2 \end{array}$$

where BS_2 can also be omitted. Such proofs are called “structured proofs”.

A structured proof can be so simple, that it has no forward reasoning steps. For this case the syntax

$$\mathbf{by} \ BS_1 \ BS_2$$

abbreviates the form **proof** BS_1 **qed** BS_2 . Again, BS_2 can be omitted which leads to the form

$$\mathbf{by} \ BS_1$$

In this form the proof consists of a single backward reasoning step which directly leads from the conclusion C to the assumptions and used external facts.

Fake Proofs

A proof can also be specified as

sorry

This is a “fake proof” which turns the proposition to a fact without actually proving it.

A fake proof can be specified at any point in backward reasoning mode, so it can be used to abort a proof script in the form $BS_1 \dots BS_n$ **sorry**.

A structured proof in forward reasoning mode cannot be aborted in this way, however, subproofs can be specified as fake proofs. This makes it possible to interactively develop a structured proof in a top-down way, by first stating all required facts with fake proofs and then replacing the fake proofs by actual proofs.

Nested Proof Contexts

The proof contexts in a structured proof can be nested. In a nested context the content of the enclosing contexts is available together with the local content. When a nested context is ended, it is removed together with all its local content.

A nested proof context is created syntactically by enclosing forward reasoning steps in braces:

$$FS_1 \dots FS_m \{ FS_{m+1} \dots FS_n \} FS_{n+1} \dots$$

Note that according to the description until now the nested context is useless, because the facts introduced by its forward reasoning steps are removed at its end and cannot contribute to the proof. How the content of a nested context can be “exported” and preserved for later use will be explained further below.

For names, nested contexts behave like a usual block structure: A name can be redefined in a nested context, then the named object in the outer context becomes inaccessible (“shadowed”) in the inner context, but becomes accessible again when the inner context ends.

When two nested contexts follow each other immediately, this has the effect of “clearing” the content of the inner contexts: the content of the first context is removed and the second context starts being empty. This can be denoted by the keyword

next

which can be thought of being equivalent to a pair } { of adjacent braces.

Moreover the syntax **proof** $method_1$ $FS_1 \dots FS_n$ **qed** $method_2$ automatically wraps the forward steps $FS_1 \dots FS_n$ in a nested context. Therefore it is possible to denote a structured proof which only consists of a sequence of nested contexts without braces as

proof $method_1$

$FS_{11} \dots FS_{1m1} \textbf{next} FS_{21} \dots FS_{2m2} \textbf{next} \dots \textbf{next} FS_{n1} \dots FS_{nmn}$
qed $method_{n+2}$

where each occurrence of **next** clears the content of the context. The goal state of the proof is maintained in the enclosing outermost context, thus it is preserved over the whole sequence of nested contexts.

1.4.4 Backward Reasoning Steps

A backward reasoning step consist of applying a proof method.

Proof Methods

Proof methods are basically denoted by method names, such as *standard*, *simp*, or *rule*. A proof method name can also be a symbol, such as $-$.

A method may have arguments, then it usually must be delimited by parentheses such as in (*rule example1*) or (*simp add: example2*), where *example1* and *example2* are fact names.

Methods can be applied to the goal state, they modify the goal state by removing and adding goals.

The effect of applying a method is determined by its implementation and must be known to the proof writer. Isabelle supports a large number of proof methods. Methods used for the proofs described in this manual are described in Section 1.5.

Method Application

A standalone method application step is denoted as

apply *method*

where *method* denotes the proof method to be applied.

The backward reasoning steps which follow **proof** and **qed** in a structured proof are simply denoted by the applied method. Hence the general form of a proof where all backward reasoning steps are method applications is

apply $method_1 \dots \textbf{apply}$ $method_n$
proof $method_{n+1}$
 $FS_1 \dots FS_m$
qed $method_{n+2}$

where $FS_1 \dots FS_m$ are forward reasoning steps.

1.4.5 Forward Reasoning Steps

A forward reasoning step consist of stating and proving a fact.

Stating a Fact

A fact is stated in the form

have *prop* $\langle proof \rangle$

where *prop* is a proposition in inner syntax and $\langle proof \rangle$ is a (sub-) proof for it. This form is similar to the specification of a theorem in a theory and has a similar effect in the local proof context.

As for a theorem the fact can be named:

have *name: prop* $\langle proof \rangle$

Note that the alternative form of a theorem using **fixes**, **assumes**, and **shows** (see Section 1.3.7) is not available for stating facts in a proof.

The subproof $\langle proof \rangle$ uses a nested context, therefore all content of the enclosing proof context is available there and can be referenced by name, as long as the name is not shadowed by a redefinition in the subproof. Note that the name given to the fact to be proven cannot be used to access it in the subproof, because it is only assigned after the proof has been finished.

Proving a Goal

A forward reasoning proof ends, if the last stated fact F_n unifies with the conclusion C . Therefore a special form of stating a fact exists, which, after proving the fact, replaces free variables by unknowns (which is called “exporting the fact”) and tries to unify it with the conclusion of a goal in the goal state. If successful, it removes the goal from the goal state:

show *prop* $\langle proof \rangle$

The syntax is the same as for **have**. If the unification with some goal conclusion is not successful the step is erroneous and the proof cannot be continued, in the interactive editor an error message is displayed.

The **show** step tries to match and remove a goal from the innermost enclosing proof context which maintains a goal state. This is one way how a fact proven in a nested context can affect an enclosing context and thus contribute to the proof there. Since the forward reasoning steps in a structured proof are wrapped in a nested context, while the goal state is maintained in the enclosing outer context, the **show** step affects the goal state of the enclosing proof.

The **have** and **show** steps are called “goal statements”, because they state the proposition *prop* as a goal which is then proven by the $\langle proof \rangle$.

Note that the proposition *prop* in a **show** statement often is the same proposition which has been specified as conclusion *C* in the proposition $\llbracket A_1; \dots; A_n \rrbracket \implies C$ which should be proven by the proof. To avoid repeating it, Isabelle automatically provides the abbreviation *?thesis* for it. So in the simplest case the last step of a forward proof can be written as

show *?thesis* $\langle proof \rangle$

If, however, the application of *method* in a structured proof **proof** *method* ... modifies the original goal, this modification is not reflected in *?thesis*. So a statement **show** *?thesis* $\langle proof \rangle$ will usually not work, because *?thesis* no more unifies with the conclusion of the modified goal. Instead, the proof writer must know the modified goal and specify its conclusion explicitly as proposition in the **show** statement. If the *method* splits the goal into several new goals, several **show** statements may be needed to remove them.

To test whether a proposition unifies with the conclusion of a goal in the goal state, a **show** statement can be specified with a fake proof:

show *prop* **sorry**

If that statement is accepted, the proposition unifies with the conclusion of a goal and removes it.

1.4.6 Facts as Proof Input

If a linear fact sequence $F_1 \implies \dots \implies F_n$ is constructed in forward reasoning mode in the form

have F_1 $\langle proof \rangle_1$
...
have F_n $\langle proof \rangle_n$

every fact F_i needs to refer to the previous fact F_{i-1} in its proof *proof_i*. This can be done by naming all facts

have *name₁*: F_1 $\langle proof \rangle_1$
...
have *name_n*: F_n $\langle proof \rangle_n$

and refer to F_{i-1} in *proof_i* by *name_{i-1}*.

Isabelle supports an alternative way by passing facts as input to a proof.

Using Input Facts in a Proof

The input facts are passed as input to the first method applied in the proof. In a proof script it is the method applied in the first **apply** step, in a structured proof **proof method** ... it is *method*.

Every proof method accepts a set of facts as input. Whether it processes them and how it uses them depends on the kind of method. Therefore input facts for a proof only work in the desired way, if a corresponding method is used at the beginning of the proof. See Section 1.5 for descriptions how methods process input facts.

Inputting Facts into a Proof

In backward reasoning mode *proof(prove)* facts can be input to the remaining proof $\langle proof \rangle$ by

using $name_1 \dots name_n \langle proof \rangle$

where the $name_i$ are names of facts or fact sets. The union of all referred facts is input to the proof following the **using** specification. In a proof script it is input to the next **apply** step. If a structured proof follows, it is input to its initial method. Since in backward reasoning mode no local facts are stated by previous steps, only external facts can be input this way.

In forward reasoning mode *proof(state)* fact input is supported with the help of the special fact set name *this*. The statement

then

inputs the facts named *this* to the proof of the following fact statement.

The statement **then** must be immediately followed by a goal statement (**have** or **show**). This is enforced by a special third proof mode *proof(chain)*. In it only a goal statement is allowed, **then** switches to this mode, the goal statement switches back to mode *proof(state)* after its proof.

Note that **then** is allowed in forward reasoning mode, although it does not state a fact. There are several other such auxiliary statements allowed in mode *proof(state)* in addition to the goal statements **have** and **show**.

The fact set *this* can be set by the statement

note $name_1 \dots name_n$

Therefore the statement sequence

note $name_1 \dots name_n$
then have *prop* $\langle proof \rangle$

inputs the union of all facts referred by $name_1 \dots name_n$ to the $\langle proof \rangle$, in the same way as **using** inputs them to the remaining proof following it.

The statement sequence

note $name_1 \dots name_n$ **then**

can be abbreviated by the statement

from $name_1 \dots name_n$

Like **then** it switches to mode $proof(chain)$ and it inputs the union of the facts referred by $name_1 \dots name_n$ to the proof of the following goal statement.

1.4.7 Fact Chaining

In both cases described for fact input until now, the facts still have been referred by names. This can be avoided by automatically using a stated fact as input to the proof of the next stated fact. That is called “fact chaining”.

Automatic Update of the Current Facts

Fact chaining is achieved, because Isabelle automatically updates the fact set *this*. Whenever a new fact is added to the proof context, the set *this* is redefined to contain (only) this fact. In particular, after every goal statement *this* names the new proven fact. Therefore the fact set *this* is also called the “current facts”.

Thus a linear sequence of facts can be constructed by

have $F_1 \langle proof \rangle_1$
then have $F_2 \langle proof \rangle_2$
 \dots
then have $F_n \langle proof \rangle_n$

Now in every $proof_i$ the fact F_{i-1} is available as input and can be used to prove F_i .

Chaining can be combined with explicit fact referral by a statement of the form

note *this* $name_1 \dots name_n$

It sets *this* to the union of *this* and the $name_1 \dots name_n$, i.e., it adds the $name_1 \dots name_n$ to *this*. In this way the current facts can be extended with other facts and then chained to the proof of the next stated fact.

The statement sequence

note *this* $name_1 \dots name_n$ **then**

can be abbreviated by the statement

with $name_1 \dots name_n$

Like **then** it switches to mode *proof(chain)* and it inputs the union of the facts referred by $name_1 \dots name_n$ together with the current facts to the proof of the following goal statement.

If a proof consists of a fact tree with several branches, every branch can be constructed this way. Before switching to the next branch the last fact must be named, so that it can later be used to prove the fact where the branches join. A corresponding proof pattern for two branches which join at fact F is

```

have  $F_{11}$   $\langle proof \rangle_{11}$ 
then have  $F_{12}$   $\langle proof \rangle_{12}$ 
...
then have  $name_1: F_{1m}$   $\langle proof \rangle_{1m}$ 
have  $F_{21}$   $\langle proof \rangle_{21}$ 
then have  $F_{22}$   $\langle proof \rangle_{22}$ 
...
then have  $F_{2n}$   $\langle proof \rangle_{2n}$ 
with  $name_1$  have  $F$   $\langle proof \rangle$ 

```

Naming and Grouping Current Facts

Since the fact set built by a **note** statement is overwritten by the next stated fact, it is possible to give it an explicit name in addition to the name *this* in the form

note $name = name_1 \dots name_n$

The *name* can be used later to refer to the same fact set again, when *this* has already been updated. Defining such names is only possible in the **note** statement, not in the abbreviated forms **from** and **with**.

The facts specified in **note**, **from**, **with**, and **using** can be grouped by separating them by **and**. Thus it is even possible to write

from $name_1$ **and** ... **and** $name_n$ **have** *prop* $\langle proof \rangle$

In the case of a **note** statement every group can be given an additional explicit name as in

note $name_1 = name_{11} \dots name_{1m1}$ **and** ... **and** $name_n = name_{n1} \dots name_{nmn}$

Accessing Input Facts in a Structured Proof

At the beginning of a structured proof the set name *this* is undefined, the name cannot be used to refer to the input facts (which are the current facts in the enclosing proof). To access the input facts they must be named before they are chained to the goal statement, then they can be referenced in the subproof by that name. For example in

```
note input = this  
then have prop  $\langle proof \rangle$ 
```

the input facts can be referenced by the name *input* in $\langle proof \rangle$.

Exporting the Current Facts of a Nested Context

At the end of a nested context (see Section 1.4.3) the current facts are automatically exported to the enclosing context, i.e. they become available there as the fact set named *this*, replacing the current facts before the nested context. This is another way how facts from a nested context can contribute to the overall proof.

Basically, only the last fact is current at the end of a context. Arbitrary facts can be exported from the nested context by explicitly making them current at its end, typically using a **note** statement:

```
... {  
  have  $f_1$ : prop1  $\langle proof \rangle_1$   
  ...  
  have  $f_n$ : propn  $\langle proof \rangle_n$   
  note  $f_1 \dots f_n$   
} ...
```

Here all facts are named and the **note** statement makes them current by referring them by their names. Note, that the names are only valid in the nested context and cannot be used to refer to the exported facts in the outer context.

The exported facts can be used in the outer context like all other current facts by directly chaining them to the next stated fact:

```
... { ... } then have prop  $\langle proof \rangle$  ...
```

or by naming them for later use, with the help of a **note** statement:

```
... { ... } note name = this ...
```

1.4.8 Assuming Facts

The Assumptions A_1, \dots, A_n of the proposition to be proven are needed as facts in the proof context to start the branches of the fact tree. However, they are not automatically inserted, that must be done explicitly.

Introducing Assumed Facts

An assumption is inserted in the proof context by a statement of the form

assume *prop*

Several assumptions can be inserted in a single **assume** statement of the form

assume *prop*₁ ... *prop*_{*n*}

As usual, the assumptions can be named and grouped by **and**.

Like goal statements an **assume** statement makes the assumed facts current, i.e. it updates the set *this* to contain the specified propositions as facts, so that they can be chained to a following goal statement. This way the fact sequence $A \implies F_1, F_1 \implies \dots$ of a proof using fact chaining can be started:

assume *A*
then have F_1
...

Alternatively, the assumed facts can be named:

assume *name:* *prop*₁ ... *prop*_{*n*}

so that they can be referred by name in the rest of the proof. Then the fact chain can be started in the form

assume *a:* *A*
from *a* **have** F_1
...

This can be useful if the proof has several branches which all start at the same assumption.

Admissible Assumed Facts

In an **assume** statement no proof is needed, since these propositions are only “assumed” to be valid. Therefore, only the propositions occurring as assumptions in the goal $\llbracket A_1; \dots; A_n \rrbracket \implies C$ to be proven are allowed here.

Actually, this condition is not checked for the **assume** statement, an arbitrary proposition can be specified by it. The condition becomes only relevant in subsequent **show** statements. When the fact proven by a show statement is tried to be unified with the conclusion of a goal, additionally each proposition stated in an **assume** statement is tested for unifying with an assumption of the same goal. If that is not satisfied, the **show** statement fails. Therefore, if a proposition is used in an **assume** statement which does not unify with an assumption in a goal, the proof cannot be completed, because all **show** statements will fail.

Exporting Facts with Assumptions

More generally, whenever a local fact F is exported from a proof context, it is combined with all locally assumed facts AF_1, \dots, AF_n to the derivation rule $\llbracket AF_1; \dots; AF_n \rrbracket \implies F$. This reflects the intention of the local assumptions: They may have been used locally to prove F without knowing whether they are valid. So outside the local context F is only known to be valid if all the assumptions are valid.

If the fact F is itself a derivation rule $\llbracket A_1; \dots; A_n \rrbracket \implies C$ then the locally assumed facts are added, resulting in the exported rule $\llbracket AF_1; \dots; AF_n; A_1; \dots; A_n \rrbracket \implies C$.

If the fact F has been proven in a **show** statement it is also exported in this way, resulting in a derivation rule. It matches with a goal if the conclusion of the exported rule unifies with the goal conclusion and if every assumption of the exported rule unifies with an assumption in the goal. This way of matching a rule with a goal is called “refinement”. So the condition for a successful **show** statement can be stated as “the exported fact must refine a goal”.

Avoiding Repeated Propositions for Assumed Facts

To make **show** statements succeed, an **assume** statement will usually repeat one or more assumptions from the proposition to be proven. This is similar to a **show** statement, which usually repeats the conclusion of that proposition. However, unlike *?thesis* for the conclusion, there is no abbreviation provided for the assumptions.

To avoid repeating propositions in **assume** statements, the proposition to be proven can be specified in the form (see Section 1.3.7)

theorem

assumes A_1 **and** ... **and** A_n

shows C

$\langle proof \rangle$

This form automatically inserts the assumptions as assumed facts in the proof context. No **assume** statements are needed, the assumptions need not be repeated, and the assumed facts are always safe for **show** statements.

The assumptions can still be named and referred by name in the proof. A proof can be started at assumption A_1 in the form

```
theorem
  assumes  $name_1: A_1$  and ... and  $name_n: A_n$ 
  shows  $C$ 
  proof  $\langle method \rangle$ 
    from  $name_1$  have  $F_1$   $\langle proof \rangle$ 
    ...
```

Additionally, the set of all assumption specified in this form of a theorem is automatically named *assms*. Since unneeded assumptions usually do not harm in a proof, each proof branch can be started in the form

```
from assms have  $F_1$ 
...
```

but it is usually clearer for the reader to specify only the relevant assumption(s) by explicit names.

In subproofs **assume** statements cannot be avoided in this way, because propositions in goal statements cannot be specified using **assumes** and **shows**. However, goal statements usually specify only a fact as proposition without assumptions. Instead of assumed facts the subproof can either use facts provided as input, or use external facts from the enclosing proof context by referring them by name.

Presuming Facts

It is also possible to use a proposition as assumed fact which does not unify with an assumption in a goal, but can be proven from them. In other words, the proof is started somewhere in the middle of the fact tree, works in forward reasoning mode, and when it reaches the conclusion the assumed fact remains to be proven. The statement

```
presume prop
```

inserts such a presumed fact into the proof context.

When a fact is exported from a context with presumed facts, they do not become a part of the exported rule. Instead, at the end of the context for each presumed fact F_p a new goal $\llbracket A_1; \dots; A_n \rrbracket \implies F_p$ is added to the enclosing goal state. So the proof has to continue after proving all original goals and is only finished when all such goals for presumed facts have been proven as well.

1.4.9 Fixing Variables

Variables occurring in the proposition of a theorem can be used in the proof as well, they are universally bound in the whole proof context, if they are not explicitly bound in the proposition, which restricts their use to the proposition itself. Thus in

theorem $\wedge x::nat. x < 3 \implies x < 5$ $\langle proof \rangle$

the variable x is restricted to the proposition and is not accessible in $\langle proof \rangle$, whereas in

theorem $(x::nat) < 3 \implies x < 5$ $\langle proof \rangle$

and

theorem $x < 3 \implies x < 5$ **for** $x::nat$ $\langle proof \rangle$

the variable x is accessible in $\langle proof \rangle$.

Local Variables

Additional local variables can be introduced (“fixed”) in a proof context in mode *proof*(*state*) by the statement

fix $x_1 \dots x_n$

As usual, types can be specified for (some of) the variables and the variables can be grouped by **and**.

If a variable name is used in a proof context without explicitly fixing it, it either refers to a variable in an enclosing context or in the proposition to be proven, or it is free. If it is explicitly fixed it names a variable which is different from all variables with the same name in enclosing contexts and the proposition to be proven.

A fixed local variable is common to the whole local context. If it occurs in several local facts it always is the same variable, it is not automatically restricted to the fact, as for toplevel theorems. Hence in

fix $x::nat$
assume $a: x < 3$
have $x < 5$ $\langle proof \rangle$

the $\langle proof \rangle$ may refer to fact a because the x is the same variable in both facts.

By convention variable names are often short consisting of one or two letters, whereas constants defined on toplevel in a theory have longer and more

descriptive names. Therefore it is usually not necessary to explicitly fix the variables in the proposition of a theorem to prevent name clashes with constants. By contrast, in a nested proof context there may be other variables with the same name in enclosing contexts, therefore it is recommended to explicitly fix all local variables.

Exporting Facts with Local Variables

Explicitly fixing variables in a proof context is not only important for avoiding name clashes. If a fact is exported from a proof context, all fixed local variables are replaced by unknowns, other variables remain unchanged. Since unification only works for unknowns, it makes a difference whether a fact uses a local variable or a variable which originates from an enclosing context or is free.

The proposition $x < 3 \implies x < 5$ can be proven by the statements

```
fix  $y::nat$ 
assume  $y < 3$ 
then show  $y < 5$  <proof>
```

because when the fact $y < 5$ is exported, the assumption is added (as described in Section 1.4.8) and then variable y is replaced by the unknown $?y$ because y has been locally fixed. The result is the rule $?y < 3 \implies ?y < 5$ which unifies with the proposition.

If, instead, y is not fixed, the sequence

```
assume  $(y::nat) < 3$ 
then have  $y < 5$  <proof>
```

still works and the local fact $y < 5$ can be proven, but it cannot be used with the **show** statement to prove the proposition $x < 3 \implies x < 5$, because the exported rule is now $y < 3 \implies y < 5$ which does not unify with the proposition, it contains a different variable instead of an unknown.

1.4.10 Term Abbreviations

A term abbreviation is a name for a proposition or a term in it.

Defining Term Abbreviations

A term abbreviation can be defined by a statement of the form

```
let  $?name = term$ 
```

Afterwards the name is “bound” to the term and can be used in place of the term in propositions and other terms, as in:

```
let ?lhs = 2*x+3
let ?rhs = 2*5+3
assume x < 5
have ?lhs ≤ ?rhs <proof>
```

The name *?thesis* (see Section 1.4.5) is a term abbreviation of this kind.

A **let** statement can define several term abbreviations in the form

```
let ?name1 = term1 and ... and ?namen = termn
```

A **let** statement can occur everywhere in mode *proof(state)*. However, it does not preserve the current facts, the fact set *this* becomes undefined by it.

Pattern Matching

Note that term abbreviations have the form of “unknowns” (see Section 1.3.7), although they are defined (“bound”). The reason is that they are actually defined by unification.

The more general form of a **let** statement is

```
let pattern = term
```

where *pattern* is a term which may contain unbound unknowns. The **let** statement unifies *pattern* and *term*, i.e., it determines terms to substitute the unknowns, so that the pattern becomes syntactically equal to *term*. If that is not possible, an error is signaled, otherwise the unknowns are bound to the substituting terms.

The **let** statement

```
let ?lhs ≤ ?rhs = 2*x+3 ≤ 2*5+3
```

binds the unknowns to the same terms as the two **let** statements above.

The pattern and the term may contain unknowns which are already bound. They are substituted by their bound terms before the pattern matching is performed. Thus the term and the pattern can be constructed with the help of abbreviation which have been defined previously. A useful example is matching a pattern with *?thesis*:

```
theorem x < 5 ⇒ 2*x+3 ≤ 2*5+3
proof method
  let ?lhs ≤ ?rhs = ?thesis
  ...
```

to reuse parts of the conclusion in the proof.

Note that the unknowns are only bound at the end of the whole **let** statement. In the form

let $pattern_1 = term_1$ **and** ... **and** $pattern_n = term_n$

the unknowns in $pattern_i$ cannot be used to build $term_{i+1}$ because they are not yet bound. In contrast, in the sequence of **let** statements

let $pattern_1 = term_1$
...
let $pattern_n = term_n$

the unknowns in $pattern_i$ can be used to build $term_{i+1}$ because they are already bound.

1.4.11 Accumulating Facts

Instead of giving individual names to facts in the proof context, facts can be collected in named fact sets. Isabelle supports the specific fact set named *calculation* and provides statements for managing it.

The fact set *calculation* is intended to accumulate current facts for later use. Therefore it is typically initialized by the statement

note $calculation = this$

and afterwards it is extended by several statements

note $calculation = calculation\ this$

After the last extension the facts in the set are chained to the next proof:

note $calculation = calculation\ this$ **then have** ...

Support for Fact Accumulation

Isabelle supports this management of *calculation* with two statements. The statement

moreover

is equivalent to **note** $calculation = this$ when it occurs the first time in a context, afterwards it behaves like **note** $calculation = calculation\ this$ but without making *calculation* current, instead, it leaves the current fact(s) unchanged. The statement

ultimately

is equivalent to **note** $calculation = calculation\ this$ **then**, i.e., it adds the current facts to the set, makes the set current, and chains it to the next goal statement.

Due to the block structure of nested proofs, the *calculation* set can be reused in nested contexts without affecting the set in the enclosing context. The first occurrence of **moreover** in the nested context initializes a fresh local *calculation* set. Therefore fact accumulation is always local to the current proof context.

Accumulating Facts in a Nested Context

Fact accumulation can be used for collecting the facts in a nested context for export (see Section 1.4.7) without using explicit names for them:

```
... {  
  have prop1 <proof>1  
  moreover have prop2 <proof>2  
  ...  
  moreover have propn <proof>n  
  moreover note calculation  
} ...
```

Accumulating Facts for Joining Branches

Fact accumulation can also be used for collecting the facts at the end of joined fact branches in a proof and inputting them to the joining step. A corresponding proof pattern for two branches which join at fact F is

```
have F11 <proof>11  
then have F12 <proof>12  
...  
then have F1m <proof>1m  
moreover have F21 <proof>21  
then have F22 <proof>22  
...  
then have F2n <proof>2n  
ultimately have F <proof>
```

The **moreover** statement starts the second branch and saves the fact F_{1m} to *calculation*. The **ultimately** statement saves the fact F_{2n} to *calculation* and then inputs the set to the proof of F .

Note that **moreover** does not chain the current facts to the following goal statement.

Using nested contexts sub-branches can be constructed and joined in the same way.

1.4.12 Equational Reasoning

Often informal proofs on paper are written in the style

$$2*(x+3) = 2*x+6 \leq 3*x+6 < 3*x+9 = 3*(x+3)$$

to derive the fact $2*(x+3) < 3*(x+3)$. Note that the formula shown above is not a wellformed proposition because of several occurrences of the toplevel relation symbols $=$, \leq and $<$. Instead, the formula is meant as an abbreviated notation of the fact sequence

$$2*(x+3) = 2*x+6, 2*x+6 \leq 3*x+6, 3*x+6 < 3*x+9, 3*x+9 = 3*(x+3)$$

which sketches a proof for $2*(x+3) < 3*(x+3)$. This way of constructing a proof is called “equational reasoning” which is a specific form of forward reasoning.

Transitivity Rules

The full proof needs additional facts which must be inserted into the sequence. From the first two facts the fact $2*(x+3) \leq 3*x+6$ is derived, then with the third fact the fact $2*(x+3) < 3*x+9$ is derived, and finally with the fourth fact the conclusion $2*(x+3) < 3*(x+3)$ is reached. The general pattern of these additional derivations can be stated as the derivation rules $\llbracket a = b; b \leq c \rrbracket \implies a \leq c$, $\llbracket a \leq b; b < c \rrbracket \implies a < c$, and $\llbracket a < b; b = c \rrbracket \implies a < c$.

Rules of this form are called “transitivity rules”. They are valid for relations such as equality, equivalence, orderings, and combinations thereof.

This leads to the general form of a proof constructed by equational reasoning: every forward reasoning step starts at a fact F_i of the form $a \ r \ b$ where r is a relation symbol. It proves an intermediate fact H_i of the form $b \ r \ c$ where r is the same or another relation symbol and uses a transitivity rule to prove the fact F_{i+1} which is $a \ r \ c$. In this way it constructs a linear sequence of facts which leads to the conclusion.

The intermediate facts H_i are usually proven from assumptions or external facts, or they may have a more complex proof which forms an own fact branch which ends at H_i and is joined with the main branch at F_{i+1} with the help of the transitivity rule.

Support for Equational Reasoning

Isabelle supports equational reasoning in the following form. It provides the statement

also

which expects that the set *calculation* contains the fact F_i and the current fact *this* is the fact H_i . It automatically selects an adequate transitivity rule, uses it to derive the fact F_{i+1} and replaces F_i in *calculation* by it. Upon its first use in a proof context **also** simply stores the current fact *this* in *calculation*. The statement

finally

behaves in the same way but additionally makes the resulting fact F_{i+1} current, i.e., puts it into the set *this*, and chains it into the next goal statement. In other words, **finally** is equivalent to **also from** *calculation*.

Note that **also** behaves like **moreover** and **finally** behaves like **ultimately**, both with additional application of the transitivity rule.

Additionally, Isabelle automatically maintains the term abbreviation \dots (which is the three-dot-symbol available for input in the Symbols panel (see Section 1.2.4) of the interactive editor in tab “Punctuation”) for the term on the right hand side of the current fact. Together, the example equational reasoning proof from above can be written

```
have  $2*(x+3) = 2*x+6$  <proof>
also have  $\dots \leq 3*x+6$  <proof>
also have  $\dots < 3*x+9$  <proof>
also have  $\dots = 3*(x+3)$  <proof>
finally show ?thesis <proof>
```

where *?thesis* abbreviates the conclusion $2*(x+3) < 3*(x+3)$. This form is quite close to the informal paper style of the proof.

Determining Transitivity Rules

To automatically determine the transitivity rule used by **also** or **finally**, Isabelle maintains the dynamic fact set (see Section 1.3.7) named *trans*. It selects a rule from that set according to the relation symbols used in the facts in *calculation* and *this*.

A transitivity rule which is not in *trans* can be explicitly specified by name in the form

```
also (name)
finally (name)
```

1.5 Proof Methods

The basic building blocks of Isabelle proofs are the proof methods which modify the goal state. If there are several goals in the goal state it depends on the specific method which goals are affected by it. In most cases only the first goal is affected.

1.5.1 The empty Method

The empty method is denoted by a single minus sign

—

If no input facts are passed to it, it does nothing, it does not alter the goal state.

Otherwise it affects all goals by inserting the input facts as assumptions after the existing assumptions. If the input facts are F_1, \dots, F_m a goal of the form $\llbracket A_1; \dots; A_n \rrbracket \Longrightarrow C$ is replaced by the goal $\llbracket A_1; \dots; A_n; F_1, \dots, F_m \rrbracket \Longrightarrow C$. The reason for this behavior will be explained below.

The empty method is useful at the beginning of a structured proof of the form

proof *method* $FS_1 \dots FS_n$ **qed**

If the forward reasoning steps $FS_1 \dots FS_n$ shall process the unmodified original goal the empty method must be specified for *method*, thus the structured proof becomes

proof — $FS_1 \dots FS_n$ **qed**

Note that it is possible to omit the *method* completely, but then it defaults to the method named *standard* which alters the goal state (see below).

1.5.2 Rule Application

As described in Section 1.4.2 the basic step in the construction of a proof is to establish the connection between a fact F_i and a fact F_{i+1} in the fact sequence. Assume that there is already a valid derivation rule $RA_i \Longrightarrow RC_i$ named r_i where RA_i unifies with F_i and RC_i unifies with F_{i+1} . Then the connection can be established by applying that rule.

The *rule* Method

A rule is applied by the method

rule name

where *name* is the name of a valid rule. The method only affects the first goal. If that goal has the form $\llbracket A_1; \dots; A_n \rrbracket \implies C$ and the rule referred by *name* has the form $\llbracket RA_1; \dots; RA_m \rrbracket \implies RC$ the method first unifies RC with the goal conclusion C . That yields the specialized rule $\llbracket RA_1'; \dots; RA_m' \rrbracket \implies RC'$ where RC' is syntactically equal to C and every RA_j' results from RA_j by substituting unknowns by the same terms as in RC' . Note that the goal normally does not contain unknowns, therefore C is not modified by the unification. If the unification fails the method cannot be executed on the goal state and an error is signaled. Otherwise the method replaces the goal by the m new goals $\llbracket A_1; \dots; A_n \rrbracket \implies RA_j'$.

If the rule has the form $RA \implies RC$ with only one assumption the method replaces the goal by the single new goal $\llbracket A_1; \dots; A_n \rrbracket \implies RA'$. If the rule is a formula RC without any assumptions the method removes the goal without introducing a new goal.

Using the *rule* Method for Backward Reasoning Steps

Assume that during a proof as described in Section 1.4.2 the intermediate fact sequence $F_{i+1} \implies \dots \implies F_n$ has already been constructed by backward reasoning and the current goal is $F_1 \implies F_{i+1}$. The backward reasoning step

apply (*rule* r_i)

will replace that goal by $F_1 \implies F_i$ and thus extend the fact sequence to $F_i \implies \dots \implies F_n$. The fact F_i is the specialized assumption RA_i' constructed by the method from the assumption RA_i of rule r_i .

Therefore the fact sequence $F_1 \implies \dots \implies F_n$ of the complete proof can be constructed by the proof script consisting of the backward reasoning steps

apply (*rule* r_{n-1})

...

apply(*rule* r_1)

Note however, that this proof script does not complete the proof, since it results in the goal $F_1 \implies F_1$. The proof method

assumption

must be used to process it. The method only affects the first goal. If that goal has the form $\llbracket A_1; \dots; A_n \rrbracket \implies C$ and one assumption A_i is syntactically equal to C the method removes the goal, otherwise an error is signaled.

Together, the full proof script has the form

```

apply (rule  $r_{n-1}$ )
...
apply(rule  $r_1$ )
apply(assumption)
done

```

If the example from Section 1.4.2 is proven this way the theorem is written together with its proof as

```

theorem  $x < 5 \implies 2*x+3 \leq 2*5 + 3$  for  $x :: nat$ 
  apply(rule example2)
  apply(rule example1)
  apply(assumption)
done

```

Note that the assumption A of the initial goal must be reached exactly by the sequence of rule applications. If it is replaced in the example by the stronger assumption $x < 3$ the rule applications will lead to the goal $x < 3 \implies x < 5$ which is trivial for the human reader but not applicable to the *assumption* method.

Using the *rule* Method for Forward Reasoning Steps

Assume that during a proof as described in Section 1.4.2 the intermediate fact sequence $F_1 \implies \dots \implies F_i$ has already been constructed by forward reasoning, so that the next step is to state fact F_{i+1} and prove it. Using method *rule* the step can be started by

```

have  $F_{i+1}$  proof (rule  $r_i$ )

```

The goal of this subproof is simply F_{i+1} , so applying the *rule* method with r_i will result in the new goal RA_i' which is F_i , as above. The subproof is not finished, since its goal state is not empty. But the goal is an already known fact. The proof method

fact name

can be used to remove that goal. The method only affects the first goal. If the fact referred by *name* unifies with it, the goal is removed, otherwise an error is signaled.

Using this method the forward reasoning step can be completed as

```

have  $F_{i+1}$  proof (rule  $r_i$ ) qed (fact  $f_i$ )

```

if F_i has been named f_i . This can be abbreviated (see Section 1.4.3) to

```

have  $F_{i+1}$  by (rule  $r_i$ ) (fact  $f_i$ )

```

Therefore the fact sequence $F_1 \implies \dots \implies F_n$ of the complete proof can be constructed by the structured proof of the form

```
proof –
assume  $f_1: F_1$ 
have  $f_2: F_2$  by (rule  $r_1$ ) (fact  $f_1$ )
...
have  $f_{n-1}: F_{n-1}$  by (rule  $r_{n-2}$ ) (fact  $f_{n-2}$ )
show  $F_n$  by (rule  $r_{n-1}$ ) (fact  $f_{n-1}$ )
qed
```

where the last fact F_n is usually the conclusion C and can be specified as *?thesis*.

If the example from Section 1.4.2 is proven this way the theorem is written together with its proof as

```
theorem  $x < 5 \implies 2*x+3 \leq 2*5 + 3$  for  $x :: nat$ 
proof –
  assume  $f_1: x < 5$ 
  have  $f_2: 2*x \leq 2*5$  by (rule example1) (fact  $f_1$ )
  show ?thesis by (rule example2) (fact  $f_2$ )
qed
```

The *fact* method can be specified in the form

fact

without naming the fact to be used. Then it selects a fact automatically. It uses the first fact from the proof context which unifies with the goal. If there is no such fact in the proof context an error is signaled.

Thus the example can be written without naming the facts as

```
theorem  $x < 5 \implies 2*x+3 \leq 2*5 + 3$  for  $x :: nat$ 
proof –
  assume  $x < 5$ 
  have  $2*x \leq 2*5$  by (rule example1) fact
  show ?thesis by (rule example2) fact
qed
```

Input Facts for the *rule* Method

If input facts F_1, \dots, F_n are passed to the *rule* method, they are used to remove assumptions from the rule applied by the method. If the rule has the form $\llbracket RA_1; \dots; RA_{n+m} \rrbracket \implies RC$ and every fact F_i unifies with assumption RA_i the first n assumptions are removed and the rule becomes

$\llbracket RA_{n+1}; \dots; RA_{n+m} \rrbracket \implies RC$. Then it is applied to the first goal in the usual way. If there are more facts than assumptions or if a fact does not unify, an error is signaled.

This allows to establish the connection from a fact F_i to F_{i+1} in a fact chain by a forward reasoning step of the form

from f_i **have** F_{i+1} **by** (*rule* r_i)

where f_i names the fact F_i . When it is input to the goal statement it is passed to the *rule* method and removes the assumption from the applied rule $RA_i \implies RC_i$, resulting in the assumption-less “rule” RC_i . When it is applied to the goal F_{i+1} it unifies and removes the goal, thus the subproof is complete.

For the fact sequence chaining can be used to write a structured proof without naming the facts:

```
proof –
assume  $F_1$ 
then have  $F_2$  by (rule  $r_1$ )
...
then have  $F_{n-1}$  by (rule  $r_{n-2}$ )
then show  $F_n$  by (rule  $r_{n-1}$ )
qed
```

If the example from Section 1.4.2 is proven this way the theorem is written together with its proof as

```
theorem  $x < 5 \implies 2*x+3 \leq 2*5 + 3$  for  $x :: nat$ 
proof –
  assume  $x < 5$ 
  then have  $2*x \leq 2*5$  by (rule example1)
  then show ?thesis by (rule example2)
qed
```

The Method *this*

Rule application can also be done by the method

this

Instead of applying a named method, it applies the input fact as rule to the first goal.

If several input facts are given, the method applies them exactly in the given order. Therefore the fact sequence can also be constructed by a structured proof of the form:


```

proof –
assume  $F_1$ 
with  $r_1$  have  $F_2$  by this
...
with  $r_{n-2}$  have  $F_{n-1}$  by this
with  $r_{n-1}$  show  $F_n$  by this
qed

```

The **with** statement inserts the explicitly named facts *before* the current facts. Therefore every goal statement for F_i gets as input the rule r_{i-1} followed by the chained fact F_{i-1} . The method *this* first applies the rule which replaces the goal by F_{i-1} . Then it applies the fact F_{i-1} as rule to this goal which removes it and finishes the subproof.

The proof

by *this*

can be abbreviated by . (a single dot).

Therefore the example from Section 1.4.2 can also be proven in the form

```

theorem  $x < 5 \implies 2*x+3 \leq 2*5 + 3$  for  $x :: nat$ 
proof –
  assume  $x < 5$ 
  with example1 have  $2*x \leq 2*5$  .
  with example2 show ?thesis .
qed

```

Adding Input Facts to the Goal

Another way of using the previous fact F_i in the proof of F_{i+1} is to add it as assumption to the goal. That is what the empty method does (see Section 1.5.1) if F_i is passed as input.

Thus the forward reasoning step can be written in the form

```

from  $f_i$  have  $F_{i+1}$ 
proof –
  assume  $F_i$ 
  then show ?thesis by (rule  $r_i$ )
qed

```

The empty method – replaces the goal F_{i+1} by the goal $F_i \implies F_{i+1}$ which is then proven by a forward reasoning step.

Of course this is much longer than the proof using the *rule* method immediately. But it shows generally how the step from F_i to F_{i+1} can be established using an arbitrary structured proof. That can be useful if no direct derivation rule r_i is available.

Automatic Rule Selection

The *rule* method can be specified in the form

rule

without naming the rule to be applied. Then it selects a rule automatically. It uses the first rule from the dynamic fact set *intro* for which the conclusion unifies with the goal conclusion. If there is no such rule in the set an error is signaled.

If the rules *example1* and *example2* would be in the *intro* set, the example proof could be written as

```
theorem  $x < 5 \implies 2*x+3 \leq 2*5 + 3$  for  $x :: nat$   
proof –  
  assume  $x < 5$   
  then have  $2*x \leq 2*5$  by rule  
  then show ?thesis by rule  
qed
```

However, the set *intro* is intended by Isabelle for a specific kind of rules called “introduction rules”. In such a rule the toplevel operator of the conclusion does not occur in any assumption, hence it is “introduced” by the rule. When an introduction rule is applied backwards, the operator is removed from the goal. This can be iterated to “deconstruct” the goal, some proofs can be written using this technique, however, the content of the set must be designed very carefully to not run into cycles.

Since in the rule *example2* the toplevel operator \leq occurs in the assumption it is not an introduction rule and should not be added to *intro*. Rule *example1* is an introduction rule but would interfere with predefined rules in the set.

The *standard* Method

The method

standard

is a method alias which can be varied for different Isabelle applications. Usually it is an alias for the *rule* method.

The *standard* method is the default, if no method is specified as the initial step in a structured proof. Thus

```
proof  $FS_1 \dots FS_n$  qed
```

is an abbreviation for

proof *standard* $FS_1 \dots FS_n$ **qed**

Note that the *standard* method will usually affect the goal by applying an introduction rule to it. That may be useful in some cases, but it has to be taken into account when writing the forward reasoning steps of the proof. For example, in Isabelle HOL there is an introduction rule in *intro* which splits a logical conjunction into two subgoals (see Section 2.3.1). Therefore the proof of a conjunction $P \wedge Q$ can be written

```
proof
show  $P$  <proof>
show  $Q$  <proof>
qed
```

because the goal is split by the *standard* method. If the empty method is used instead, the proof has to be of the form

```
proof –
show  $P \wedge Q$  <proof>
qed
```

because the goal is not modified.

In the abbreviated form **by** *method* of a structured proof the method cannot be omitted, but the proof **by** *standard* can be abbreviated to

..

(two dots). It can be used as complete proof for a proposition which can be proven by a single automatic rule application. Hence, if the rules *example1* and *example2* would be in the *intro* set, the example proof could be further abbreviated as

```
theorem  $x < 5 \implies 2*x+3 \leq 2*5 + 3$  for  $x :: nat$ 
proof –
  assume  $x < 5$ 
  then have  $2*x \leq 2*5$  ..
  then show ?thesis ..
qed
```

1.5.3 Composed Proof Methods

Proof methods can be composed from simpler methods with the help of “method expressions”. A method expression has one of the following forms:

- m_1, \dots, m_n : a sequence of methods which are applied in their order,

- $m_1; \dots; m_n$: a sequence of methods where each is applied to the goals created by the previous method,
- $m_1 | \dots | m_n$: a sequence of methods where only the first applicable method is applied,
- $m[n]$: the method m is applied to the first n goals,
- $m?$: the method m is applied if it is applicable,
- $m+$: the method m is applied once and then repeated as long as it is applicable.

Parentheses are used to structure and nest composed methods.

Composed methods can be used to combine backward reasoning steps to a single step. Using composed methods the example backward reasoning proof from Section 1.5.2 can be written as

```
theorem  $x < 5 \implies 2*x+3 \leq 2*5 + 3$  for  $x :: nat$ 
  apply(rule example2,rule example1,assumption)
done
```

In particular, it is possible to apply an arbitrarily complex backward reasoning step as the first method in a structured proof. Using composed methods the first example forward reasoning proof can be written

```
theorem  $x < 5 \implies 2*x+3 \leq 2*5 + 3$  for  $x :: nat$ 
proof –
  assume  $x < 5$ 
  have  $2*x \leq 2*5$  by (rule example1,fact)
  show ?thesis by (rule example2,fact)
qed
```

1.5.4 The Simplifier

A common proof technique is “rewriting”. If it is known that a term a is equal to a term b , some occurrences of a in a proposition can be replaced by b without changing the validity of the proposition.

Equality of two terms a and b can be expressed by the proposition $a = b$. If that proposition has been proven to be valid, i.e., is a fact, a can be substituted by b and vice versa in goals during a proof.

The *subst* Method

Rewriting is performed by the method

subst name

where *name* references an equality fact. The method only affects the first goal. If the referenced fact has the form $a = b$ the method replaces the first occurrence of a in the goal conclusion by b . The order of the terms in the equality fact matters, the method always substitutes the term on the left by that on the right.

If the equality contains unknowns unification is used: a is unified with every sub-term of the goal conclusion, the first match is replaced by b' which is b after substituting unknowns in the same way as in a . If there is no match of a in the goal conclusion an error is signaled.

For a goal $\llbracket A_1; \dots; A_n \rrbracket \implies C$ the method only rewrites in the conclusion C . The first match in the assumptions $A_1 \dots A_n$ can be substituted by the form

subst (asm) name

If not only the first match shall be substituted, a number of the match or a range of numbers may be specified in both forms as in

subst (asm) (i..j) name

The equality fact can also be a meta equality of the form $a \equiv b$. Therefore the method can be used to expand constant definitions. After the definition

definition *inc* $x \equiv x + 1$

the method *subst inc-def* will rewrite the first occurrence of a function application (*inc t*) in the goal conclusion to $(t + 1)$. Remember from Section 1.3.4 that the defining equation is automatically named *inc-def*. Note the use of unification to handle the actual argument term t .

The equality fact may be conditional, i.e., it may be a derivation rule with assumptions of the form $\llbracket RA_1; \dots; RA_m \rrbracket \implies a = b$. When the *subst* method applies a conditional equation of this form to a goal $\llbracket A_1; \dots; A_n \rrbracket \implies C$, it adds the goals $\llbracket A_1; \dots; A_n \rrbracket \implies RA_i'$ to the goal state after rewriting, where RA_i' result from RA_i by the unification of a in C . These goals are inserted before the original goal, so the next method application will usually process the goal $\llbracket A_1; \dots; A_n \rrbracket \implies RA_1'$.

As an example if there are theorems

theorem *eq1*: $n = 10 \implies n+3 = 13$ **for** $n::nat$ *<proof>*

theorem *eq2*: $n = 5 \implies 2*n = 10$ **for** $n::nat$ *<proof>*

the method *subst* (2) *eq2* replaces the goal $(x::nat) < 5 \implies 2*x+3 \leq 2*5 + 3$ by the goals

$$\begin{aligned} x < 5 &\implies 5 = 5 \\ x < 5 &\implies 2 * x + 3 \leq 10 + 3 \end{aligned}$$

where the first is trivial (but still must be removed by applying a rule). The second goal is replaced by the method *subst* (2) *eq1* by

$$\begin{aligned} x < 5 &\implies 10 = 10 \\ x < 5 &\implies 2 * x + 3 \leq 13 \end{aligned}$$

Note that the method *subst eq2* would unify $2*n$ with the first match $2*x$ in the original goal and replace it by

$$\begin{aligned} x < 5 &\implies x = 5 \\ x < 5 &\implies 10 + 3 \leq 2 * 5 + 3 \end{aligned}$$

where the first goal cannot be proven because it is invalid.

Simplification

If the term b in an equation $a = b$ is in some sense “simpler” than a , the goal will also become simpler by successful rewriting with the equation. If there are several such equations a goal can be replaced by successively simpler goals by rewriting with these equations. This technique can contribute to the goal’s proof and is called “simplification”.

Basically, simplification uses a set of equations and searches an equation in the set where the left hand side unifies with a sub-term in the goal, then substitutes it. This step is repeated until no sub-term in the goal unifies with a left hand side in an equation in the set.

It is apparent that great care must be taken when populating the set of equations, otherwise simplification may not terminate. If two equations $a = b$ and $b = a$ are in the set simplification will exchange matching terms forever. If an equation $a = a+0$ is in the set, a term matching a will be replaced by an ever growing sum with zeroes.

Simplification with a set of definitional equations from constant definitions (see Section 1.3.4) always terminates. Since constant definitions cannot be recursive, every substitution removes one occurrence of a defined constant from the goal. Simplification terminates if no defined constant from the set remains in the goal. Although the resulting goal usually is larger than the original goal, it is simpler in the sense that it uses fewer defined constants.

If the set contains conditional equations, simplification may produce additional goals. Then simplification is applied to these goals as well. Together, simplification may turn a single complex goal into a large number of simple

goals, but it cannot reduce the number of goals. Therefore simplification is usually complemented by methods which remove trivial goals like $x = x$, $A \implies A$, and *True*. Such an extended simplification may completely solve and remove the goal to which it is applied.

The *simp* Method

Isabelle supports simplification by the method

simp

which is also called “the simplifier”. It uses the dynamic fact set *simp* as the set of equations, which is also called “the simpset”. The method only affects the first goal. If no equation in *simp* is applicable to it or it is not modified by the applicable equations an error is signaled.

The *simp* method simplifies the whole goal, i.e., it applies rewriting to the conclusion and to all assumptions.

The simpset may contain facts which are not directly equations, but can be converted to an equation. In particular, an arbitrary derivation rule $\llbracket A_1; \dots; A_n \rrbracket \implies C$ can always be converted to the equation $\llbracket A_1; \dots; A_n \rrbracket \implies C = \text{True}$. The simplifier performs this conversion if no other conversion technique applies, therefore the simpset may actually contain arbitrary facts.

The *simp* method also detects several forms of trivial goals and removes them. Thus a complete proof may be performed by a single application of the simplifier in the form

by *simp*

In Isabelle HOL (see Section 2) the simpset is populated with a large number of facts which make the simplifier a very useful proof tool. Actually all examples of facts used in the previous sections can be proven by the simplifier:

theorem *example1*: $(x::\text{nat}) < c \implies n*x \leq n*c$ **by** *simp*

theorem *example2*: $(x::\text{nat}) \leq c \implies x + m \leq c + m$ **by** *simp*

theorem $(x::\text{nat}) < 5 \implies 2*x+3 \leq 2*5 + 3$ **by** *simp*

theorem *eq1*: $n = 10 \implies n+3 = 13$ **for** $n::\text{nat}$ **by** *simp*

theorem *eq2*: $n = 5 \implies 2*n = 10$ **for** $n::\text{nat}$ **by** *simp*

Configuring the Simplifier

The simplifier can be configured by modifying the equations it uses. The form

simp add: $name_1 \dots name_n$

uses the facts $name_1, \dots, name_n$ in addition to the facts in the simpset for its rewriting steps. The form

simp del: $name_1 \dots name_n$

uses only the facts from the simpset without the facts $name_1, \dots, name_n$, and the form

simp only: $name_1 \dots name_n$

uses only the facts $name_1, \dots, name_n$. The three forms can be arbitrarily combined.

As usual, a theorem may be added permanently to the simpset as described in Section 1.3.7 by specifying it as

theorem [*simp*]: *prop* <*proof*>

and the defining equation of a definition can be added by

definition *name::type* **where** [*simp*]: *name* \equiv *term*

Adding own constant definitions to the simplifier is a common technique to expand the definition during simplification. However, this may also have a negative effect: If an equation has been specified using the defined constant, it is no more applicable for rewriting after expanding the definition. Note that the facts in the simpset and the facts provided by *add*:, *del*:, and *only*: are not simplified themselves, the defined constant will not be expanded there.

Therefore it is usually not recommended to add defining equations to the simpset permanently. Instead, they can be specified by *add*: when they really shall be expanded during simplification.

Input Facts for the *simp* Method

As usual, facts may be input to the *simp* method. Like the empty method (see Section 1.5.1) it inserts these facts as assumptions into the goal, before it starts simplification. Since simplification is also applied to the assumptions, the input facts will be simplified as well.

As a possible effect of this behavior, after simplifying an input fact and the goal conclusion the results may unify, leading to the situation where the goal is removed by the *assumption* method (see Section 1.5.2). This is also done by the simplifier, hence in this way the input fact may contribute to prove the goal.

The *simp-all* Method

The method

simp-all

behaves like the *simp* method but processes all goals. It inserts input facts to all goals in the goal state and simplifies them. If it fails for all goals an error is signaled. Otherwise it simplifies only the goals for which it does not fail. If it achieves to remove all goals the proof is finished.

The *simp-all* method can be configured by *add:*, *del:*, and *only:* in the same way as the *simp* method.

The *simp-all* method is useful, if first a method *method* is applied to the goal which splits it into several subgoals which all can be solved by simplification. Then the complete proof can be written as

by *method simp-all*

Debugging the Simplifier

If the simplifier fails, it may be difficult to find out the reason. There are several debugging techniques which may help.

The content of the simpset can be displayed by the command

print-simpset

which may be written in the proof text in modes *proof(prove)* and *proof(state)* and outside of proofs. In the interactive editor the result is displayed in the Output panel (see Section 1.2.2).

There is also a simplifier trace which displays the successful rewrite steps. It is activated by the command

declare *[[simp-trace-new depth=n]]*

outside a theorem or definition. The number *n* should be atleast 2. When the cursor is positioned on an application of the *simp* method the button “Show trace” can be used in the Simplifier Trace Panel to display the trace in a separate window. See the documentation for more information about how to use the trace.

Another technique is to replace the *simp* method by a sequence of *subst* method applications and explicitly specify the equations which should have been used. To do this for a structured proof, replace it by a proof script for the *subst* methods.

1.5.5 Other Automatic Proof Methods

Isabelle provides several other proof methods which internally perform several steps, like the simplifier.

Automatic Methods

The following list contains automatic methods other than *simp*:

- *blast* mainly applies logical rules and can be used to solve complex logical formulas.
- *clarify* is similar but does not split goals and does not follow unsafe paths. It can be used to show the problem if *blast* fails.
- *auto* combines logical rule application with simplification. It processes all goals and leaves those it cannot solve.
- *clarsimp* combines *clarify* with simplification. It processes only the first goal and usually does not split goals.
- *fastforce* uses more techniques than *auto*, but processes only the first goal.
- *force* uses even more techniques and tries to solve the first goal.

The methods which do simplification can be configured like the simplifier by adding specifications *simp add:*, *simp del:*, and *simp only:*. For example, additional simplification rules can be specified for the *auto* method in the form

auto simp add: name₁ ... name_n

For more information about these methods see the Isabelle documentation.

Trying Methods

Instead of manually trying several automatic methods it is possible to specify the command

try

anywhere in mode *proof*(*prove*), i.e. at the beginning of a proof or in a proof script. It will try many automatic proof methods and describe the result in the Output window. It may take some time until results are displayed, in particular, if the goal is invalid and cannot be proven.

If **try** finds a proof for one or more goals it displays them as single proof methods, which, by clicking on them can be copied to the cursor position in the text area. The **try** command must be removed manually.

If **try** tells you that the goal can be “directly solved” by some fact, you can prove it by the *fact* method, but that also means that there is already a fact of the same form and your theorem is redundant.

It may also be the case that **try** finds a counterexample, meaning that the goal is invalid and cannot be proven.

Chapter 2

Isabelle HOL

2.1 Type Definitions

- typedef
- datatype
- record

2.2 Terms

- let
- case
- if

2.3 Types

2.3.1 Boolean Values

- $\wedge, \vee, \neg, \longrightarrow$
- $=, \neq, \longleftrightarrow$
- \forall, \exists
- intro

- 2.3.2** Sets
- 2.3.3** Natural Numbers
- 2.3.4** Tuples
- 2.3.5** Lists
- 2.3.6** Fixed-Size Binary Numbers